
Bootstrap a Kubernetes Cluster Using Kubeadm

[Edition 20]

[Last Update 211115]

For any issues/help contact : support@k21academy.com

K21 Academy

Contents

1	Introduction	4
2	Documentation.....	5
2.1	Kubernetes Documentation.....	5
3	Pre-Requisite	6
4	K8S Cluster Setup Overview	7
4.1	Ports Required between Master & Worker Node.....	7
4.2	Create Master & Worker Node Linux Machines.....	9
5	Create 3 virtual Machines to set up Kubeadm cluster.....	10
5.1	Create Master Node In East Region	10
5.2	Create Worker Node One In Central US	18
5.3	Create Worker Node Two In Central US.....	23
5.4	Virtual Network Peering	29
5.5	Cost overview in Pay-As-You-Go Subscription	36
6	Installing Docker, Kubeadm and Other Kubectl Packages	37
7	Configure Firewall, Update Iptables Settings, Disable SELinux and Disable SWAP	41
7.1	Configure Firewall.....	41
7.2	Update Iptables Settings.....	42
7.3	Disable SELinux	43
7.4	Disable SWAP	43
8	Kubeadm To Create And Initialise a Cluster	44
9	Using Kubeadm To Join Worker Nodes To The Cluster	45
10	(Optional) Unsubscribe Pay As you Go In Azure.....	49
11	TroubleShooting	52
11.1	Getting Warning while Configure cgroup driver used by kubelet on control-plane	52
11.2	Getting Error while Connecting node to Master node	52
11.3	Running kubectl get nodes and getting server localhost:8080.....	53
11.4	Error when Running Kubeadm init Command.....	53
11.5	Using Expired Token When Joining Node	54
11.6	Error when Running Kubeadm init Command.....	54
11.7	Unable to Ping Master node to Worker after Installing Docker	56
11.8	Kubelet is not running	62
11.9	Running kubectl get nodes and getting server 10.1.0.4:6443 was refused	63
11.10	Getting error while running Update command after adding the Kubernetes apt repository.....	64
11.11	Getting x509 error while build cluster on prem linux server	65
12	Summary.....	66

K21 Academy

1 INTRODUCTION

A **Kubernetes cluster** is a set of node machines for running containerized applications. If you're running **Kubernetes**, you're running a **cluster**. At a minimum, a **cluster** contains a control plane and one or more compute machines, or nodes.

This guide Covers:

- Bootstrap (Install & Configure) a Kubernetes Cluster Using Kubeadm

K21 Academy

2 DOCUMENTATION

2.1 Kubernetes Documentation

1. Installing Kubeadm <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>
2. Bootstrapping clusters with kubeadm
<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/>
3. Creating a single control-plane cluster with kubeadm
<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>
4. Cancel your Azure subscription
<https://docs.microsoft.com/en-us/azure/cost-management-billing/manage/cancel-azure-subscription>

3 PRE-REQUISITE

Ensure that you have completed following three activity guides (or you have an Ubuntu Server)

- Create account (Trial or Paid) on Azure Cloud.

Note: Follow Activity Guide

*Register_For_Azure_Cloud_Account_Accessing_Console_ed** from portal*

K21 Academy

4 K8S CLUSTER SETUP OVERVIEW

4.1 Ports Required between Master & Worker Node

Ideally in production setup, communication between Master & Worker nodes should be allowed only for required ports as per screenshot given below. You can also refer to K8S documentation at <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

Control-plane node(s)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	6443*	Kubernetes API server	All
TCP	Inbound	2379-2380	etcd server client API	kube-apiserver, etcd
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	10251	kube-scheduler	Self
TCP	Inbound	10252	kube-controller-manager	Self

Worker node(s)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	30000-32767	NodePort Services†	All

In our labs, we are going to open all ports (as this is play environment) between Master & Worker Node in Azure as given in screenshot below and also covered later in respective section in this guide.

Home > Virtual machines >

Virtual machines

Default Directory

+ Add ▾ 🕒 Reservations ...

Filter by name...

Name ↑↓

BackupVM ...

master ...

worker-01 ...

worker-02 ...

master | Networking

Virtual machine

Search (Ctrl+ /)

Attach network i

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Settings

Networking

Connect

Disks

Size

Security

Advisor recommendations

Extensions

Continuous delivery

Availability + scaling

Configuration

Identity

Properties

IP configuration

Network Interface

Inbound port rule

Network security

Impacts 0 subnets

Add inbound p

Priority

300

65000

65001

65500

Add inbound security rule

master-nsg

Basic

Source * ⓘ

Any

Source port ranges * ⓘ

*

Destination * ⓘ

Any

Destination port ranges * ⓘ

*

Protocol * ⓘ

Any TCP UDP ICMP

Action *

Allow Deny

Priority * ⓘ

310

Name * ⓘ

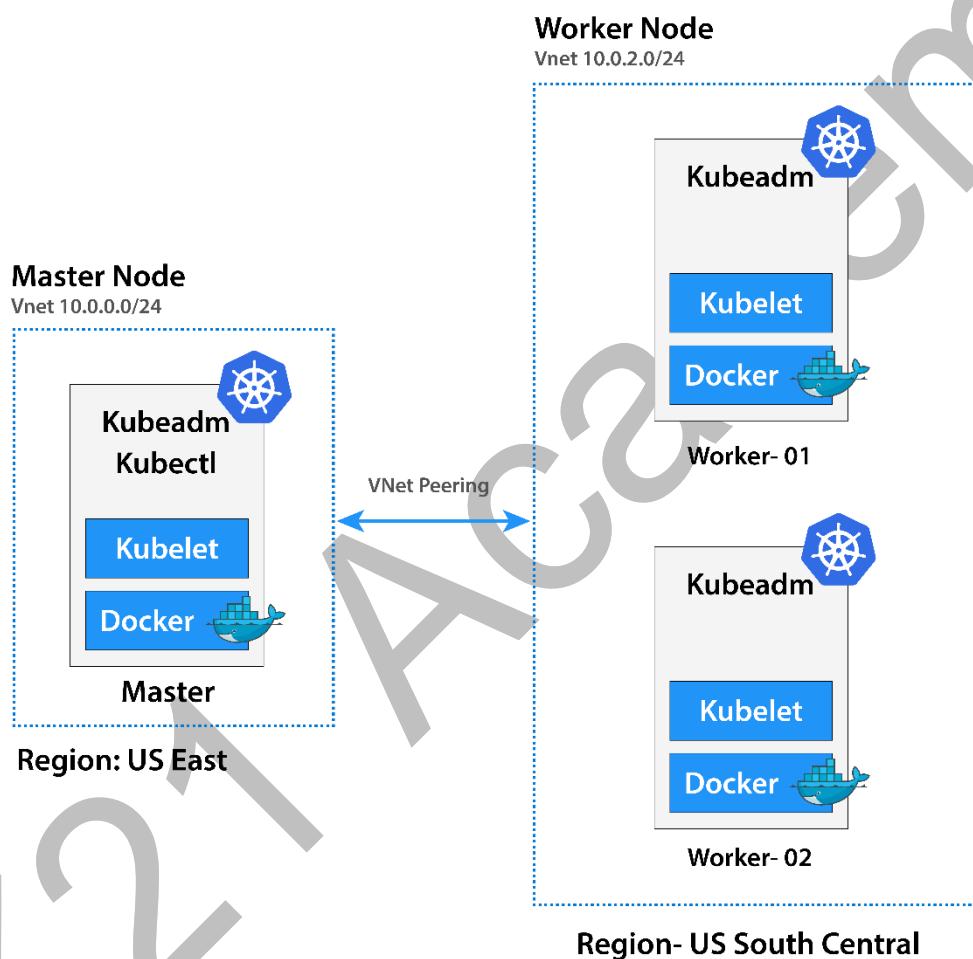
Port_All

Description

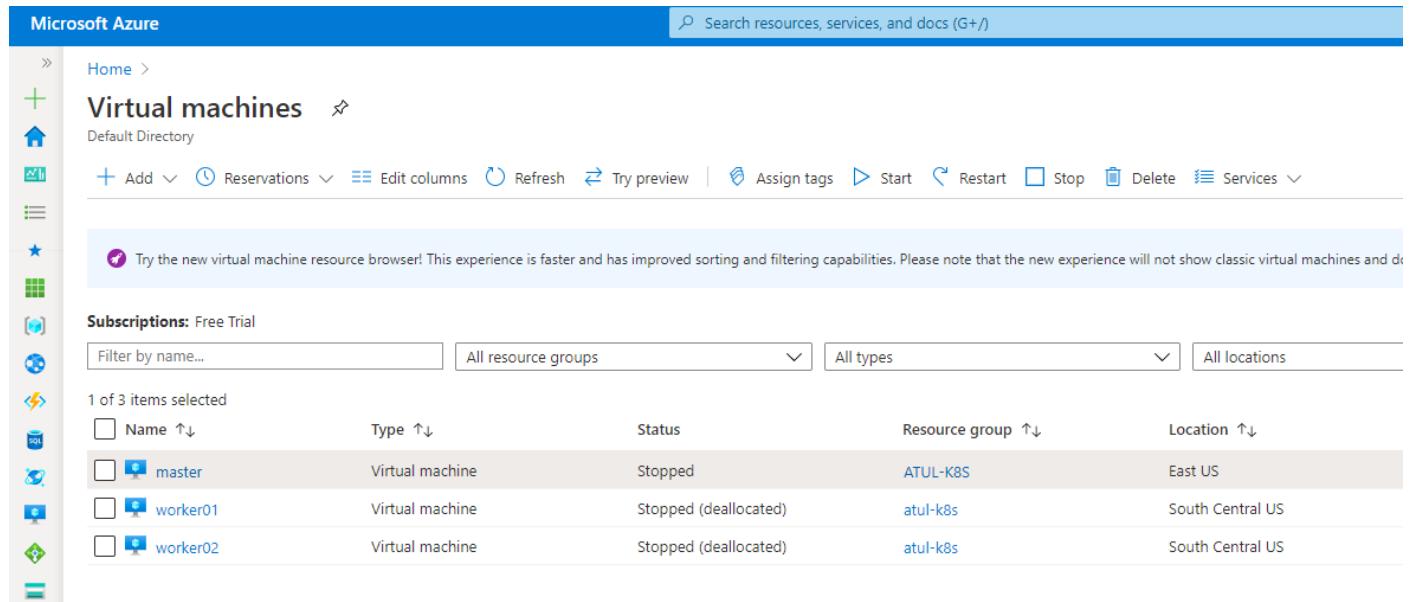
Add

4.2 Create Master & Worker Node Linux Machines

Ideally to ensure latency between Master & Worker Node is kept to minimum, All cluster members must be created in single Region. In FREE Azure trial account, there is a limit on number of Azure machine you can create hence we create Master Node in one Region and Worker Node in another region and peer these two networks using Azure VNET peering so master & worker nodes can communicate as per screenshot below (Peering is covered later in this guide)



5 CREATE 3 VIRTUAL MACHINES TO SET UP KUBEADM CLUSTER



The screenshot shows the Microsoft Azure Virtual machines dashboard. It displays three virtual machines: 'master' (Virtual machine, Stopped, ATUL-K8S, East US), 'worker01' (Virtual machine, Stopped (deallocated), atul-k8s, South Central US), and 'worker02' (Virtual machine, Stopped (deallocated), atul-k8s, South Central US). A search bar at the top right says 'Search resources, services, and docs (G+/-)'. The left sidebar has icons for Home, Virtual machines, Default Directory, and other Azure services.

Important Note: In this exercise we are going to create three Virtual Machines. In Azure Free tier account we can't create 3 virtual machine in a Single region due to service limit so we are creating One Master node in East US Region and Two Worker node (worker-1, worker-2 in South central US Region then we connect Vnets using VCN Peering so for now you don't need to convert your account to pay-as you go)

To Bootstrap a Kubernetes Cluster Using Kubeadm the recommended is to Create 3 new ubuntu VM in Azure Cloud.

Note: One will Work as Master node and other two will work as a Worker Nodes.

- master
- worker-1
- worker-2

Note: Resource Group should be same for all Machines i.e: kubeadm

5.1 Create Master Node In East Region

1. Create a Virtual Machine

Home >

Virtual machines

Default Directory

 Add  Reservations  Edit columns  Refresh  Try preview |  Assign tags  Start  Restart  Stop  Delete  Services

 Try the new virtual machine resource browser! This experience is faster and has improved sorting and filtering capabilities. Please note that the new experience will not show classic virtual machines and does not include support for some columns such as maintenance status.

Subscriptions: Free Trial

Filter by name...	All resource groups	All types	All locations	All tags	No grouping
5 items					
<input type="checkbox"/> Name ↑↓	Type ↑↓	Status	Resource group ↑↓	Location ↑↓	Source
<input type="checkbox"/>                                              <img alt					

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Azure for Students

Resource group * ⓘ

(New) Resource group

[Create new](#)

Instance details

Virtual machine name * ⓘ

A resource group is a container that holds related resources for an Azure solution.

Region * ⓘ

Name *

kubeadm

Availability options ⓘ

Image * ⓘ

[Browse all public and private images](#)

2. Instance details

- Virtual machine name: **master <Name of your virtual machine>**
- Region: (US) East US
- Availability Options: Leave default
- Image: Leave default

Create a virtual machine

[Basics](#) [Disks](#) [Networking](#) [Management](#) [Advanced](#) [Tags](#) [Review + create](#)

Create a virtual machine that runs Linux or Windows. Select an image from Azure marketplace or use your own customized image. Complete the Basics tab then Review + create to provision a virtual machine with default parameters or review each tab for full customization. [Learn more](#) ↗

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Azure for Students

Resource group * ⓘ

(New) kubeadm

[Create new](#)

Instance details

Virtual machine name * ⓘ

master

Region * ⓘ

(US) East US

Availability options ⓘ

No infrastructure redundancy required

Image * ⓘ

Ubuntu Server 18.04 LTS

[Browse all public and private images](#)

Azure Spot instance ⓘ

Yes No

Note: Use latest Image - Ubuntu Server 18.04 LTS – Gen 1

Microsoft Azure [Upgrade](#) Search resources, services, and docs (G+)

Home > Virtual machines >

Create a virtual machine

Basics Disks Networking Management Advanced Tags Review + create

Create a virtual machine that runs Linux or Windows. Select an image from Azure marketplace or use your own customized image. Complete the Basics tab then Review + create to provision a virtual machine with default parameters or review each tab for full customization. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription *	Free Trial
Resource group *	(New) atul_k8s
	Create new

Instance details

Virtual machine name * master

Region * (US) East US

Availability options No infrastructure redundancy required

Image * Ubuntu Server 18.04 LTS - Gen1

Azure Spot instance Yes No

Size * Standard_D2s_v3 - 2 vcpus, 8 GiB memory (₹4,632.03/month)

[Select size](#)

Administrator account

[Review + create](#) < Previous Next : Disks >

K21Academy

Create a virtual machine

Size * ⓘ

Standard_D2s_v3 - 2 vcpus, 8 GiB memory (₹4,632.03/month)

Select size

Administrator account

Authentication type ⓘ

SSH public key Password

Username * ⓘ

ubuntu



Password * ⓘ



Confirm password * ⓘ



Inbound port rules

Select which virtual machine network ports are accessible from the public internet. You can specify more limited or granular network access on the Networking tab.

Public inbound ports * ⓘ

None Allow selected ports

Select inbound ports *

SSH (22)



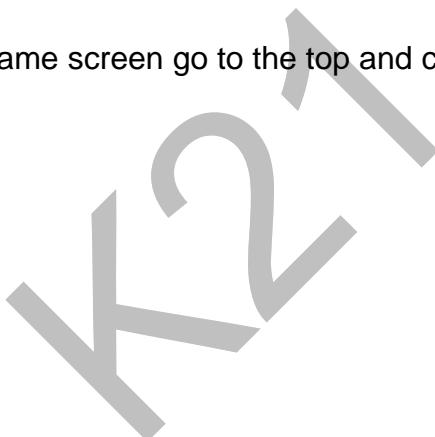
⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to create rules to limit inbound traffic to known IP addresses.

Review + create

< Previous

Next : Disks >

Same screen go to the top and click>Select **Networking**



Create a virtual machine

Basics

Disks

Networking

Management

Advanced

Tags

Review + create

Create a virtual machine that runs Linux or Windows. Select an image from Azure marketplace or use your own customized image. Complete the Basics tab then Review + create to provision a virtual machine with default parameters or review each tab for full customization. [Learn more ↗](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Azure for Students

Resource group * ⓘ

(New) kubeadm

[Create new](#)

Instance details

Virtual machine name * ⓘ

master

Region * ⓘ

(US) East US

Availability options ⓘ

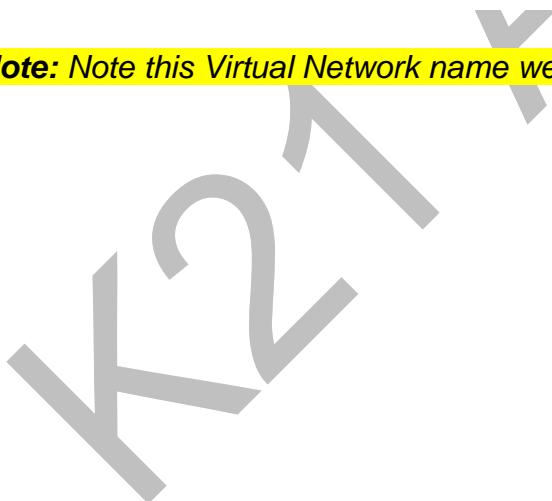
No infrastructure redundancy required

Image * ⓘ

Ubuntu Server 18.04 LTS

[Browse all public and private images](#)

Note: Note this Virtual Network name we gonna need this in later sections.



- Click on **Review + Create**

Basics Disks **Networking** Management Advanced Tags Review + create

Define network connectivity for your virtual machine by configuring network interface card (NIC) settings. You can control ports, inbound and outbound connectivity with security group rules, or place behind an existing load balancing solution.

[Learn more](#)

Network interface

When creating a virtual machine, a network interface will be created for you.

Virtual network * ⓘ

(new) kubeadm-vnet

[Create new](#)

Subnet * ⓘ

(new) default (10.0.1.0/24)

Public IP ⓘ

(new) master-ip

[Create new](#)

NIC network security group ⓘ

None Basic Advanced

Public inbound ports * ⓘ

None Allow selected ports

Select inbound ports *

SSH (22)

⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to create rules to limit inbound traffic to known IP addresses.

[Review + create](#)

< Previous

Next : Management >

Home > Virtual machines >

Create a virtual machine

 Validation passed

Basics Disks Networking Management Advanced Tags **Review + create**

 **You have set SSH port(s) open to the internet.** This is only recommended for testing. If you want to change this setting, go back to Basics tab.

Basics

Subscription	Azure for Students
Resource group	(new) kubeadm
Virtual machine name	master
Region	East US
Availability options	No infrastructure redundancy required
Image	Ubuntu Server 18.04 LTS
Size	Standard D2s v3 (2 vcpus, 8 GiB memory)
Authentication type	Password
Username	ubuntu
Public inbound ports	SSH
Azure Spot	No

Create

< Previous

Next >

Download a template for automation

- Open all ports for further labs

Home > Virtual machines >

Virtual machines

Default Directory

+ Add Reservations ...

Filter by name...

Name ↑↓

BackupVM ...

master ...

worker-01 ...

worker-02 ...

master | Networking

Search (Ctrl+ /)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Settings

Networking

Connect

Disks

Size

Security

Advisor recommendations

Extensions

Continuous delivery

Availability + scaling

Configuration

Identity

Properties

Attach network interface

master442

IP configuration ipconfig1 (Primary)

Network Interface

Virtual network/subnet Accelerated network

Inbound port rules

Network security Impact to subnets

Add inbound rule

Priority

300

65000

65001

65500

...

Add

Add inbound security rule

master-nsg

Basic

Source * Any

Source port ranges * *

Destination * Any

Destination port ranges * *

Protocol * Any TCP UDP ICMP

Action * Allow Deny

Priority * 310

Name * Port_All

Description

Add

5.2 Create Worker Node One In Central US

1. Create a Virtual Machine

Home >

Virtual machines

Default Directory

+ Add

Reservations ...

Edit columns Refresh Try preview Assign tags Start Restart Stop Delete Services ...

Try the new virtual machine resource browser! This experience is faster and has improved sorting and filtering capabilities. Please note that the new experience will not show classic virtual machines and does not include support for some columns such as maintenance status.

Subscriptions: Free Trial

Filter by name... All resource groups All types All locations All tags No grouping

5 items

Name ↑↓	Type ↑↓	Status	Resource group ↑↓	Location ↑↓	Source	Maintenance status	Subscription ↑↓	...
...
...
...
...

2. Select Resource Group which you created earlier

3. Instance details

- Virtual machine name: **worker-01 <Name of your virtual machine>**

- Region: (US) South Central US
- Availability Options: Leave default
- Image: Leave default

Create a virtual machine

Basics Disks Networking Management Advanced Tags Review + create

Create a virtual machine that runs Linux or Windows. Select an image from Azure marketplace or use your own customized image. Complete the Basics tab then Review + create to provision a virtual machine with default parameters or review each tab for full customization. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription *	Azure for Students
Resource group *	kubeadm
	Create new

Instance details

Virtual machine name *	worker-01
Region *	(US) South Central US
Availability options	No infrastructure redundancy required
Image *	Ubuntu Server 18.04 LTS
Browse all public and private images	
Azure Spot instance	<input type="radio"/> Yes <input checked="" type="radio"/> No

Home > Virtual machines >

Create a virtual machine

Size * ⓘ

Standard_D2s_v3 - 2 vcpus, 8 GiB memory (₹5,307.53/month)

Select size

Administrator account

Authentication type ⓘ

SSH public key Password

Username * ⓘ

ubuntu

Password * ⓘ

.....

Confirm password * ⓘ

.....

Inbound port rules

Select which virtual machine network ports are accessible from the public internet. You can specify more limited or granular network access on the Networking tab.

Public inbound ports * ⓘ

None Allow selected ports

Select inbound ports *

SSH (22)

⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to create rules to limit inbound traffic to known IP addresses.

[Review + create](#)

< Previous

Next : Disks >

Same screen go to the top and click>Select **Networking**

Home > Virtual machines >

Create a virtual machine

Basics Disks Networking Management Advanced Tags Review + create

Define network connectivity for your virtual machine by configuring network interface card (NIC) settings. You can control ports, inbound and outbound connectivity with security group rules, or place behind an existing load balancing solution.

[Learn more](#)

Network interface

When creating a virtual machine, a network interface will be created for you.

Virtual network * ⓘ <input type="text" value="(new) kubeadmvnet485"/> ▼ Create new	Subnet * ⓘ <input type="text" value="(new) default (10.0.2.0/24)"/> ▼ Create new
Public IP ⓘ <input type="text" value="(new) worker-01-ip"/> ▼ Create new	
NIC network security group ⓘ <input type="radio"/> None <input checked="" type="radio"/> Basic <input type="radio"/> Advanced	
Public inbound ports * ⓘ <input type="radio"/> None <input checked="" type="radio"/> Allow selected ports	
Select inbound ports * <input type="text" value="SSH (22)"/> ▼	

⚠️ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to create rules to limit inbound traffic to known IP addresses.

Review + create

< Previous

Next : Management >

Note: Note this Virtual Network name we gonna need this in later sections.

- Click on **Review + Create**

[Home](#) > [Virtual machines](#) >

Create a virtual machine

 Validation passed

Basics Disks Networking Management Advanced Tags **Review + create**

 You have set SSH port(s) open to the internet. This is only recommended for testing. If you want to change this setting, go back to Basics tab.

Basics

Subscription	Azure for Students
Resource group	kubeadm
Virtual machine name	worker-01
Region	South Central US
Availability options	No infrastructure redundancy required
Image	Ubuntu Server 18.04 LTS
Size	Standard D2s v3 (2 vcpus, 8 GiB memory)
Authentication type	Password
Username	ubuntu
Public inbound ports	SSH
Azure Spot	No

[Disks](#)

[Create](#)

[< Previous](#)

[Next >](#)

[Download a template for automation](#)

- Open all ports for further labs

The screenshot shows the Azure portal interface for managing virtual machines. On the left, the 'Virtual machines' blade lists several VMs, with 'worker-01' selected and highlighted by a red box. The main pane displays the 'Networking' settings for 'worker-01'. Under 'Inbound port rules', there is a button labeled 'Add inbound p...' which is also highlighted by a red box. A large red arrow points from this button to the 'Add inbound security rule' dialog box on the right.

Add inbound security rule

worker-01-nsg

Basic

Source * Any

Source port ranges * * (highlighted by a red box)

Destination * Any

Destination port ranges * * (highlighted by a red box)

Protocol * Any (TCP, UDP, ICMP) (highlighted by a red box)

Action * Allow (highlighted by a red box)

Priority * 310

Name * Port_All (highlighted by a red box)

Description

Add (highlighted by a red box)

5.3 Create Worker Node Two In Central US

Create a virtual machine

Basics Disks Networking Management Advanced Tags Review + create

Create a virtual machine that runs Linux or Windows. Select an image from Azure marketplace or use your own customized image. Complete the Basics tab then Review + create to provision a virtual machine with default parameters or review each tab for full customization. [Learn more ↗](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Azure for Students

Resource group * ⓘ

kubeadm

[Create new](#)

Instance details

Virtual machine name * ⓘ

worker-02

Region * ⓘ

(US) South Central US

Availability options ⓘ

No infrastructure redundancy required

Image * ⓘ

Ubuntu Server 18.04 LTS

[Browse all public and private images](#)

Azure Spot instance ⓘ

Yes No

K21

Home > Virtual machines >

Create a virtual machine

Size * ⓘ

Standard_D2s_v3 - 2 vcpus, 8 GiB memory (₹5,307.53/month)



Select size

Administrator account

Authentication type ⓘ

SSH public key Password

Username * ⓘ

ubuntu



Password * ⓘ

.....



Confirm password * ⓘ

.....



Inbound port rules

Select which virtual machine network ports are accessible from the public internet. You can specify more limited or granular network access on the Networking tab.

Public inbound ports * ⓘ

None Allow selected ports

Select inbound ports *

SSH (22)

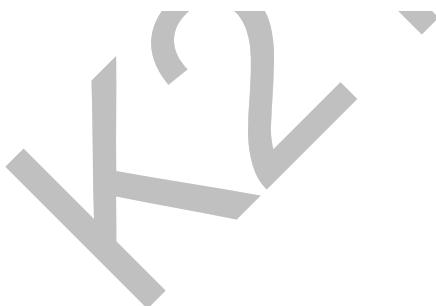


⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to create rules to limit inbound traffic to known IP addresses.

Review + create

< Previous

Next : Disks >



Create a virtual machine

Basics Disks **Networking** Management Advanced Tags Review + create

Define network connectivity for your virtual machine by configuring network interface card (NIC) settings. You can control ports, inbound and outbound connectivity with security group rules, or place behind an existing load balancing solution.

[Learn more](#)

Network interface

When creating a virtual machine, a network interface will be created for you

Virtual network * ⓘ

kubeadmvnet485 ▾

[Create new](#)

Subnet * ⓘ

default (10.0.2.0/24) ▾

[Manage subnet configuration](#)

Public IP ⓘ

(new) worker-02-ip ▾

[Create new](#)

NIC network security group ⓘ

None Basic Advanced

Public inbound ports * ⓘ

None Allow selected ports

Select inbound ports *

SSH (22) ▾

⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to create rules to limit inbound traffic to known IP addresses.

Review + create

< Previous

Next : Management >



Home > Virtual machines >

Create a virtual machine

✓ Validation passed

Basics Disks Networking Management Advanced Tags **Review + create**

⚠ You have set SSH port(s) open to the internet. This is only recommended for testing. If you want to change this setting, go back to Basics tab.

Basics

Subscription	Azure for Students
Resource group	kubeadm
Virtual machine name	worker-02
Region	South Central US
Availability options	No infrastructure redundancy required
Image	Ubuntu Server 18.04 LTS
Size	Standard D2s v3 (2 vcpus, 8 GiB memory)
Authentication type	Password
Username	ubuntu
Public inbound ports	SSH
Azure Spot	No

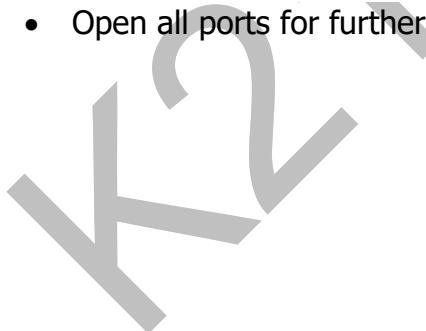
Create

< Previous

Next >

Download a template for automation

- Open all ports for further labs



Home > Virtual machines >

Virtual machines

Default Directory

+ Add Reservations ...

Filter by name...

<input type="checkbox"/> Name ↑↓
<input type="checkbox"/> BackupVM ...
<input type="checkbox"/> master ...
<input type="checkbox"/> worker-01 ...
<input type="checkbox"/> worker-02 ...

worker-02 | Networking

Virtual machine

Search (Ctrl+ /)

Attach network

worker-02106

IP configuration ⓘ ipconfig1 (Primary)

Network Interface Virtual network/subnet Accelerated network

Inbound port rules

Add inbound port rule

Priority

300

65000

65001

65500

Protocol Any TCP UDP ICMP

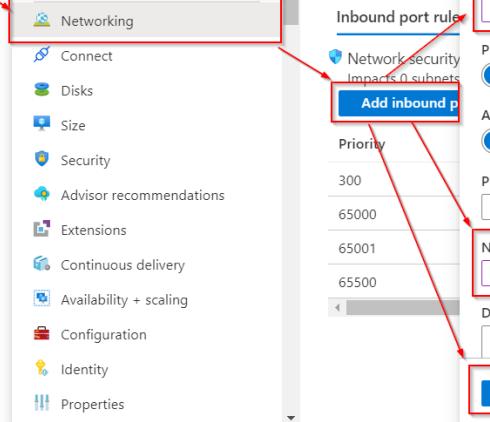
Action Allow Deny

Priority 310

Name Port_All

Description

Add



- Final Screen after Creating all the Machine

Home >

Virtual machines

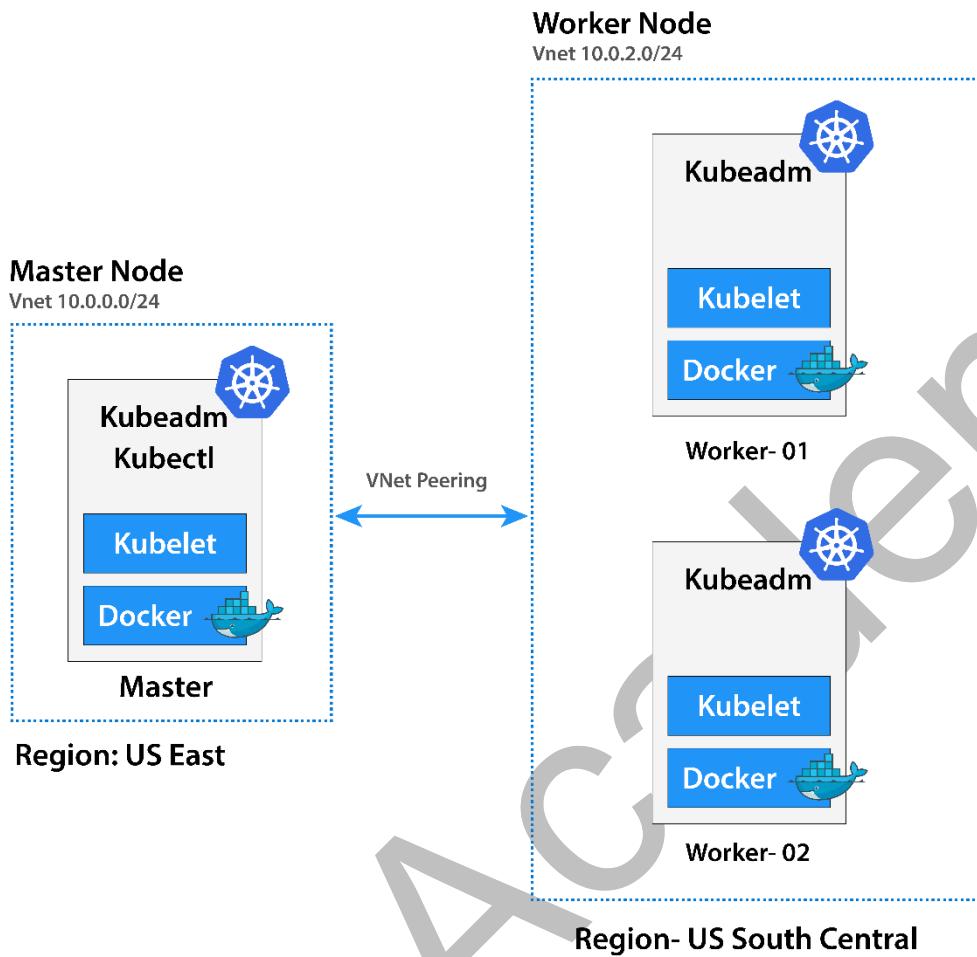
Default Directory

+ Add Reservations Edit columns Refresh Assign tags Start Restart Stop Delete Services

Subscriptions: 1 of 2 selected – Don't see a subscription? [Open Directory + Subscription settings](#)

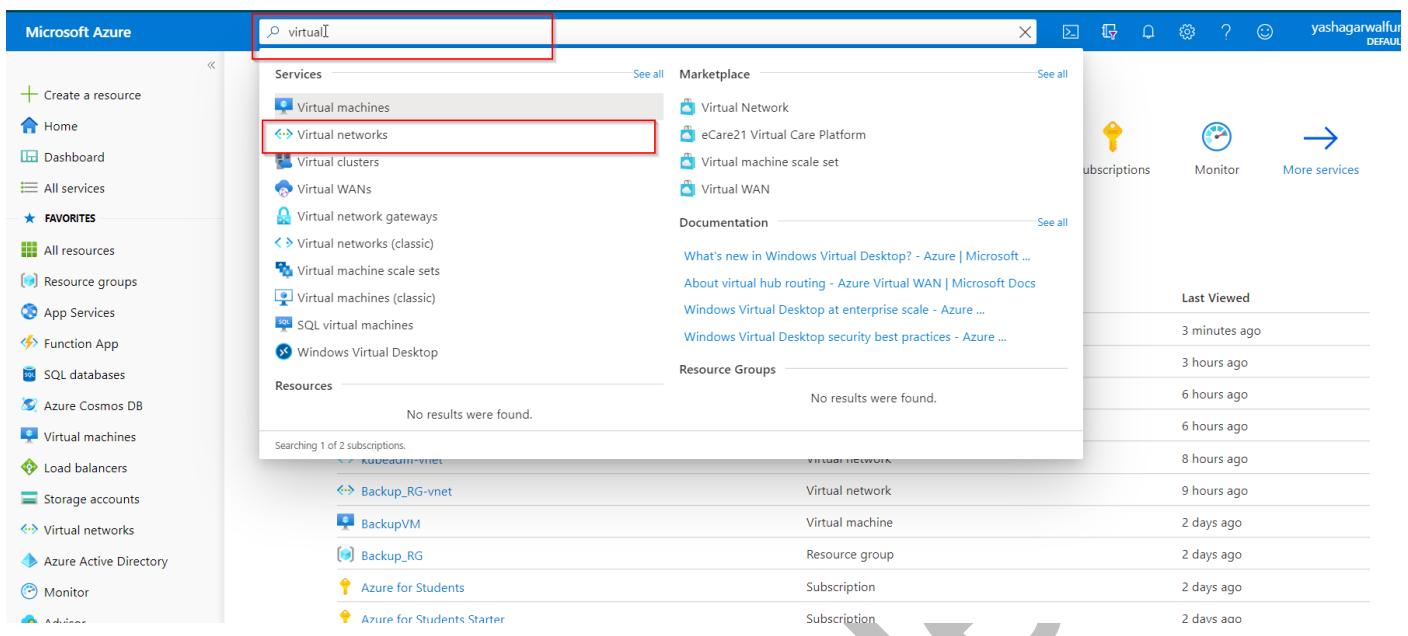
Filter by name...	Azure for Students	All resource groups	All types	All locations	All tags	No grouping
1 of 4 items selected						
<input type="checkbox"/> Name ↑↓	Type ↑↓	Status	Resource group ↑↓	Location ↑↓	Source	Maintenance status
<input type="checkbox"/> BackupVM	Virtual machine	Stopped (deallocated)	Backup_RG	East US	Marketplace	-
<input type="checkbox"/> master	Virtual machine	Running	kubeadm	East US	Marketplace	-
<input type="checkbox"/> worker-01	Virtual machine	Running	kubeadm	South Central US	Marketplace	-
<input type="checkbox"/> worker-02	Virtual machine	Running	KUBADM	South Central US	Marketplace	-

5.4 Virtual Network Peering



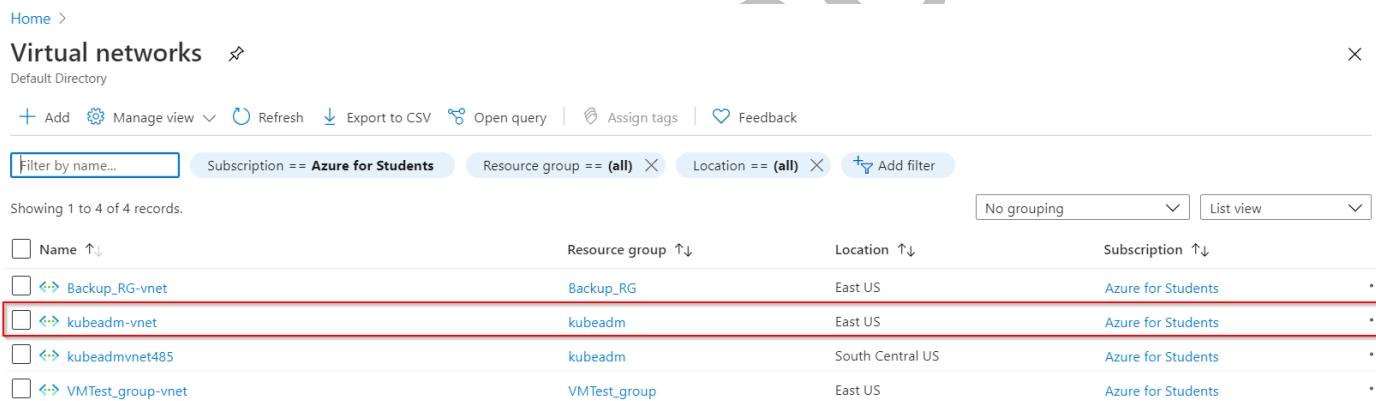
Important Note: Now We will perform virtual peering so the machines in different regions can talk to each other.

1. In search bar search for virtual network



The screenshot shows the Microsoft Azure portal interface. A search bar at the top contains the text "virtual". Below the search bar, the "Services" section is expanded, showing categories like Virtual machines, Virtual networks, Virtual clusters, etc. The "Virtual networks" category is highlighted with a red box. To the right of the search results, there's a sidebar with "Last Viewed" items such as "3 minutes ago", "6 hours ago", etc. At the bottom, there's a list of resources under "Virtual network" and "Subscription".

2. Select master virtual machine Virtual Network



The screenshot shows the "Virtual networks" blade in the Azure portal. It lists four virtual networks: "Backup_RG-vnet", "BackupVM", "Backup_RG", and "Azure for Students". The "Azure for Students" entry is highlighted with a red box. The blade includes a search bar, filter options (Subscription, Resource group, Location), and a table view with columns for Name, Resource group, Location, and Subscription.

Name	Resource group	Location	Subscription
Backup_RG-vnet	Backup_RG	East US	Azure for Students
Azure for Students	kubeadm	East US	Azure for Students
BackupVM	kubeadm	South Central US	Azure for Students
Backup_RG	VMTest_group	East US	Azure for Students

3. In Overview section go for Peering and click Add

Virtual networks

Default Directory

[+ Add](#) [Manage view](#) ...

Filter by name...

Name ↑↓

↳ kubeadm-vnet

↳ kubeadmvnet485

kubeadm-vnet | Peerings

Virtual network

Search (Ctrl+/)

[+ Add](#) [Refresh](#)

[Overview](#)

[Activity log](#)

[Access control \(IAM\)](#)

[Tags](#)

[Diagnose and solve problems](#)

Settings

[Address space](#)

[Connected devices](#)

[Subnets](#)

[DDoS protection](#)

[Firewall](#)

[Security](#)

[DNS servers](#)

[Peerings](#)

[Service endpoints](#)

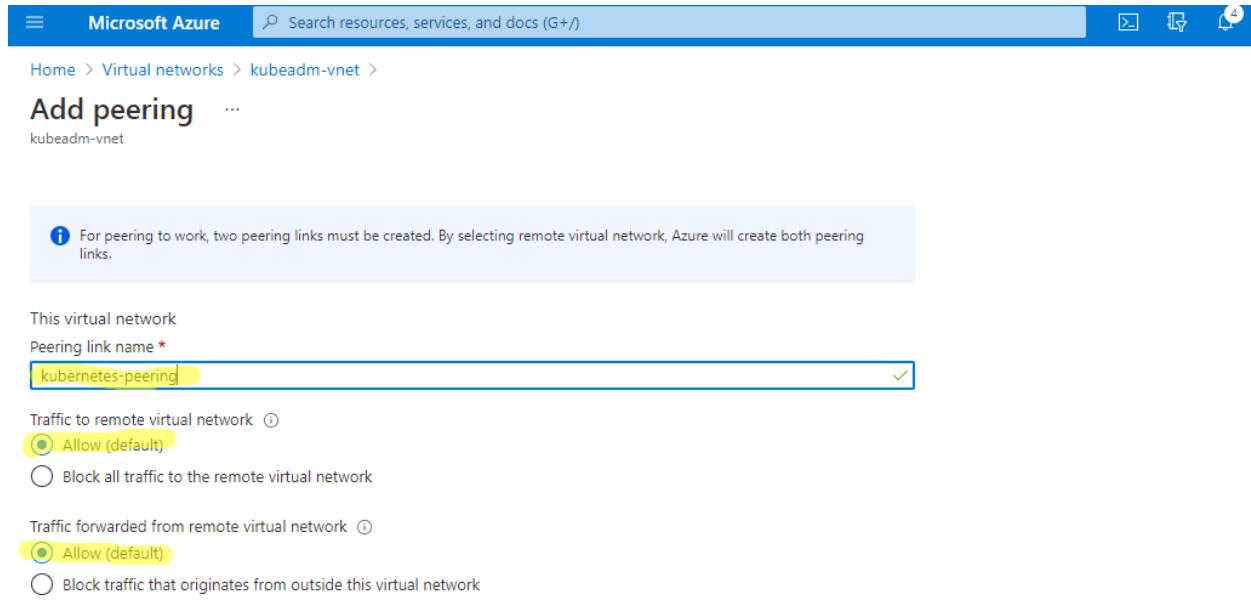
[Private endpoints](#)

[Properties](#)

4. On Add Peering screen Fill:

- **Peering link name:** kubeernetes-peering
- **Traffic to remote virtual network:** choose Allow (default)
- **Traffic forwarded from remote virtual network:** choose Allow (default)
- **Virtual network gateway:** Choose None (default)
- **Virtual Network:** Kubeadmvnet485 (*This is the other network that you want to peer, in our case network where our Worker node VMs are created*)

Note: Select same options as given in the Image.

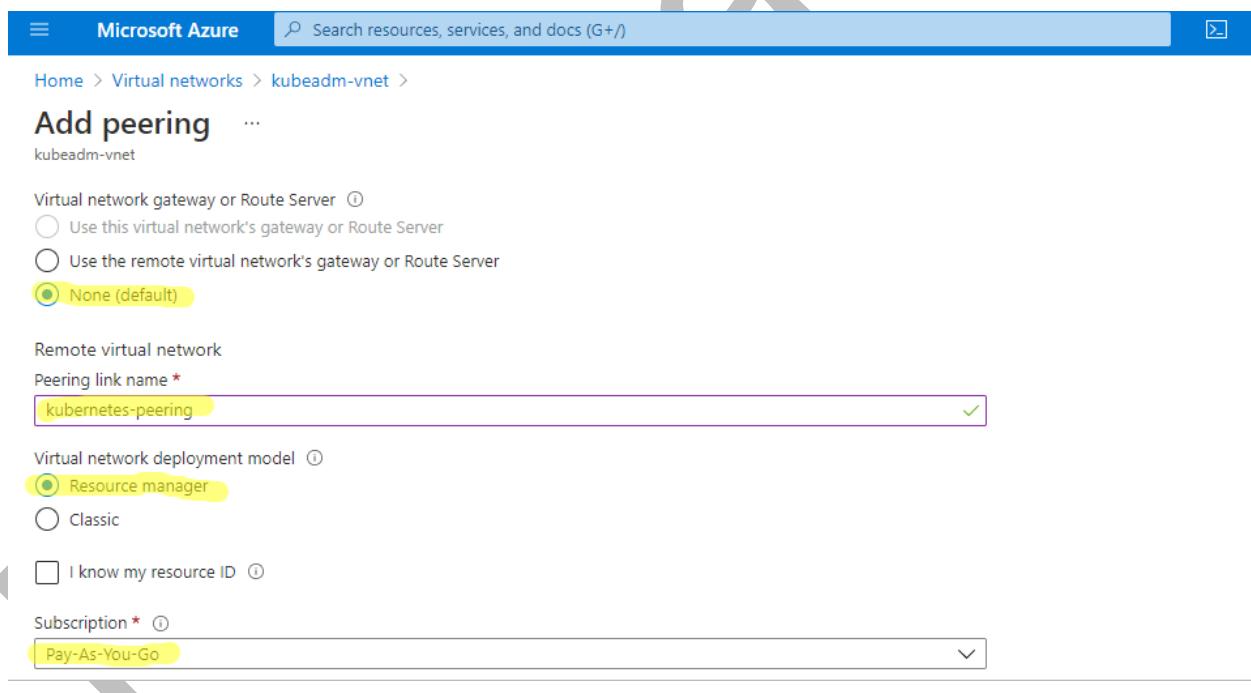


This virtual network

Peering link name *****
 ✓

Traffic to remote virtual network ⓘ
 Allow (default)
 Block all traffic to the remote virtual network

Traffic forwarded from remote virtual network ⓘ
 Allow (default)
 Block traffic that originates from outside this virtual network



Virtual network gateway or Route Server ⓘ
 Use this virtual network's gateway or Route Server
 Use the remote virtual network's gateway or Route Server
 None (default)

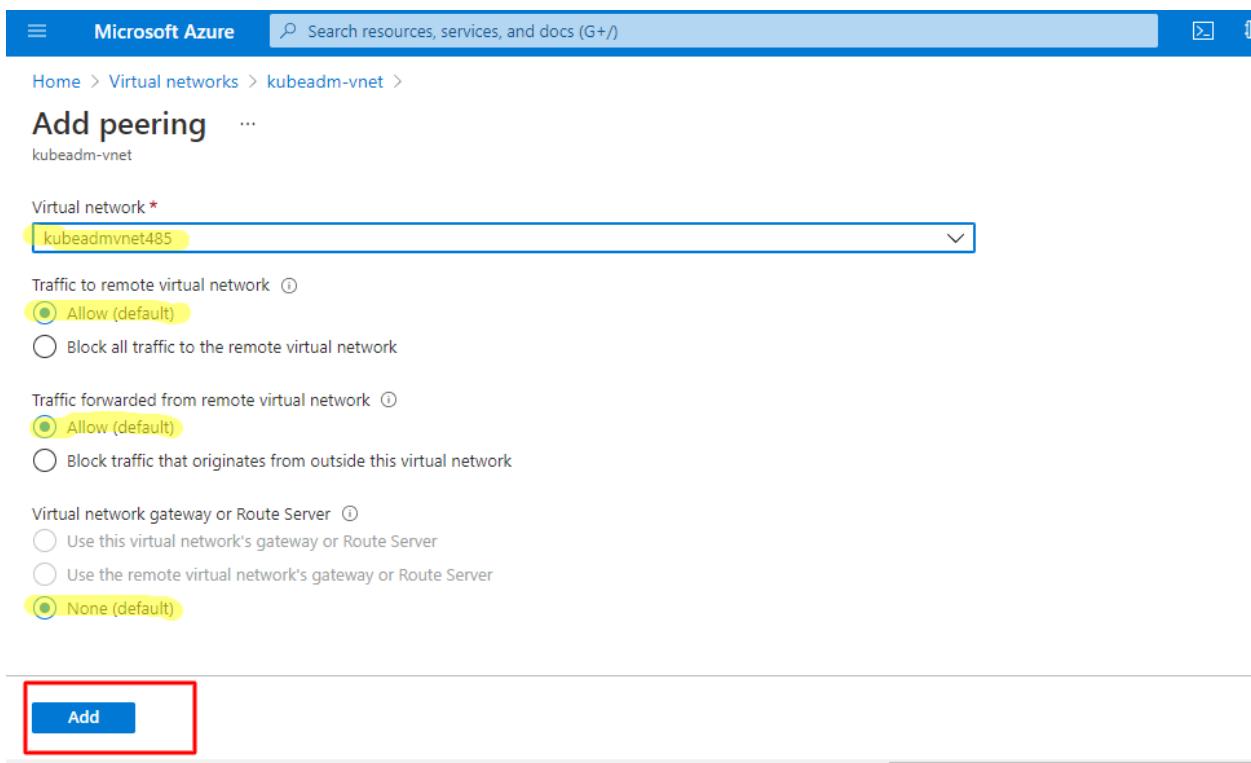
Remote virtual network

Peering link name *****
 ✓

Virtual network deployment model ⓘ
 Resource manager
 Classic

I know my resource ID ⓘ

Subscription ***** ⓘ
 ✓



Microsoft Azure Search resources, services, and docs (G+)

Home > Virtual networks > kubeadm-vnet >

Add peering

kubeadm-vnet

Virtual network *

kubeadmvnet485

Traffic to remote virtual network ⓘ

Allow (default)

Block all traffic to the remote virtual network

Traffic forwarded from remote virtual network ⓘ

Allow (default)

Block traffic that originates from outside this virtual network

Virtual network gateway or Route Server ⓘ

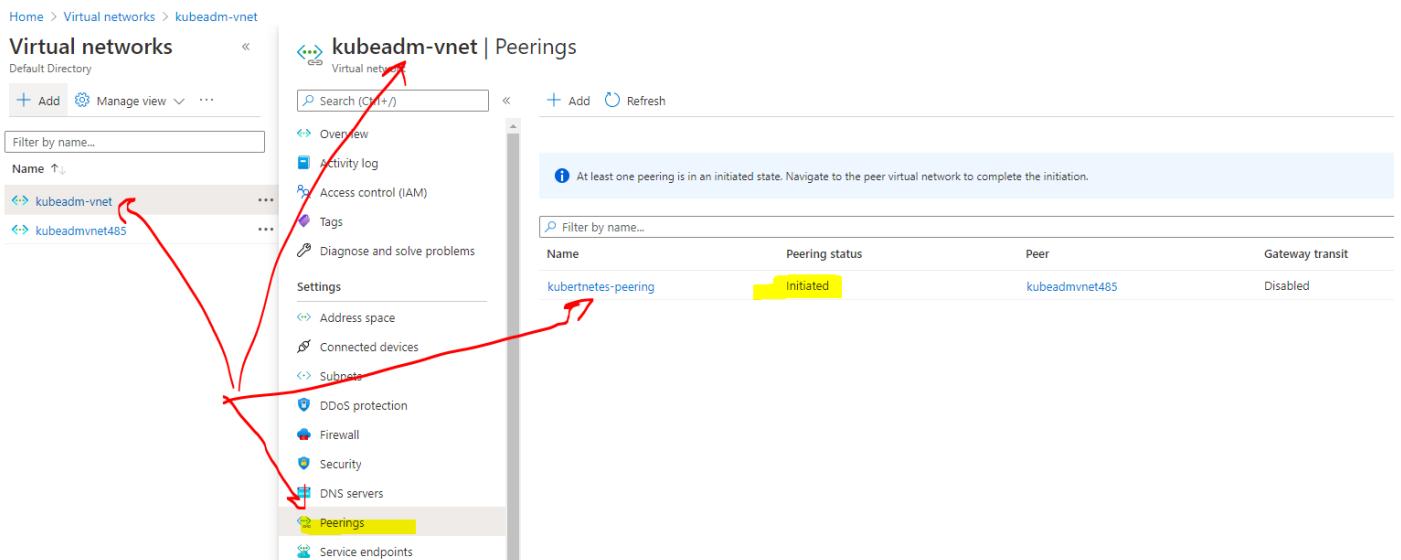
Use this virtual network's gateway or Route Server

Use the remote virtual network's gateway or Route Server

None (default)

Add

5. Click on add, you will see Peering status as **Initiated**



Home > Virtual networks > kubeadm-vnet

Virtual networks

Default Directory

+ Add Manage view ...

Filter by name...

Name	Peering status	Peer	Gateway transit
kubeadm-vnet	Initiated	kubeadmvnet485	Disabled
kubeadmvnet485	Pending	kubeadm-vnet	Enabled

kubeadm-vnet | Peering

Virtual network

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Address space

Connected devices

Subnets

DDoS protection

Firewall

Security

DNS servers

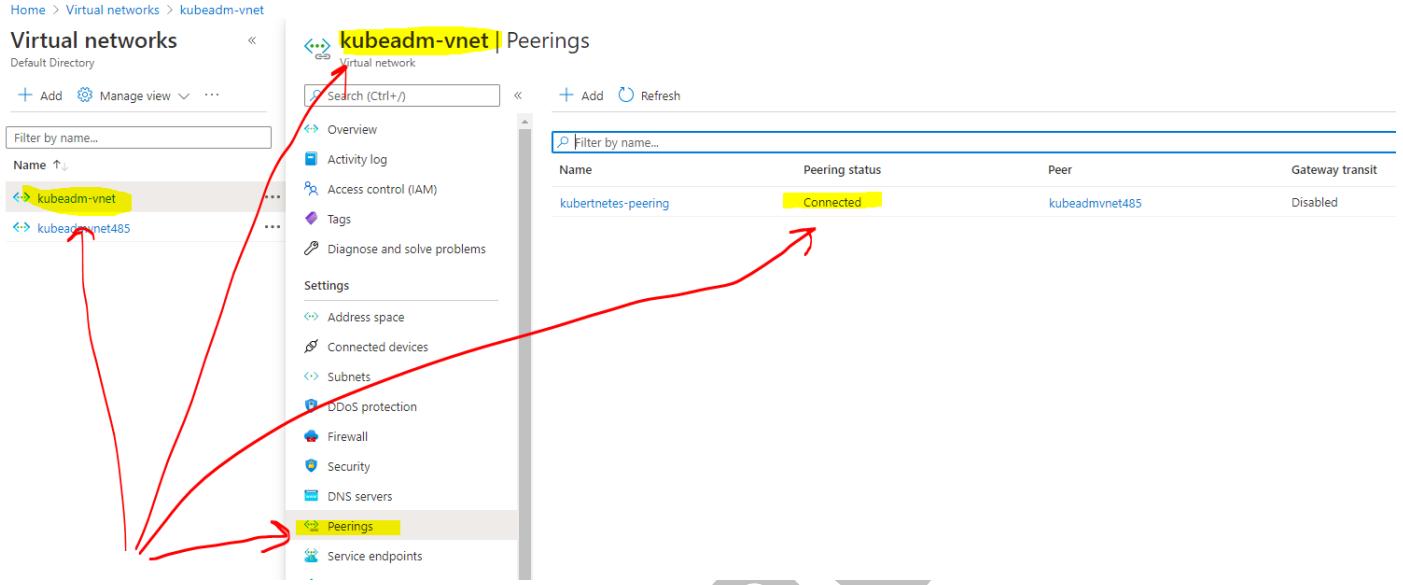
Peering

Service endpoints

At least one peering is in an initiated state. Navigate to the peer virtual network to complete the initiation.

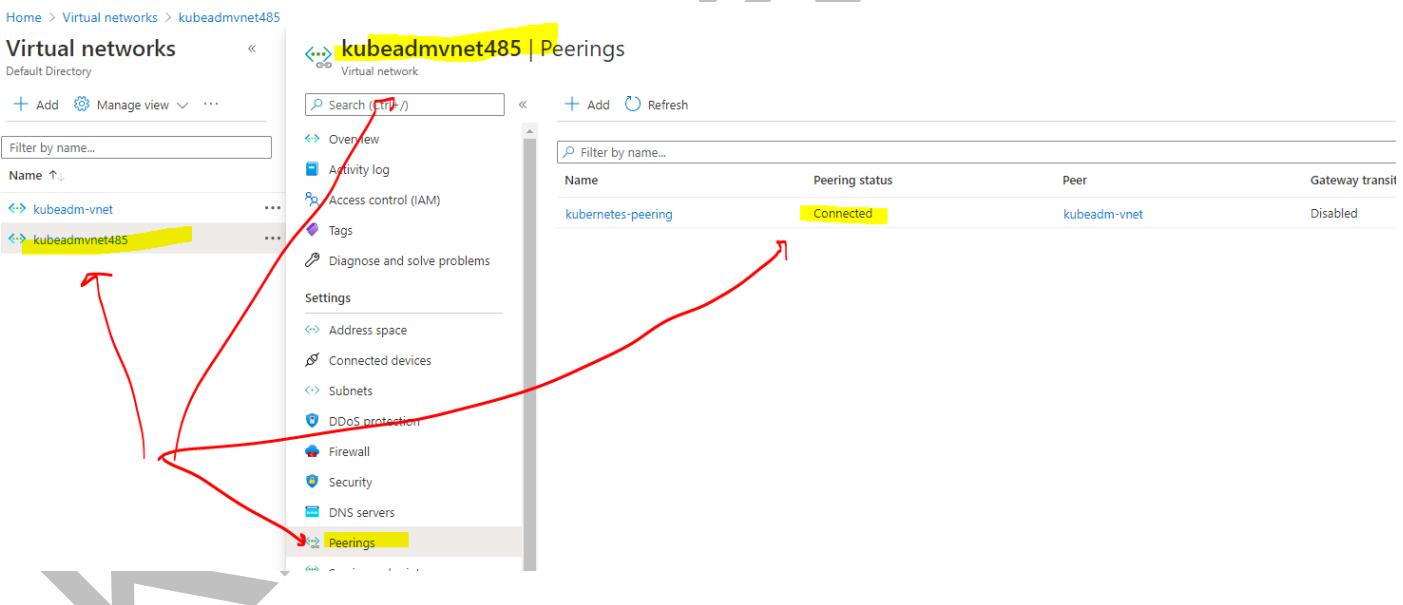
Name	Peering status	Peer	Gateway transit
kubernetes-peering	Initiated	kubeadmvnet485	Disabled

6. Ensure that status of Peering is Connected



The screenshot shows the Azure portal interface for managing virtual networks. On the left, there's a list of virtual networks: 'kubeadm-vnet' and 'kubeadmvnet485'. The 'kubeadmvnet485' item is highlighted with a yellow box and has a red arrow pointing to it from the bottom-left. On the right, the 'Peering' section is displayed for 'kubeadm-vnet'. It shows one peering entry:

Name	Peering status	Peer	Gateway transit
kubernetes-peering	Connected	kubeadmvnet485	Disabled

The screenshot shows the Azure portal interface for managing virtual networks. On the left, there's a list of virtual networks: 'kubeadm-vnet' and 'kubeadmvnet485'. The 'kubeadmvnet485' item is highlighted with a yellow box and has a red arrow pointing to it from the bottom-left. On the right, the 'Peering' section is displayed for 'kubeadmvnet485'. It shows one peering entry:

Name	Peering status	Peer	Gateway transit
kubernetes-peering	Connected	kubeadm-vnet	Disabled

7. Test connectivity by pinging machine's IP in different network

- a) Without peering these machines won't be able to ping each other over their private IPs

```
ubuntu@master:~$ ping 10.0.2.4
PING 10.0.2.4 (10.0.2.4) 56(84) bytes of data.
```

```
ubuntu@worker-01:~$ ping 10.0.1.4
PING 10.0.1.4 (10.0.1.4) 56(84) bytes of data.
```

- b) After peering enabled machines in different Network will be able to ping each other over their private IPs

Note: Here 10.0.2.4 is Private IP of worker node that is accessible from master node and 10.0.1.4 is Private IP of Master node that is accessible from worker node

```
ubuntu@master:~$ ping 10.0.2.4
PING 10.0.2.4 (10.0.2.4) 56(84) bytes of data.
64 bytes from 10.0.2.4: icmp_seq=1 ttl=64 time=37.7 ms
64 bytes from 10.0.2.4: icmp_seq=2 ttl=64 time=37.5 ms
64 bytes from 10.0.2.4: icmp_seq=3 ttl=64 time=37.6 ms
64 bytes from 10.0.2.4: icmp_seq=4 ttl=64 time=37.6 ms
64 bytes from 10.0.2.4: icmp_seq=5 ttl=64 time=38.2 ms
64 bytes from 10.0.2.4: icmp_seq=6 ttl=64 time=37.4 ms
64 bytes from 10.0.2.4: icmp_seq=7 ttl=64 time=37.4 ms
64 bytes from 10.0.2.4: icmp_seq=8 ttl=64 time=37.3 ms
64 bytes from 10.0.2.4: icmp_seq=9 ttl=64 time=37.4 ms
64 bytes from 10.0.2.4: icmp_seq=10 ttl=64 time=37.7 ms
64 bytes from 10.0.2.4: icmp_seq=11 ttl=64 time=37.6 ms
64 bytes from 10.0.2.4: icmp_seq=12 ttl=64 time=37.3 ms
64 bytes from 10.0.2.4: icmp_seq=13 ttl=64 time=37.5 ms
64 bytes from 10.0.2.4: icmp_seq=14 ttl=64 time=37.8 ms
64 bytes from 10.0.2.4: icmp_seq=15 ttl=64 time=37.5 ms
64 bytes from 10.0.2.4: icmp_seq=16 ttl=64 time=37.5 ms
3 packets transmitted, 0 received, 100% packet loss, time 2026ms
ubuntu@worker-01:~$ clear
ubuntu@worker-01:~$ ping 10.0.1.4
PING 10.0.1.4 (10.0.1.4) 56(84) bytes of data.
--- 10.0.1.4 ping statistics ---
333 packets transmitted, 0 received, 100% packet loss, time 749545ms
ubuntu@worker-01:~$ ping 10.0.1.4
PING 10.0.1.4 (10.0.1.4) 56(84) bytes of data.
64 bytes from 10.0.1.4: icmp_seq=1593 ttl=64 time=38.0 ms
64 bytes from 10.0.1.4: icmp_seq=1594 ttl=64 time=37.3 ms
64 bytes from 10.0.1.4: icmp_seq=1595 ttl=64 time=37.5 ms
64 bytes from 10.0.1.4: icmp_seq=1596 ttl=64 time=37.5 ms
64 bytes from 10.0.1.4: icmp_seq=1597 ttl=64 time=40.9 ms
64 bytes from 10.0.1.4: icmp_seq=1598 ttl=64 time=37.4 ms
64 bytes from 10.0.1.4: icmp_seq=1599 ttl=64 time=37.9 ms
64 bytes from 10.0.1.4: icmp_seq=1600 ttl=64 time=37.5 ms
64 bytes from 10.0.1.4: icmp_seq=1601 ttl=64 time=37.6 ms
```

This completes VNet peering

5.5 Cost overview in Pay-As-You-Go Subscription

When you create an account on azure cloud, you get free trial services for 1 month. After that you can create another free account using different email and payment option or you can choose Pay-As-You-Go service.

Here is the average cost overview, when you create a Kubernetes cluster with 1 master & 2 worker nodes, in Pay-As-You-Go service.

<input type="checkbox"/> Invoice ID	Billing Period	Invoice date	Total amount	Status	Type
<input type="checkbox"/>	9/28/2021 - 10/27/2021	10/29/2021	₹650.54	 Paid	Azure services
<input type="checkbox"/>	8/28/2021 - 9/27/2021	9/29/2021	₹381.26	 Paid	Azure services
<input type="checkbox"/>	7/28/2021 - 8/27/2021	8/29/2021	₹1,023.23	 Paid	Azure services
<input type="checkbox"/>	6/28/2021 - 7/27/2021	7/29/2021	₹2,402.34	 Paid	Azure services
<input type="checkbox"/>	5/28/2021 - 6/27/2021	6/29/2021	₹2,696.57	 Paid	Azure services
<input type="checkbox"/>	4/28/2021 - 5/27/2021	5/29/2021	₹2,043.14	 Paid	Azure services

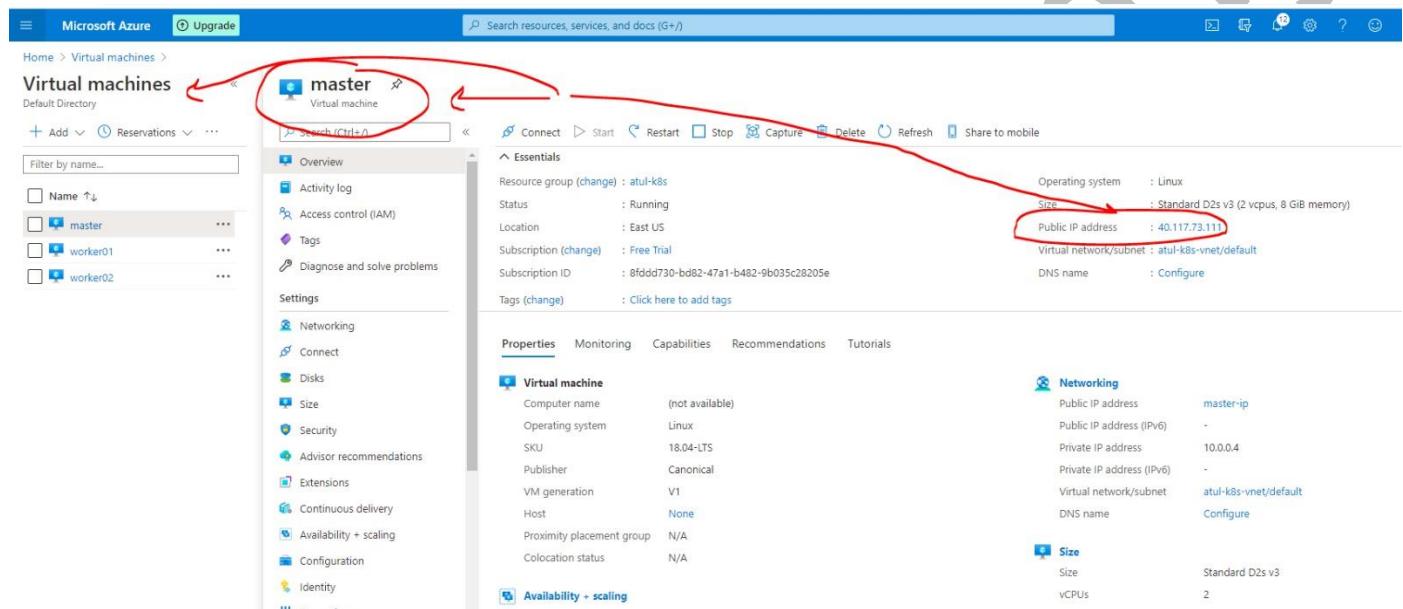
6 INSTALLING DOCKER, KUBEADM AND OTHER KUBECTL PACKAGES

Note: First perform All Section 6 Steps on the **Master node** then repeat same on the **both worker Nodes too.**

1. SSH to the virtual machine with the username and password you used while creating the VM

```
$ ssh root@publicipaddress
```

Note: To get the Public IP Address, go to the master node and copy the Public IP Address.



The screenshot shows the Microsoft Azure portal interface. In the left sidebar, under 'Virtual machines', the 'master' virtual machine is selected and highlighted with a red circle. On the right, the 'Overview' tab is active. In the 'Essentials' section, the 'Public IP address' field is circled in red and contains the value '40.117.73.111'. Other details shown include Resource group (atul-k8s), Status (Running), Location (East US), Subscription (Free Trial), and Tags.

Note: Linux or Mac user can use ssh command. Windows user can use Putty.

2. Switch to root user in case you aren't logged as root

```
$ sudo -i
```

```
ubuntu@kubeadm-master:~$  
ubuntu@kubeadm-master:~$ sudo su  
root@kubeadm-master:/home/ubuntu#
```

3. Install docker package using the following command

```
$ apt-get update && apt-get install -y docker.io
```

```
root@kubeadm-master:/home/ubuntu# apt-get update && apt-get install -y docker.io  
Hit:1 http://azure.archive.ubuntu.com/ubuntu bionic InRelease  
Get:2 http://azure.archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]  
Get:3 http://azure.archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
```

Verify the docker version installed

```
$ docker --version
```

4. Update the apt package index and install packages needed to use the Kubernetes apt repository:

```
$ apt-get update && apt-get install -y apt-transport-https ca-certificates curl
```

```
root@Master:~# sudo apt-get update
Hit:1 http://azure.archive.ubuntu.com/ubuntu bionic InRelease
Get:2 http://azure.archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
Get:3 http://azure.archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
Hit:5 http://security.ubuntu.com/ubuntu bionic-security InRelease
Get:4 https://packages.cloud.google.com/apt kubernetes-xenial InRelease [9383 B]
Get:6 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 Packages [48.5 kB]
Fetched 221 kB in 0s (459 kB/s)
Reading package lists... Done
root@Master:~# sudo apt-get install -y apt-transport-https ca-certificates curl
Reading package lists... Done
Building dependency tree
Reading state information... Done
ca-certificates is already the newest version (20210119~18.04.1).
ca-certificates set to manually installed.
apt-transport-https is already the newest version (1.6.14).
The following package was automatically installed and is no longer required:
  linux-headers-4.15.0-151
Use 'sudo apt autoremove' to remove it.
The following additional packages will be installed:
  libcurl4
The following packages will be upgraded:
  curl libcurl4
2 upgraded, 0 newly installed, 0 to remove and 13 not upgraded.
Need to get 378 kB of archives.
After this operation, 0 B of additional disk space will be used.
Get:1 http://azure.archive.ubuntu.com/ubuntu bionic-updates/main amd64 curl amd64 7.58.0-2ubuntu3.14 [159 kB]
Get:2 http://azure.archive.ubuntu.com/ubuntu bionic-updates/main amd64 libcurl4 amd64 7.58.0-2ubuntu3.14 [219 kB]
Fetched 378 kB in 0s (16.8 MB/s)
(Reading database ... 77278 files and directories currently installed.)
Preparing to unpack .../curl_7.58.0-2ubuntu3.14_amd64.deb ...
Unpacking curl (7.58.0-2ubuntu3.14) over (7.58.0-2ubuntu3.13) ...
Preparing to unpack .../libcurl4_7.58.0-2ubuntu3.14_amd64.deb ...
Unpacking libcurl4:amd64 (7.58.0-2ubuntu3.14) over (7.58.0-2ubuntu3.13) ...
Setting up libcurl4:amd64 (7.58.0-2ubuntu3.14) ...
Setting up curl (7.58.0-2ubuntu3.14) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Processing triggers for libc-bin (2.27-3ubuntu1.4) ...
```

5. Download the Google Cloud public signing key:

```
$ sudo curl -fsSLo /usr/share/keyrings/kubernetes-archive-keyring.gpg
https://packages.cloud.google.com/apt/doc/apt-key.gpg
```

6. Add the Kubernetes apt repository:

Note: Copy-Paste or type Next Command Carefully.
Copy it and paste it on Notepad and make it as a single line command.

```
$ echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

```
root@Master:~# echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main
```

7. Update apt package index, install kubelet, kubeadm and kubectl, and pin their version:

```
$ KUBE_VERSION=1.21.0
```

```
$ apt-get update
```

```
root@Master:~# KUBE_VERSION=1.21.0
root@Master:~# apt-get update
Hit:1 http://azure.archive.ubuntu.com/ubuntu bionic InRelease
Hit:2 http://azure.archive.ubuntu.com/ubuntu bionic-updates InRelease
Hit:3 http://azure.archive.ubuntu.com/ubuntu bionic-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu bionic-security InRelease
Get:5 https://packages.cloud.google.com/apt kubernetes-xenial InRelease [9383 B]
Err:5 https://packages.cloud.google.com/apt kubernetes-xenial InRelease
      The following signatures couldn't be verified because the public key is not available: NO_PUBKEY FEEA9169307EA071 NO_PUBKEY 8B57C5C2836F4BEB
Reading package lists... Done
W: GPG error: https://packages.cloud.google.com/apt kubernetes-xenial InRelease: The following signatures couldn't be verified because the public key is not available: NO_PUBKEY FEEA9169307EA071 NO_PUBKEY 8B57C5C2836F4BEB
E: The repository 'https://apt.kubernetes.io/ kubernetes-xenial InRelease' is not signed.
N: Updating from such a repository can't be done securely, and is therefore disabled by default.
N: See apt-secure(8) manpage for repository creation and user configuration details.
root@Master:~#
```

```
$ apt-get install -y kubelet=${KUBE_VERSION}-00 kubeadm=${KUBE_VERSION}-00
kubectl=${KUBE_VERSION}-00 kubernetes-cni=0.8.7-00
```

```
root@Master:~# apt-get install -y kubelet=${KUBE_VERSION}-00 kubeadm=${KUBE_VERSION}-00 kubectl=${KUBE_VERSION}-00 kubernetes-cni=0.8.7-00
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following package was automatically installed and is no longer required:
  linux-headers-4.15.0-151
Use 'apt autoremove' to remove it.
The following additional packages will be installed:
  conntrack cri-tools socat
The following NEW packages will be installed:
  conntrack cri-tools kubeadm kubelet kubernetes-cni socat
0 upgraded, 7 newly installed, 0 to remove and 13 not upgraded.
Need to get 70.4 MB of archives.
After this operation, 309 MB of additional disk space will be used.
Get:1 http://azure.archive.ubuntu.com/ubuntu bionic/main amd64 conntrack amd64 1:1.4.4+snapshot20161117-6ubuntu2 [30.6 kB]
Get:2 http://azure.archive.ubuntu.com/ubuntu bionic/main amd64 socat amd64 1.7.3.2-2ubuntu2 [342 kB]
Get:3 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 cri-tools amd64 1.13.0-01 [8775 kB]
Get:4 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 kubernetes-cni amd64 0.8.7-00 [25.0 MB]
Get:5 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 kubelet amd64 1.21.0-00 [18.8 MB]
Get:6 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 kubectl amd64 1.21.0-00 [8972 kB]
Get:7 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 kubeadm amd64 1.21.0-00 [8544 kB]
Fetched 70.4 MB in 2s (30.1 MB/s)
Selecting previously unselected package conntrack.
(Reading database ... 77278 files and directories currently installed.)
Preparing to unpack .../0-conntrack_1%3a1.4.4+snapshot20161117-6ubuntu2_amd64.deb ...
Unpacking conntrack (1:1.4.4+snapshot20161117-6ubuntu2) ...
Selecting previously unselected package cri-tools.
Preparing to unpack .../1-cri-tools_1.13.0-01_amd64.deb ...
Unpacking cri-tools (1.13.0-01) ...
Selecting previously unselected package kubernetes-cni.
Preparing to unpack .../2-kubernetes-cni 0.8.7-00 amd64.deb ...
```

```
$ apt-mark hold kubelet kubeadm kubectl
```

8. Configure cgroup driver and storage driver.

```
$ cat > /etc/docker/daemon.json <<EOF
```

```
{
```

```
  "exec-opts": ["native.cgroupdriver=systemd"],
```

```
"log-driver": "json-file",
"storage-driver": "overlay2"
}
EOF
$ mkdir -p /etc/systemd/system/docker.service.d
```

```
root@worker01:~# cat > /etc/docker/daemon.json <<EOF
> {
>   "exec-opts": ["native.cgroupdriver=systemd"],
>   "log-driver": "json-file",
>   "storage-driver": "overlay2"
> }
> EOF
root@worker01:~# mkdir -p /etc/systemd/system/docker.service.d
root@worker01:~#
```

9. Restart the docker

```
$ systemctl daemon-reload
$ systemctl restart docker && systemctl enable docker
```

10. Verify Storage driver and cgroup driver

```
$ docker info | grep -i "storage"
$ docker info | grep -i "cgroup"

root@Master:~# docker info | grep -i "storage"
Storage Driver: overlay2
WARNING: No swap limit support
root@Master:~# docker info | grep -i "cgroup"
WARNING: No swap limit support
Cgroup Driver: systemd
Cgroup Version: 1
```

11. Start the kubelet service is required on all the nodes

```
$ systemctl enable kubelet && systemctl start kubelet
```

7 CONFIGURE FIREWALL, UPDATE IPTABLES SETTINGS, DISABLE SELINUX AND DISABLE SWAP

NOTE: This section is optional. In Ubuntu all this steps are preconfigured so need to perform in ubuntu OS. But if you are using other linux distribution. You May require to use the following steps.

7.1 Configure Firewall

To fulfil their tasks, nodes, containers, and pods must be able to communicate across the cluster. Add the following ports by entering the listed commands.

On the Master Node enter

```
$ sudo firewall-cmd --permanent --add-port=6443/tcp
$ sudo firewall-cmd --permanent --add-port=2379-2380/tcp
$ sudo firewall-cmd --permanent --add-port=10250/tcp
$ sudo firewall-cmd --permanent --add-port=10251/tcp
$ sudo firewall-cmd --permanent --add-port=10252/tcp
$ sudo firewall-cmd --permanent --add-port=10255/tcp
$ sudo firewall-cmd --reload
```

Each time a port is added the system confirms with a ‘success’ message.

```
Master@Master:~$ sudo firewall-cmd --permanent --add-port=6443/tcp
Warning: ALREADY_ENABLED: 6443:tcp
success
Master@Master:~$ sudo firewall-cmd --permanent --add-port=2379-2380/tcp
Warning: ALREADY_ENABLED: 2379-2380:tcp
success
Master@Master:~$ sudo firewall-cmd --permanent --add-port=10250/tcp
Warning: ALREADY_ENABLED: 10250:tcp
success
Master@Master:~$ sudo firewall-cmd --permanent --add-port=10251/tcp
Warning: ALREADY_ENABLED: 10251:tcp
success
Master@Master:~$ sudo firewall-cmd --permanent --add-port=10252/tcp
Warning: ALREADY_ENABLED: 10252:tcp
success
Master@Master:~$ sudo firewall-cmd --permanent --add-port=10255/tcp
Warning: ALREADY_ENABLED: 10255:tcp
success
Master@Master:~$ sudo firewall-cmd --reload
success
Master@Master:~$ ...
```

Enter the following commands on each worker node:

```
$ sudo firewall-cmd --permanent --add-port=10251/tcp
$ sudo firewall-cmd --permanent --add-port=10255/tcp
$ firewall-cmd --reload
```

7.2 Update Iptables Settings

To ensure packets are properly processed by IP tables during filtering and port forwarding. Set the net.bridge.bridge-nf-call-iptables to '1' in your sysctl config file.

```
$ cat <<EOF > /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
```

```
$ sysctl --system
```

```
root@Master:~# cat <<EOF > /etc/sysctl.d/k8s.conf
> net.bridge.bridge-nf-call-ip6tables = 1
> net.bridge.bridge-nf-call-iptables = 1
> EOF
root@Master:~# sysctl --system
* Applying /etc/sysctl.d/10-console-messages.conf ...
kernel.printk = 4 4 1 7
* Applying /etc/sysctl.d/10-ipv6-privacy.conf ...
net.ipv6.conf.all.use_tempaddr = 2
net.ipv6.conf.default.use_tempaddr = 2
* Applying /etc/sysctl.d/10-kernel-hardening.conf ...
kernel.kptr_restrict = 1
* Applying /etc/sysctl.d/10-link-restrictions.conf ...
fs.protected_hardlinks = 1
fs.protected_symlinks = 1
* Applying /etc/sysctl.d/10-lxd-inotify.conf ...
fs.inotify.max_user_instances = 1024
* Applying /etc/sysctl.d/10-magic-sysrq.conf ...
kernel.sysrq = 176
* Applying /etc/sysctl.d/10-network-security.conf ...
net.ipv4.conf.default.rp_filter = 1
net.ipv4.conf.all.rp_filter = 1
net.ipv4.tcp_syncookies = 1
* Applying /etc/sysctl.d/10-ptrace.conf ...
kernel.yama.ptrace_scope = 1
* Applying /etc/sysctl.d/10-zero-page.conf ...
vm.mmap_min_addr = 65536
* Applying /usr/lib/sysctl.d/50-default.conf ...
net.ipv4.conf.all.promote_secondaries = 1
net.core.default_qdisc = fq_codel
* Applying /etc/sysctl.d/99-cloudimg-ipv6.conf ...
net.ipv6.conf.all.use_tempaddr = 0
net.ipv6.conf.default.use_tempaddr = 0
* Applying /etc/sysctl.d/99-sysctl.conf ...
* Applying /etc/sysctl.d/k8s.conf ...
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
* Applying /etc/sysctl.conf ...
```

7.3 Disable SELinux

SELinux needs to be set to permissive mode. Because containers need to access the host filesystem in most of the times. If you, set to permissive mode it will effectively disables its security functions.

Use the following command to disable SELinux:

```
$ sudo setenforce 0  
$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

7.4 Disable SWAP

To work kubelet properly we need to disable SWAP

To disable SWAP use the following commands.

```
$ sudo sed -i '/swap/d' /etc/fstab  
$ sudo swapoff -a
```

8 KUBEADM TO CREATE AND INITIALISE A CLUSTER

- Initialising the control-plane node run the below command on the **(master node)**

```
$ kubeadm init --kubernetes-version=${KUBE_VERSION}
```

```
root@Master:~# kubeadm init --kubernetes-version=${KUBE_VERSION}
[init] Using Kubernetes version: v1.21.0
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local master] and IPs
[10.96.0.1 10.0.0.4]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [localhost master] and IPs [10.0.0.4 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [localhost master] and IPs [10.0.0.4 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Starting the kubelet
[control-plane] Using manifest folder "/etc/kubernetes/manifests"
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-manager"
[control-plane] Creating static Pod manifest for "kube-scheduler"
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods from directory "/etc/kubernetes/manifests". This can take up to 4m0s
```

- If cluster initialisation has succeeded you will see a cluster join command. Copy and save that for future reference. This command would be used by the worker nodes to join the cluster

```
[bootstrap-token] configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized **successfully!**

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 10.0.0.4:6443 --token 40ptpl.10eqqlw6k4i3zrrq
--discovery-token-ca-cert-hash sha256:aab219d719ae51e847df39df3e0344462224e8162aa91ae07f30fa8379b6278
root@Master:~#
```

9 USING KUBEADM TO JOIN WORKER NODES TO THE CLUSTER

- Run the above received kubeadm join command (on both the **worker nodes**)

Note: This is above cluster command, you will get your command in your cluster so use that command not this command

```
$ kubeadm join 10.0.0.4:6443 --token 40tpl.10eqglw6k4i3zrrq \
--discovery-token-ca-cert-hash
sha256:aab219d719ae51e847df39df3e0344462224e8162aaf91ae07f30fa8379b6278
```

```
root@worker01:~# kubeadm join 10.0.0.4:6443 --token 40tpl.10eqglw6k4i3zrrq \
>      --discovery-token-ca-cert-hash sha256:aab219d719ae51e847df39df3e0344462224e8162aaf91ae07f30fa8379b6278
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

```
root@worker01:~#
```

Note: (Optional Step) If you forgot to add Save these tokens then you can always create a new token then use them to join Worker node

```
$ kubeadm token create --print-join-command
```

```
root@master:/home# kubeadm token create --print-join-command
W0807 11:35:35.839914    16801 configset.go:202] WARNING: kubeadm cannot validate component configs for API groups [kubel
et.config.k8s.io kubeProxy.config.k8s.io]
kubeadm join 10.0.4.4:6443 --token 8n46zi.u0vwo5h2yia42c83      --discovery-token-ca-cert-hash sha256:776e8fd4f4b456fa3b5
c3aaa6b78180df67c545309fe816fafdca7decba74a3f
root@master:/home#
```

Note: Everytime you are starting the Master you have to set these Environment Variables

- To start using the cluster set the environment variable on the (**master node**)

```
$ cp /etc/kubernetes/admin.conf $HOME/
$ chown $(id -u):$(id -g) $HOME/admin.conf
$ export KUBECONFIG=$HOME/admin.conf
```

```
root@kubeadm-master:/home/ubuntu# cp /etc/kubernetes/admin.conf $HOME/
root@kubeadm-master:/home/ubuntu# chown $(id -u):$(id -g) $HOME/admin.conf
root@kubeadm-master:/home/ubuntu# export KUBECONFIG=$HOME/admin.conf
root@kubeadm-master:/home/ubuntu#
```

Note: (Optional Step) To permanently set this environment variable, add it to your .bashrc file in your Root user.

```
$ echo 'export KUBECONFIG=$HOME/admin.conf' >> $HOME/.bashrc
```

```
root@master:/home/master# echo 'export KUBECONFIG=$HOME/admin.conf' >> $HOME/.bashrc

#
# colored GCC warnings and errors
#export GCC_COLORS='error=01;31:warning=01;35:note=01;36:caret=01;32:locus=01:quote=01'

# some more ls aliases
alias ll='ls -alF'
alias la='ls -A'
alias l='ls -CF'

# Add an "alert" alias for long running commands. Use like so:
# sleep 10; alert
alias alert='notify-send --urgency=low -i "$( [ $? = 0 ] && echo terminal || echo error)" "$(history|tail -n1|sed -e '\''s/^\s*[0-9]\+\s*//;s/[;&]\s*alert$/'\'')"'

# Alias definitions.
# You may want to put all your additions into a separate file like
# ~/.bash_aliases, instead of adding them here directly.
# See /usr/share/doc/bash-doc/examples in the bash-doc package.

if [ -f ~/.bash_aliases ]; then
    . ~/.bash_aliases
fi

# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
if ! shopt -oq posix; then
    if [ -f /usr/share/bash-completion/bash_completion ]; then
        . /usr/share/bash-completion/bash_completion
    elif [ -f /etc/bash_completion ]; then
        . /etc/bash_completion
    fi
fi
root@master:/home/master#
```

3. Verify the cluster status by executing kubectl command on the master node **(You should see 3 machines here)**

```
$ kubectl get nodes
```

```
root@Master:~# kubectl get nodes
NAME      STATUS      ROLES      AGE      VERSION
master    NotReady   control-plane,master   8m59s   v1.21.0
worker01  NotReady   <none>     3m41s   v1.21.0
root@Master:~#
```

4. Install CNI so that pods can communicate across nodes and also Cluster DNS to start functioning. Apply weave CNI (Container Network Interface) on the master node

```
$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n")"
```

```
root@kubeadm-master:/home/ubuntu# kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"

serviceaccount/weave-net created
clusterrole.rbac.authorization.k8s.io/weave-net created
clusterrolebinding.rbac.authorization.k8s.io/weave-net created
role.rbac.authorization.k8s.io/weave-net created
rolebinding.rbac.authorization.k8s.io/weave-net created
daemonset.apps/weave-net created
[root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu#
```

5. Wait for few minutes and verify the cluster status by executing kubectl command on the master node and see that nodes come to ready state (*You should see 3 machines here*)

```
$ kubectl get nodes
```

```
root@Master:~# kubectl get nodes
NAME     STATUS   ROLES          AGE      VERSION
master   Ready    control-plane,master   13m    v1.21.0
worker01 Ready    <none>        8m2s    v1.21.0
root@Master:~#
```

6. Verify the status of the system pods like coredns, weave-net, kube-proxy and all other master node system processes

```
$ kubectl get pods -n kube-system
```

```
root@Master:~# kubectl get pods -n kube-system
NAME           READY   STATUS    RESTARTS   AGE
coredns-558bd4d5db-8jkhm   1/1     Running   0          14m
coredns-558bd4d5db-v7w6n   1/1     Running   0          14m
etcd-master     1/1     Running   0          14m
kube-apiserver-master   1/1     Running   1          14m
kube-controller-manager-master  1/1     Running   0          14m
kube-proxy-6qhm4       1/1     Running   0          8m56s
kube-proxy-xgdkm       1/1     Running   0          14m
kube-scheduler-master   1/1     Running   0          14m
weave-net-9n2tz        2/2     Running   1          2m56s
weave-net-t6gb6        2/2     Running   1          2m56s
root@Master:~#
```

7. Git clone the code files to be used for further labs

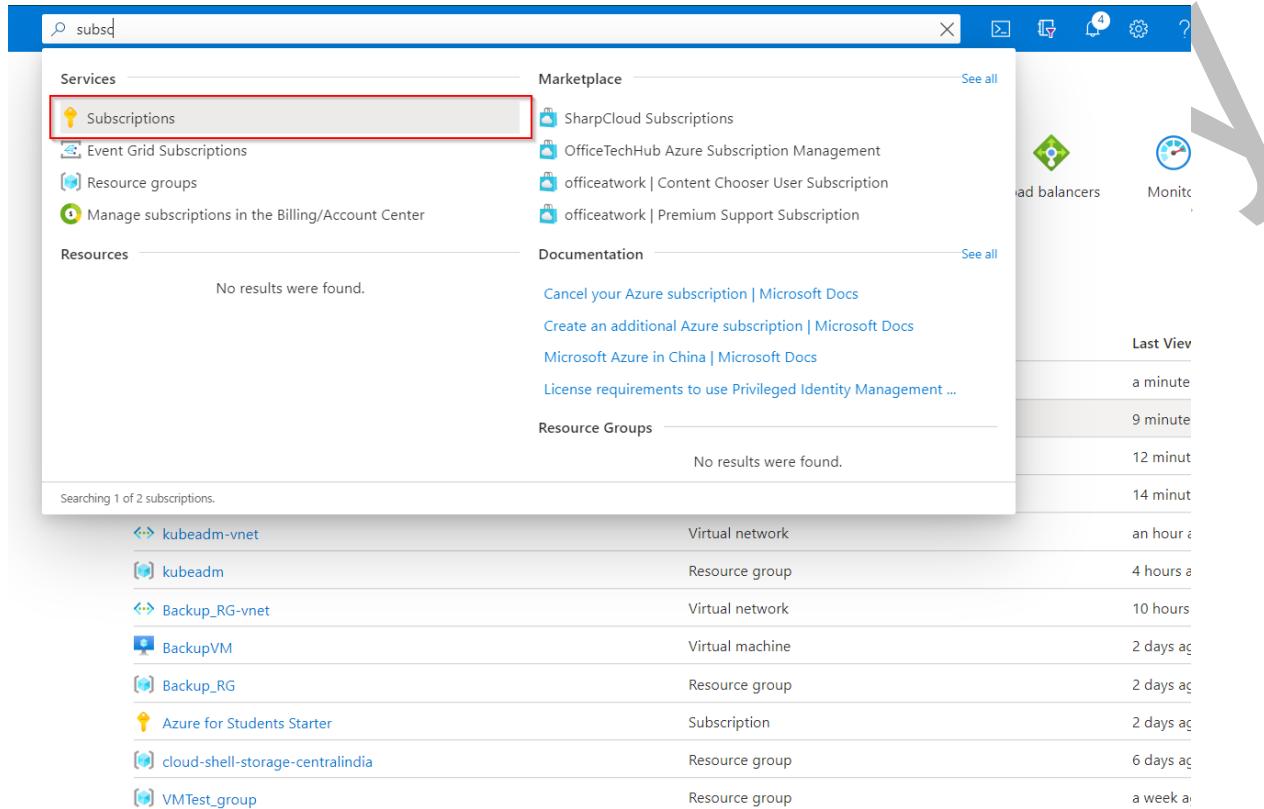
```
$ git clone https://github.com/k21academyuk/Kubernetes  
$ cd Kubernetes
```

K21 Academy

10 (OPTIONAL) UNSUBSCRIBE PAY AS YOU GO IN AZURE

If you already upgraded to pay as you go then you can unsubscribe/deactivate the pay as you go and use delete all the resources.

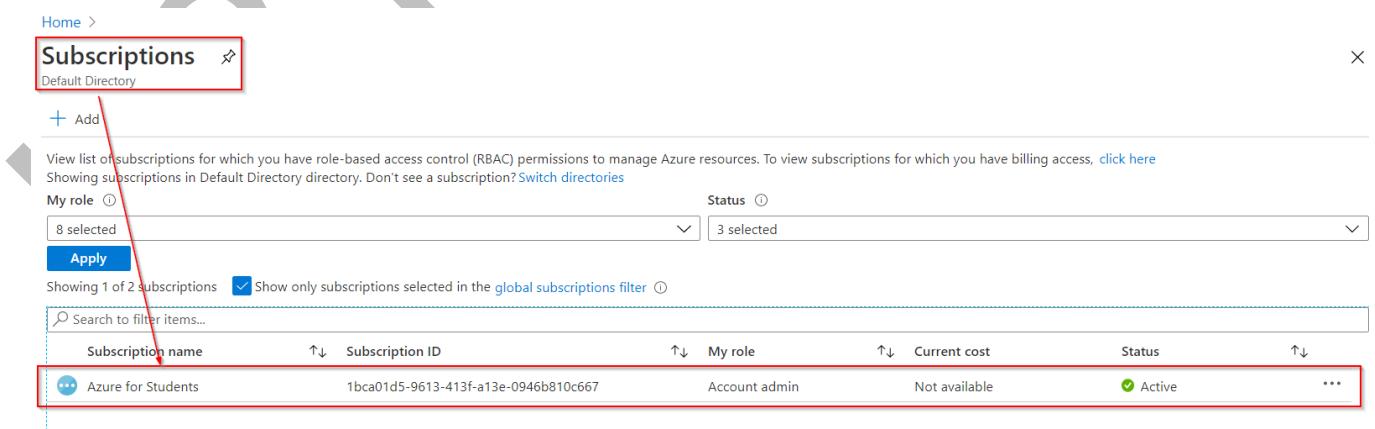
1. In Azure Search bar search for Azure Subscription



The screenshot shows the Azure portal search results for the query "subsc". The "Subscriptions" item under the "Services" section is highlighted with a red box. The results also include "Event Grid Subscriptions", "Resource groups", and "Manage subscriptions in the Billing/Account Center". Under "Resources", it says "No results were found.". In the "Marketplace" section, there are links to "SharpCloud Subscriptions", "OfficeTechHub Azure Subscription Management", "officeatwork | Content Chooser User Subscription", and "officeatwork | Premium Support Subscription". The "Documentation" section includes links to "Cancel your Azure subscription", "Create an additional Azure subscription", "Microsoft Azure in China", and "License requirements to use Privileged Identity Management ...". The "Resource Groups" section shows a list of resources: "kubeadm-vnet" (Virtual network), "kubeadm" (Resource group), "Backup_RG-vnet" (Virtual network), "BackupVM" (Virtual machine), "Backup_RG" (Resource group), "Azure for Students Starter" (Subscription), "cloud-shell-storage-centralindia" (Resource group), and "VMTest_group" (Resource group). A grey arrow points from the "Subscriptions" link in the search bar to the highlighted "Subscriptions" item in the results.

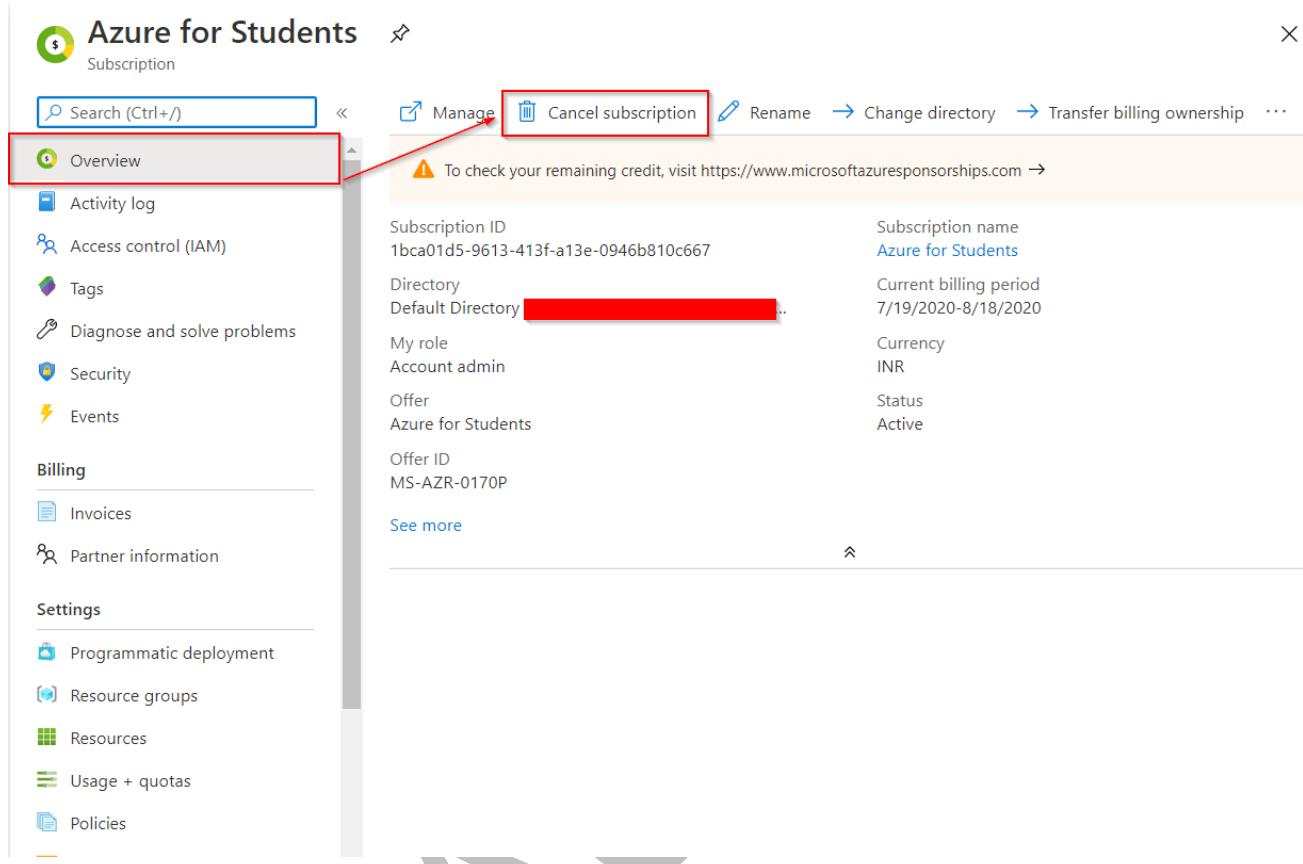
Note: In screenshots the account is Student account not as paid account but screens are same

2. Select the subscription that you want to cancel.



The screenshot shows the "Subscriptions" page in the Azure portal. The "Subscriptions" button in the top navigation bar is highlighted with a red box. The page displays a list of subscriptions: "Azure for Students" (Subscription ID: 1bca01d5-9613-413f-a13e-0946b810c667, Status: Active). A red box highlights this row. The "Apply" button in the "My role" section is also highlighted with a red box. A red arrow points from the "Azure for Students" row to the "Apply" button. The "Status" column shows "Active" for the selected subscription.

3. Select Overview, and then select Cancel subscription.



The screenshot shows the Azure portal interface for managing a subscription. The top navigation bar includes a search bar, 'Manage', 'Cancel subscription' (which is highlighted with a red box), 'Rename', 'Change directory', 'Transfer billing ownership', and more. Below the navigation is a warning message: 'To check your remaining credit, visit https://www.microsoftazuresponsorships.com →'. The main content area displays subscription details:

Subscription ID	1bca01d5-9613-413f-a13e-0946b810c667	Subscription name	Azure for Students
Directory	Default Directory [REDACTED]	Current billing period	7/19/2020-8/18/2020
My role	Account admin	Currency	INR
Offer	Azure for Students	Status	Active
Offer ID	MS-AZR-0170P		

Below the details, there's a 'See more' link. The left sidebar contains links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Security, Events, Billing (Invoices, Partner information), Settings (Programmatic deployment, Resource groups, Resources, Usage + quotas, Policies), and Help.

Are you sure you want to cancel subscription Azure for Students?



Warning! If you cancel this subscription, all of its data will be deleted after the retention period ([learn more](#)). Make sure you back up the information you want to save before continuing.

Your resources may be in use for critical workloads. As a precaution, we strongly recommend you review these resources and delete them prior to cancelling your subscription. If you've already reviewed and/or deleted your resources and are still seeing this message, you can click the 'Ignore and cancel' button below and proceed with cancelling your subscription.

[Review resources](#)

[Ignore and cancel](#)

4. Now please create a new account and you can continue without pay-as-you go. Follow Create Azure account guide to create an account

5. FAQ

What happens after I cancel my subscription?

After you cancel, billing is stopped immediately. However, it can take up to 10 minutes for the cancellation to show in the portal. If you cancel in the middle of a billing period, we send the final invoice on your typical invoice date after the period ends.

After you cancel, your services are disabled. That means your virtual machines are de-allocated, temporary IP addresses are freed, and storage is read-only.

After your subscription is canceled, Microsoft waits 30 - 90 days before permanently deleting your data in case you need to access it or you change your mind. We don't charge you for retaining the data. To learn more, see [Microsoft Trust Center - How we manage your data](#).

K21Acade^

11 TROUBLESHOOTING

11.1 Getting Warning while Configure cgroup driver used by kubelet on control-plane

Issue: Getting warning while like below Image

```

Unpacking kubectl (1.18.6-00) ...
Selecting previously unselected package kubeadm.
Preparing to unpack .../6-kubeadm_1.18.6-00_amd64.deb ...
Unpacking kubeadm (1.18.6-00) ...
Setting up conntrack (1:1.4.4+snapshot20161117-6ubuntu2) ...
Setting up kubernetes-cni (0.8.6-00) ...
Setting up cri-tools (1.13.0-01) ...
Setting up socat (1.7.3.2-2ubuntu2) ...
Setting up kubelet (1.18.6-00) ...
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /lib/systemd/system/kubelet.service.
Setting up kubectl (1.18.6-00) ...
Setting up kubeadm (1.18.6-00) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
root@kubeadm-master:/home/azureadmin# apt-mark hold kubelet kubeadm kub
kubeadm set on hold.
kubeadm set on hold.
kubeadm set on hold.
root@kubeadm-master:/home/azureadmin# docker info | grep -i cgroup
Cgroup Driver: cgroupfs
WARNING: No swap limit support
root@kubeadm-master:/home/azureadmin#

```

Fix: It's just a warning plz ignore this warning

11.2 Getting Error while Connecting node to Master node

Issue: getting error while running connecting node to master node command

```

root@master:~# kubeadm join 10.0.4.4:6443 --token x0ovic.bu8xnx9154fe77/on \
> --discovery-token-ca-cert-hash sha256:776e8fd4f4b456fa3b5c3aaa6b78180df67c545309fe816fafdca7decba74a3f
W0728 10:50:03.928021 44127 join.go:346] [preflight] WARNING: JoinControlPlane.controlPlane settings will be ignored when control-plane flag is not set.
[preflight] Running pre-flight checks
[WARNING Service-Docker]: docker service is not enabled, please run 'systemctl enable docker.service'
[WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the Docker cgroup driver. The recommended driver is "systemd". Please follow the guide at https://kuber
.net/docs/setup/cri/
error execution phase preflight: [preflight] Some fatal errors occurred:
[ERROR DirAvailable--etc-kubernetes-manifests]: /etc/kubernetes/manifests is not empty
[ERROR FileAvailable--etc-kubernetes-kubelet.conf]: /etc/kubernetes/kubelet.conf already exists
[ERROR Port-10250]: Port 10250 is in use
[ERROR FileAvailable--etc-kubernetes-pki-ca.crt]: /etc/kubernetes/pki/ca.crt already exists
[preflight] If you know what you are doing, you can make a check non-fatal with `--ignore-preflight-errors=...`.
To see the stack trace of this error execute with --v=5 or higher

```

Reason: If you are getting error like above then you are running connecting command in master node instead of worker node

Fix: Please run above Command in Worker Nodes.

11.3 Running kubectl get nodes and getting server localhost:8080

Issue: If you are running any command like kubectl get nodes and getting error like below

```
root@kubernetes-master:/home/AzureUser# kubectl get nodes
The connection to the server localhost:8080 was refused - did you specify the
right host or port?
root@kubernetes-master:/home/AzureUser# kubectl apply -f "
https://cloud.weave.works/k8s/net?k8s-version=\$\(kubectl version | base64 | tr -d '\n')
The connection to the server localhost:8080 was refused - did you specify the
right host or port?
The connection to the server localhost:8080 was refused - did you specify the
right host or port?
```

4:33 AM

Reason: Environment variable not set on the master node

Fix: set the environment variable on the master nod.

To start using the cluster set the environment variable on the master node

```
$ cp /etc/kubernetes/admin.conf $HOME/
$ chown $(id -u):$(id -g) $HOME/admin.conf
$ export KUBECONFIG=$HOME/admin.conf
```

```
root@kubeadm-master:/home/ubuntu# cp /etc/kubernetes/admin.conf $HOME/
root@kubeadm-master:/home/ubuntu# chown $(id -u):$(id -g) $HOME/admin.conf
root@kubeadm-master:/home/ubuntu# export KUBECONFIG=$HOME/admin.conf
root@kubeadm-master:/home/ubuntu#
```

11.4 Error when Running Kubeadm init Command

Issue: If you are running kubeadm init command and getting error like below

```
root@master:/home/ubuntu# kubeadm init
[init] Using Kubernetes version: v1.20.4
[preflight] Running pre-flight checks
[WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the Docker cgroup driver. The recommended driver is "systemd". Please follow the guide at
https://kubernetes.io/docs/setup/cri/
error execution phase preflight: [preflight] Some fatal errors occurred:
[ERROR Port-6443]: Port 6443 is in use
[ERROR Port-10259]: Port 10259 is in use
[ERROR Port-10257]: Port 10257 is in use
[ERROR FileAvailable--etc-kubernetes-manifests-kube-apiserver.yaml]: /etc/kubernetes/manifests/kube-apiserver.yaml already exists
[ERROR FileAvailable--etc-kubernetes-manifests-kube-controller-manager.yaml]: /etc/kubernetes/manifests/kube-controller-manager.yaml already exists
[ERROR FileAvailable--etc-kubernetes-manifests-kube-scheduler.yaml]: /etc/kubernetes/manifests/kube-scheduler.yaml already exists
[ERROR FileAvailable--etc-kubernetes-manifests-etcd.yaml]: /etc/kubernetes/manifests/etc.yaml already exists
[ERROR Port-10250]: Port 10250 is in use
[ERROR Port-2379]: Port 2379 is in use
[ERROR Port-2380]: Port 2380 is in use
[ERROR DirAvailable--var-lib-etcd]: /var/lib/etcd is not empty
[preflight] If you know what you are doing, you can make a check non-fatal with `--ignore-preflight-errors=...`  
To see the stack trace of this error execute with --v=5 or higher
```

Reason: Because you are running kubeadm init command again

Fix: please check if kubernetes cluster running or not already by

```
kubectl get nodes
```

if its not running please run below command and run **kubeadm init** command again.

```
kubeadm reset
```

11.5 Using Expired Token When Joining Node

Issue: Can't join Worker Node to master node using previous created tokens.

Reason: Tokens created init command by default expires after the 24 hours so you have to create a new token to join a new worker node to master node.

Fix: Create new tokens using below command and use new tokens.

```
$ kubeadm token create --print-join-command
```

11.6 Error when Running Kubeadm init Command

When you are doing Optional guide [Bonus] [Optional] Connect Windows Worker Node to Ubuntu Master Node we have installed **Flannel CNI** so due to this this is creating issue.

Issue: two CNI plugin installed weave & flannel and flannel is not working.

```
sc = failed to set up sandbox container "060839caaab9b17e26aa72487610e395cedeca6b63a07afdc90e
o set up pod "nginx_default" network: open /run/flannel/subnet.env: no such file or directory
root@master:~# kubectl get pods -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE
coredns-74ff55c5b-hlgs5           0/1    Completed  0          4d23h
coredns-74ff55c5b-ntncp          0/1    Completed  0          4d23h
etcd-master                        1/1    Running   5          4d23h
kube-apiserver-master            1/1    Running   6          4d23h
kube-controller-manager-master   1/1    Running   5          4d23h
kube-flannel-ds-2vdwf           0/1    CrashLoopBackOff 180        4d22h
kube-flannel-ds-56wdj            0/1    CrashLoopBackOff 181        4d22h
kube-flannel-ds-t6cvq            0/1    CrashLoopBackOff 180        4d22h
kube-proxy-78qj2                 1/1    Running   5          4d23h
kube-proxy-tq8zk                 1/1    Running   5          4d23h
kube-proxy-xr2mp                 1/1    Running   5          4d23h
kube-scheduler-master            1/1    Running   5          4d23h
weave-net-2wr5q                  2/2    Running   11         4d22h
weave-net-xch7t                  2/2    Running   11         4d22h
weave-net-zx2b9                  2/2    Running   11         4d22h
root@master:~# kubectl describe pod kube-flannel-ds-2vdwf -n kube-system
Name:           kube-flannel-ds-2vdwf
Namespace:      kube-system
...
root@master:/home/ubuntu# kubectl get pods --output=wide
NAME                               READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NOD
E   READINESS GATES
nginx                            0/1    ContainerCreating  0          2m47s  <none>       worker-01  <none>
<none>
test-replicaset-7m8fl            0/1    ContainerCreating  0          34m    <none>       worker-02  <none>
<none>
test-replicaset-rtqpg            0/1    ContainerCreating  0          34m    <none>       worker-02  <none>
<none>
```

Reason: flannel CNI plug-in not working so pod not scheduling on the worker nodes because pods are using flannel CNI

Fix: Remove Flannel from the cluster from the master node

```
$ kubectl delete -f kube-flannel.yml
$ kubectl delete -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

Install Weave again

```
$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

Perfrom below stepson all the nodes master and worker

```
$ rm /etc/cni/net.d/10-flannel.conflist
$ systemctl restart kubelet
```

Now the problem is solved

```

IP: 10.40.0.3
Containers:
  nginx:
    Container ID: docker://4c2d4c461ede0e528701948d81e7099617ee5a73ca9dc291f0a03482cc727ac5
    Image: nginx
    Image ID: docker-pullable://nginx@sha256:be093310363114f4f5074df54ae2748123118ba04f4bc076ac4d3268dc4d1ff1
    Port: <none>
    Host Port: <none>
    State: Running
    Started: Sat, 27 Mar 2021 16:38:33 +0000
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-p446g (ro)
Conditions:
  Type          Status
  Initialized   True
  Ready         True
  ContainersReady True
  PodScheduled  True
Volumes:
  default-token-p446g:
    Type: Secret (a volume populated by a Secret)
    SecretName: default-token-p446g
    Optional: false
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
           node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason          Age   From            Message
  ----   -----          --   ----            -----
  Warning FailedCreatePodSandBox 5m27s (x6403 over 4h) kubelet  (combined from similar events): Failed to create pod sandbox: rpc error: code = Unknown desc = failed to set up sandbox container "a390bbba259556ef0b01ae0ec1333cdc7484faf80b3e5d357fa760c9a1fdd72f" network for pod "nginx": networkPlugin cni failed to set up pod "nginx_default" network: open /run/flannel/subnet.env: no such file or directory
  Normal   Pulling          39s   kubelet        Pulling image "nginx"
  Normal   Pulled           34s   kubelet        Successfully pulled image "nginx" in 4.371419363s
  Normal   Created          29s   kubelet        Created container nginx
  Normal   Started          29s   kubelet        Started container nginx
root@master:/home/ubuntu# 

```

11.7 Unable to Ping Master node to Worker after Installing Docker

Issue: When we create 3 machine for Kubernetes by default they can ping to each other but after installing the docker on the master machine unable to ping machines.

```

root@master:~# ping 172.17.0.4
PING 172.17.0.4 (172.17.0.4) 56(84) bytes of data.
64 bytes from 172.17.0.4: icmp_seq=2 ttl=64 time=44.0 ms
64 bytes from 172.17.0.4: icmp_seq=3 ttl=64 time=33.0 ms
64 bytes from 172.17.0.4: icmp_seq=4 ttl=64 time=27.0 ms
^C64 bytes from 172.17.0.4: icmp_seq=5 ttl=64 time=27.5 ms
64 bytes from 172.17.0.4: icmp_seq=6 ttl=64 time=27.3 ms
64 bytes from 172.17.0.4: icmp_seq=7 ttl=64 time=29.9 ms
^C
--- 172.17.0.4 ping statistics ---
7 packets transmitted, 6 received, 14% packet loss, time 6024ms
rtt min/avg/max/mdev = 27.051/31.509/44.047/5.978 ms
root@master:~# ping 172.17.0.5
PING 172.17.0.5 (172.17.0.5) 56(84) bytes of data.
64 bytes from 172.17.0.5: icmp_seq=1 ttl=64 time=32.8 ms
64 bytes from 172.17.0.5: icmp_seq=2 ttl=64 time=32.4 ms
64 bytes from 172.17.0.5: icmp_seq=3 ttl=64 time=32.4 ms
^C
--- 172.17.0.5 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 32.492/32.625/32.886/0.278 ms
root@master:~# ping google.com
PING google.com (172.217.12.238) 56(84) bytes of data.
64 bytes from iad30s15-in-f14.1e100.net (172.217.12.238): icmp_seq=1 ttl=111 time=5.34 ms
64 bytes from iad30s15-in-f14.1e100.net (172.217.12.238): icmp_seq=2 ttl=111 time=5.22 ms
^C
--- google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 5.228/5.286/5.345/0.093 ms

```

```
PING 172.16.0.4 (172.16.0.4) 56(84) bytes of data.
64 bytes from 172.16.0.4: icmp_seq=1 ttl=64 time=27.8 ms
64 bytes from 172.16.0.4: icmp_seq=2 ttl=64 time=27.1 ms
^C
--- 172.16.0.4 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 27.134/27.470/27.806/0.336 ms
root@Worker-1:~# ping 172.17.0.5
PING 172.17.0.5 (172.17.0.5) 56(84) bytes of data.
64 bytes from 172.17.0.5: icmp_seq=1 ttl=64 time=2.34 ms
64 bytes from 172.17.0.5: icmp_seq=2 ttl=64 time=1.48 ms
64 bytes from 172.17.0.5: icmp_seq=3 ttl=64 time=1.31 ms
^C
--- 172.17.0.5 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 1.313/1.715/2.345/0.452 ms
root@Worker-1:~# ping google.com
PING google.com (142.250.113.139) 56(84) bytes of data.
^C64 bytes from 142.250.113.139: icmp_seq=1 ttl=108 time=7.92 ms

--- google.com ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 7.926/7.926/7.926/0.000 ms
root@Worker-1:~# ping google.com
PING google.com (142.250.113.101) 56(84) bytes of data.
64 bytes from 142.250.113.101: icmp_seq=1 ttl=105 time=9.32 ms
64 bytes from 142.250.113.101: icmp_seq=2 ttl=105 time=9.39 ms
64 bytes from 142.250.113.101: icmp_seq=3 ttl=105 time=9.35 ms
64 bytes from 142.250.113.101: icmp_seq=4 ttl=105 time=9.36 ms
64 bytes from 142.250.113.101: icmp_seq=5 ttl=105 time=9.35 ms
64 bytes from 142.250.113.101: icmp_seq=6 ttl=105 time=9.32 ms
64 bytes from 142.250.113.101: icmp_seq=7 ttl=105 time=9.68 ms
^C
--- google.com ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 15027ms
rtt min/avg/max/mdev = 9.320/9.398/9.685/0.173 ms
root@Worker-1:~#
```

```
ubuntu@Worker-2:~$ sudo -i
root@Worker-2:~# ping 172.16.0.4
PING 172.16.0.4 (172.16.0.4) 56(84) bytes of data.
64 bytes from 172.16.0.4: icmp_seq=1 ttl=64 time=33.0 ms
64 bytes from 172.16.0.4: icmp_seq=2 ttl=64 time=32.6 ms
^C
--- 172.16.0.4 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 32.641/32.821/33.001/0.180 ms
root@Worker-2:~# ping 172.17.0.4
PING 172.17.0.4 (172.17.0.4) 56(84) bytes of data.
64 bytes from 172.17.0.4: icmp_seq=1 ttl=64 time=1.37 ms
64 bytes from 172.17.0.4: icmp_seq=2 ttl=64 time=1.37 ms
64 bytes from 172.17.0.4: icmp_seq=3 ttl=64 time=1.38 ms
^C
--- 172.17.0.4 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 1.374/1.378/1.388/0.043 ms
root@Worker-2:~# ping google.com
PING google.com (172.217.14.174) 56(84) bytes of data.
64 bytes from dfw28s22-in-f14.1e100.net (172.217.14.174): icmp_seq=1 ttl=116 time=8.32 ms
64 bytes from dfw28s22-in-f14.1e100.net (172.217.14.174): icmp_seq=2 ttl=116 time=8.47 ms
^C
--- google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 8.328/8.399/8.470/0.071 ms
root@Worker-2:~#
```

After installing the docker on master machine



```
Selecting previously unselected package runc.
Preparing to unpack .../2-runc_1.0.0-rc93~Ubuntu18.04.2_amd64.deb ...
Unpacking runc (1.0.0-rc93~Ubuntu18.04.2) ...
Selecting previously unselected package containerd.
Preparing to unpack .../3-containerd_1.4.4~Ubuntu18.04.2_amd64.deb ...
Unpacking containerd (1.4.4~Ubuntu18.04.2) ...
Selecting previously unselected package docker.io.
Preparing to unpack .../4-docker.io_20.10.2~Ubuntu18.04.2_amd64.deb ...
Unpacking docker.io (20.10.2~Ubuntu18.04.2) ...
Selecting previously unselected package ubuntu-fan.
Preparing to unpack .../5-ubuntu-fan_0.12.10_all.deb ...
Unpacking ubuntu-fan (0.12.10) ...
Setting up runc (1.0.0-rc93~Ubuntu18.04.2) ...
Setting up containerd (1.4.4~Ubuntu18.04.2) ...
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /lib/systemd/system/containerd.service.
Setting up bridge-utils (1.5-15ubuntu1) ...
Setting up ubuntu-fan (0.12.10) ...
Created symlink /etc/systemd/system/multi-user.target.wants/ubuntu-fan.service → /lib/systemd/system/ubuntu-fan.service.
Setting up pigg (2.4-1) ...
Setting up docker.io (20.10.2~Ubuntu18.04.2) ...
Adding group 'docker' (GID 116) ...
Done.
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /lib/systemd/system/docker.service.
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /lib/systemd/system/docker.socket.
Processing triggers for systemd (237-3ubuntu10.47) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Processing triggers for ureadahead (0.100.0-21) ...
root@master:~# ping 172.17.0.5
PING 172.17.0.5 (172.17.0.5) 56(84) bytes of data.
From 172.17.0.1 icmp_seq=1 Destination Host Unreachable
From 172.17.0.1 icmp_seq=2 Destination Host Unreachable
From 172.17.0.1 icmp_seq=3 Destination Host Unreachable
^C
--- 172.17.0.5 ping statistics ---
6 packets transmitted, 0 received, +3 errors, 100% packet loss, time 5099ms
pipe 4
root@master:~# apt-get update && apt install -y docker.io
```

```
root@master:~# ifconfig
dockero: flags=40000UP,BROADCAST,MULTICAST> mtu 1500
  inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    ether 02:42:65:31:f0:20 txqueuelen 0 (Ethernet)
      RX packets 0 bytes 0 (0.0 B)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 0 bytes 0 (0.0 B)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
  inet 172.16.0.4 netmask 255.255.255.0 broadcast 172.16.0.255
    ether 00:0d:3a:7b:c8:54 txqueuelen 1000 (Ethernet)
      RX packets 154141 bytes 193565826 (193.5 MB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 40858 bytes 7940474 (7.9 MB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
  inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
      loop txqueuelen 1000 (Local Loopback)
      RX packets 4539 bytes 566139 (566.1 KB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 4539 bytes 566139 (566.1 KB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
root@Worker-1:~# ifconfig
docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.18.0.1 netmask 255.255.0.0 broadcast 172.18.255.255
        ether 02:42:fc:c6:14:a1 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.4 netmask 255.255.255.0 broadcast 172.17.0.255
        inet6 fe80::20d:3aff:fe5d:e47c prefixlen 64 scopeid 0x20<link>
        ether 00:0d:3a:5d:e4:7c txqueuelen 1000 (Ethernet)
    RX packets 96584 bytes 115191527 (115.1 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 32918 bytes 7056836 (7.0 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
        loop txqueuelen 1000 (Local Loopback)
    RX packets 6385 bytes 687091 (687.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 6385 bytes 687091 (687.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```



Reason: When we install docker on the machine by default the docker0 takes default IP that is 172.17.0.1/16 so this IP range is conflicting with the other two worker nodes private IP that is 172.17.0.4 and 172.17.0.5 that's why when we are ping to master to worker node it pings using the docker0 instead of host because both in same IP range.

```

ubuntu@master:~$ ping 172.17.0.4
PING 172.17.0.4 (172.17.0.4) 56(84) bytes of data.
From 172.17.0.1 icmp_seq=1 Destination Host Unreachable
From 172.17.0.1 icmp_seq=2 Destination Host Unreachable
From 172.17.0.1 icmp_seq=3 Destination Host Unreachable
From 172.17.0.1 icmp_seq=4 Destination Host Unreachable
From 172.17.0.1 icmp_seq=5 Destination Host Unreachable
From 172.17.0.1 icmp_seq=6 Destination Host Unreachable
^C
--- 172.17.0.4 ping statistics ---
7 packets transmitted, 0 received, +6 errors, 100% packet loss, time 6145ms
pipe 4
ubuntu@master:~$ ifconfig
docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
        ether 02:42:b9:ba:9a:37 txqueuelen 0 (Ethernet)
        RX packets 0 bytes 0 (0.0 B)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 0 bytes 0 (0.0 B)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.16.0.4 netmask 255.255.255.0 broadcast 172.16.0.255
        inetb fe80::20d3:aff:fe7b:c854 prefixlen 64 scopeid 0x20<link>
        ether 00:0d:3a:7b:c8:54 txqueuelen 1000 (Ethernet)
        RX packets 88661 bytes 111897087 (111.8 MB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 22430 bytes 4663415 (4.6 MB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
        loop txqueuelen 1000 (Local Loopback)
        RX packets 2888 bytes 355031 (355.0 KB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 2888 bytes 355031 (355.0 KB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ubuntu@master:~$
```

Fix: First, explicitly set the address of docker0 in your **/etc/docker/daemon.json** Use any network that doesn't conflict your worker node networks.

Reference: <https://superuser.com/questions/1336567/installing-docker-ce-in-ubuntu-18-04-breaks-internet-connectivity-of-host/1380836#1380836?newreg=74f46627c99344e2b382b6ab3efd49ed>

If **/etc/docker/daemon.json** not present then create new file

```
$ vi /etc/docker/daemon.json
```

Add any Ip range in **bip field**

```
{
  "bip": "172.31.0.1/16"
}
```

```
{
  "bip": "172.31.0.1/16"
}

~  

~  

~
```

Now restart the docker

```
$ service docker restart
```

```
root@master:~# vim /etc/docker/daemon.json
root@master:~# service docker restart
root@master:~# ifconfig
```

Now the Ip of docker0 is changed now try to ping again from master to worker machine. Problem is resolved.

```
root@master:~# ifconfig
docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
  inet 172.31.0.1 netmask 255.255.0.0 broadcast 172.31.255.255
    ether 02:42:65:31:f0:20 txqueuelen 0 (Ethernet)
      RX packets 0 bytes 0 (0.0 B)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 0 bytes 0 (0.0 B)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
  inet 172.16.0.4 netmask 255.255.255.0 broadcast 172.16.0.255
    inet6 fe80::20d:3aff:fe7b:c854 prefixlen 64 scopeid 0x20<link>
      ether 00:0d:3a:7b:c8:54 txqueuelen 1000 (Ethernet)
      RX packets 158373 bytes 195609150 (195.6 MB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 45216 bytes 8933809 (8.9 MB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
  inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
      loop txqueuelen 1000 (Local Loopback)
      RX packets 5129 bytes 641702 (641.7 KB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 5129 bytes 641702 (641.7 KB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

root@master:~# ping 172.17.0.4
PING 172.17.0.4 (172.17.0.4) 56(84) bytes of data.
64 bytes from 172.17.0.4: icmp_seq=1 ttl=64 time=27.7 ms
64 bytes from 172.17.0.4: icmp_seq=2 ttl=64 time=27.3 ms
64 bytes from 172.17.0.4: icmp_seq=3 ttl=64 time=27.2 ms
^C
--- 172.17.0.4 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 27.291/27.452/27.704/0.180 ms
root@master:~#
```

11.8 Kubelet is not running

Issue: Kubelet is not running.

```
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-manager"
[control-plane] Creating static Pod manifest for "kube-scheduler"
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods from directory "/etc/kubernetes/manifests". This can take up to 4m0s
[kubelet-check] Initial timeout of 40s passed.
[kubelet-check] It seems like the kubelet isn't running or healthy.
[kubelet-check] The HTTP call equal to 'curl -sSL http://localhost:10248/healthz' failed with error: Get "http://localhost:10248/healthz": dial tcp [::1]:10248: connect : connection refused.
[kubelet-check] It seems like the kubelet isn't running or healthy.
[kubelet-check] The HTTP call equal to 'curl -sSL http://localhost:10248/healthz' failed with error: Get "http://localhost:10248/healthz": dial tcp [::1]:10248: connect : connection refused.
[kubelet-check] It seems like the kubelet isn't running or healthy.
[kubelet-check] The HTTP call equal to 'curl -sSL http://localhost:10248/healthz' failed with error: Get "http://localhost:10248/healthz": dial tcp [::1]:10248: connect : connection refused.
[kubelet-check] It seems like the kubelet isn't running or healthy.
[kubelet-check] The HTTP call equal to 'curl -sSL http://localhost:10248/healthz' failed with error: Get "http://localhost:10248/healthz": dial tcp [::1]:10248: connect : connection refused.
[kubelet-check] It seems like the kubelet isn't running or healthy.
[kubelet-check] The HTTP call equal to 'curl -sSL http://localhost:10248/healthz' failed with error: Get "http://localhost:10248/healthz": dial tcp [::1]:10248: connect : connection refused.
[kubelet-check] It seems like the kubelet isn't running or healthy.
[kubelet-check] The HTTP call equal to 'curl -sSL http://localhost:10248/healthz' failed with error: Get "http://localhost:10248/healthz": dial tcp [::1]:10248: connect : connection refused.

Unfortunately, an error has occurred:
timed out waiting for the condition

This error is likely caused by:
- The kubelet is not running
- The kubelet is unhealthy due to a misconfiguration of the node in some way (required cgroups disabled)

If you are on a systemd-powered system, you can try to troubleshoot the error with the following commands:
- 'systemctl status kubelet'
- 'journalctl -xeu kubelet'

Additionally, a control plane component may have crashed or exited when started by the container runtime.
To troubleshoot, list all containers using your preferred container runtimes CLI.

Here is one example how you may list all Kubernetes containers running in docker:
- 'docker ps -a | grep kube | grep -v pause'
Once you have found the failing container, you can inspect its logs with:
- 'docker logs CONTAINERID'

error execution phase wait-control-plane: couldn't initialize a Kubernetes cluster
To see the stack trace of this error execute with --v=5 or higher
[root@master ~]#
```

Reason: Kubernetes has removed docker as a runtime in Kubernetes 1.22, due to this dockershim is no longer working as container run time. That's why Kubelet is not running.

Fix: Run the following commands on Master and all the Worker nodes.

```
$ apt-get remove -y docker.io kubelet kubeadm kubectl kubernetes-cni
$ apt-get autoremove -y
$ kubeadm reset -f
```

Please refer to edition-15+ guides for Kubernetes installation.

11.9 Running kubectl get nodes and getting server 10.1.0.4:6443 was refused

Issue: If you are running any command like kubectl get nodes and getting error like below.

```
root@master:/home/ubuntu# kubectl get nodes
The connection to the server 10.1.0.4:6443 was refused - did you specify the right host or port?
```

Reason: Environment variable not set on the master node.

Fix: set the environment variable on the master nod.

To start using the cluster set the environment variable on the master node

```
$ cp /etc/kubernetes/admin.conf $HOME/
$ chown $(id -u):$(id -g) $HOME/admin.conf
$ export KUBECONFIG=$HOME/admin.conf
```

```
root@kubeadm-master:/home/ubuntu# cp /etc/kubernetes/admin.conf $HOME/
root@kubeadm-master:/home/ubuntu# chown $(id -u):$(id -g) $HOME/admin.conf
root@kubeadm-master:/home/ubuntu# export KUBECONFIG=$HOME/admin.conf
root@kubeadm-master:/home/ubuntu#
```

11.10 Getting error while running

Update command after adding the Kubernetes apt repository

Issue: Getting error while running Update command after adding the Kubernetes apt repository.

```
root@worker-02:~# echo "deb[signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb[signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main
root@worker-02:~# KUBE_VERSION=1.21.0
root@worker-02:~# apt-get update
E: Type 'deb[signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]' is not known on line 1 in source list /etc/apt/sources.list.d/kubernetes.list
E: The list of sources could not be read.
```

Reason: The command for adding the Kubernetes apt repository is a single line command.

When you copy the command for adding the Kubernetes apt repository, it does not get copied in a single line.

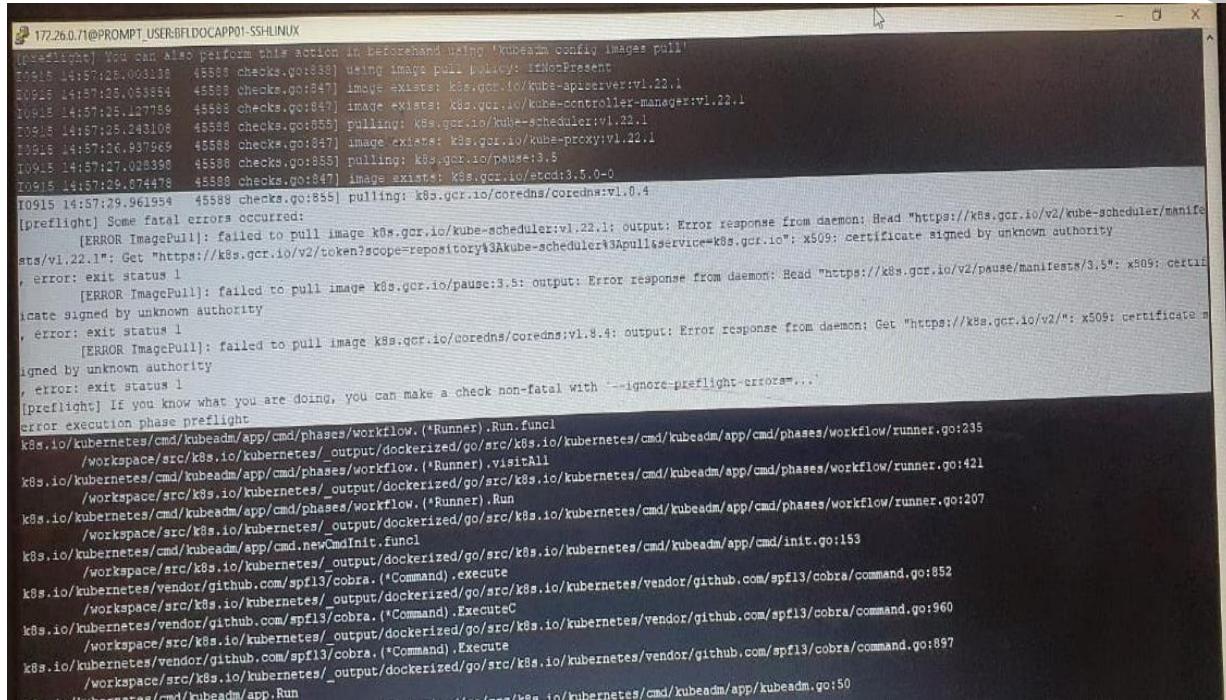
Fix: Try copying the command and paste it somewhere else like in notepad then make it a single line command. Give spaces where marked red then run it.

Note: Copy-Paste or type Next Command Carefully

```
$ echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

11.11 Getting x509 error while build cluster on prem linux server

Issue: Getting x509 error while build cluster on prem linux server.



```

172.26.0.71@PROMPT_USER:~/FILDOCAPP01-SSHLINUX
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
[0915 14:57:25.003188 45588 checks.go:838] using image pull policy: IfNotPresent
[0915 14:57:25.053894 45588 checks.go:847] image exists: k8s.gcr.io/kube-scheduler:v1.22.1
[0915 14:57:25.127759 45588 checks.go:847] image exists: k8s.gcr.io/kube-controller-manager:v1.22.1
[0915 14:57:25.243108 45588 checks.go:858] pulling: k8s.gcr.io/kube-scheduler:v1.22.1
[0915 14:57:26.937969 45588 checks.go:847] image exists: k8s.gcr.io/kube-proxy:v1.22.1
[0915 14:57:27.028398 45588 checks.go:855] pulling: k8s.gcr.io/pause:3.5
[0915 14:57:29.074478 45588 checks.go:847] image exists: k8s.gcr.io/etcd:3.5.0-0
[0915 14:57:29.961954 45588 checks.go:855] pulling: k8s.gcr.io/coredns/coredns:v1.8.4
[0915 14:57:29.961954 45588 checks.go:855] failed to pull image k8s.gcr.io/coredns/coredns:v1.8.4: output: Error response from daemon: Head "https://k8s.gcr.io/v2/kube-scheduler/manifests/v1.22.1": Get "https://k8s.gcr.io/v2/token?scope=repository%2Akube-scheduler%2Apull-service=k8s.gcr.io": x509: certificate signed by unknown authority
, error: exit status 1
[ERROR ImagePull]: failed to pull image k8s.gcr.io/pause:3.5: output: Error response from daemon: Head "https://k8s.gcr.io/v2/pause/manifests/3.5": x509: certificate signed by unknown authority
, error: exit status 1
[ERROR ImagePull]: failed to pull image k8s.gcr.io/etcd:3.5.0-0: output: Error response from daemon: Head "https://k8s.gcr.io/v2/etcd:3.5.0-0": x509: certificate signed by unknown authority
, error: exit status 1
[ERROR ImagePull]: failed to pull image k8s.gcr.io/pause:3.5: output: Error response from daemon: Head "https://k8s.gcr.io/v2/pause/manifests/3.5": x509: certificate signed by unknown authority
, error: exit status 1
[ERROR ImagePull]: failed to pull image k8s.gcr.io/etcd:3.5.0-0: output: Error response from daemon: Head "https://k8s.gcr.io/v2/etcd:3.5.0-0": x509: certificate signed by unknown authority
, error: exit status 1
[ERROR ImagePull]: failed to pull image k8s.gcr.io/coredns/coredns:v1.8.4: output: Error response from daemon: Get "https://k8s.gcr.io/v2/": x509: certificate signed by unknown authority
, error: exit status 1
[preflight] If you know what you are doing, you can make a check non-fatal with --ignore-preflight-errors=...
error execution phase preflight
k8s.io/kubernetes/cmd/kubeadm/app/cmd/phases/workflow.(*Runner).Run.func1
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/cmd/kubeadm/app/cmd/phases/workflow/runner.go:235
k8s.io/kubernetes/cmd/kubeadm/app/cmd/phases/workflow.(*Runner).visitAll
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/cmd/kubeadm/app/cmd/phases/workflow/runner.go:421
k8s.io/kubernetes/cmd/kubeadm/app/cmd/phases/workflow.(*Runner).Run
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/cmd/kubeadm/app/cmd/phases/workflow/runner.go:207
k8s.io/kubernetes/cmd/kubeadm/app/cmd/newCmInit.func1
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/cmd/kubeadm/app/cmd/init.go:153
k8s.io/kubernetes/vendor/github.com/spf13/cobra.(*Command).execute
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/vendor/github.com/spf13/cobra/command.go:852
k8s.io/kubernetes/vendor/github.com/spf13/cobra.(*Command).ExecuteC
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/vendor/github.com/spf13/cobra/command.go:960
k8s.io/kubernetes/vendor/github.com/spf13/cobra.(*Command).Execute
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/vendor/github.com/spf13/cobra/command.go:897
k8s.io/kubernetes/vendor/github.com/spf13/cobra.Run
    /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/cmd/kubeadm/app/kubeadm.go:50
    /k8s.io/kubernetes/cmd/kubeadm/app.Run
        /workspace/src/k8s.io/kubernetes/_output/dockerized/go/src/k8s.io/kubernetes/cmd/kubeadm/app/kubeadm.go:153

```

Fix: Fix this error by adding registry cert to docker.

12 SUMMARY

In this guide we Covered:

- Bootstrap A Kubernetes Cluster Using Kubeadm

K21 Academy