
Pod Scheduling (Node Selector Node Affinity, Anti-Affinity, Taint & Toleration)

[Edition 3]

[Last Update 211018]

For any issues/help contact : support@k21academy.com

Contents

1	Introduction	3
2	Documentation.....	4
2.1	Kubernetes Documentation	4
2.2	Linux Commands and VIM Commands.....	4
3	Pre-Requisite	5
4	Constraining Pods with Node Selector	6
4.1	Adding Label to the nodes in the cluster	6
4.2	Create Deployment With Node Constraints	7
4.3	Clean-Up Resources & Labels	8
5	Constraining Pods with Node Affinity	9
5.1	Create Deployment with Node Affinity preferred constraint	9
5.2	Creating Pod with Node Affinity required constraint.....	11
5.3	Verify pod scheduling.....	12
5.4	Clean-Up Resources & Label added in this Exercise.....	12
6	Constraining Pods with Node Anti-Affinity	13
6.1	Label nodes “worker2” and “worker1”	13
6.2	Create Deployment with Node Anti-Affinity preferred constraint	14
6.3	Delete the deployment.....	15
6.4	Creating Pod with Node Anti-Affinity required constraint.....	15
6.5	Verify Pod Scheduling	16
6.6	Clean-Up Resources & Label added in this Exercise.....	16
7	Advanced Scheduling with Taint and Tolerations	18
7.1	Tainting a Node to Simulate Advanced Scheduling	18
7.2	Creating Pod without Toleration.....	19
7.3	Creating a Pod with Toleration.....	20
7.4	Simulate eviction of Pod using NoSchedule effect	21
8	Troubleshooting	23
8.1	Nodes are remain tainted even after untaint them	23
9	Summary.....	24

1 INTRODUCTION

Assigning Pods to Nodes

You can constrain a Pod to only be able to run on particular Node(s), or to prefer to run on particular nodes.

nodeSelector is the simplest recommended form of node selection constraint. **nodeSelector** is a field of **PodSpec**. It specifies a map of key-value pairs. For the pod to be eligible to run on a node, the node must have each of the indicated key-value pairs as labels (it can have additional labels as well).

Node affinity

Node affinity is conceptually similar to **nodeSelector** -- it allows you to constrain which nodes your pod is eligible to be scheduled on, based on labels on the node.

Pod affinity and anti-affinity

Inter-pod affinity and anti-affinity allow you to constrain which nodes your pod is eligible to be scheduled based on labels on pods that are already running on the node rather than based on labels on nodes.

Taint and Tolerations

Node affinity, is a property of Pods that attracts them to a set of nodes (either as a preference or a hard requirement). Taints are the opposite -- they allow a node to repel a set of pods.

Tolerations are applied to pods, and allow (but do not require) the pods to schedule onto nodes with matching taints.

Taints and tolerations work together to ensure that pods are not scheduled onto inappropriate nodes. One or more taints are applied to a node; this marks that the node should not accept any pods that do not tolerate the taints.

This guide Covers:

- Constraining pods with node selector
- Constraining pods with node affinity
- Constraining pods with node anti-affinity
- Constraining pods with taint and toleration

2 DOCUMENTATION

2.1 Kubernetes Documentation

1. DaemonSet Controller

<https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

2. Assign Pods to Nodes NodeSelector

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-pods-nodes/>

<https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/>

3. Assign Pods to Nodes using Node Affinity

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-pods-nodes-using-node-affinity/>

4. Advanced Scheduling in Kubernetes

<https://kubernetes.io/blog/2017/03/advanced-scheduling-in-kubernetes/#:~:text=Node%20affinity%2Fanti%2Daffinity%20allows,to%20pods%20in%20other%20services%3F>

5. Taints and Tolerations

<https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/>

2.2 Linux Commands and VIM Commands

Note: If you are new to Linux and wanted to learn basic linux for Kubernetes, then drop us a mail at support@k21academy.com and get Bonus **Linux for Beginners** course free.

1. Basic Linux Commands

<https://maker.pro/linux/tutorial/basic-linux-commands-for-beginners>

<https://www.hostinger.in/tutorials/linux-commands>

2. Basic VIM Commands

<https://coderwall.com/p/adv71w/basic-vim-commands-for-getting-started>

3. Popular VIM Commands

<https://www.keycdn.com/blog/vim-commands>

3 PRE-REQUISITE

Ensure that you have completed following three activity guides (or you have an Ubuntu Server)

- **Note:** Follow Activity Guide *AG_Bootstrap_Kubernetes_Cluster_Using_Kubeadm_Guide_ed*** from portal
- **Note:** Follow Activity Guide *AG_Deploy_App_On_Pod_&_Basic_Networking_ed*** from portal
- **Note:** Follow Activity Guide *AG_Deploying_Scalable_and_Configuring_Autoscaling_For_Stateless_Application_ed*** from portal
- **Note:** Follow Activity Guide *AG_Configuring_NFS_Storage_Persistence_Volume_ed*** from portal

4 CONSTRAINING PODS WITH NODE SELECTOR

Note: In below Sections we are going to use YAML files no need write complete yaml file because in CKA exam you can official Kubernetes documentation use Below GIT url to clone repo and use yaml files

```
$ git clone https://github.com/k21academyuk/Kubernetes
$ cd Kubernetes
```

```
root@master:~# cd Kubernetes/
root@master:~/Kubernetes# ls
Dockerfile
README.md
__pycache__
adapter-configmap.yaml
adapter-pod.yaml
app.py
apple.yaml
banana.yaml
config-map.yaml
configmap-pod.yaml
counter-pod.yaml
cron.yaml
daemonset.yaml
demo-pod.yaml
docker-compose.yaml
docker-registry-secret.yaml
dockerfile-mq
elasticsearch-rbac.yaml
elasticsearch-stfullset-oci.yaml
elasticsearch-stfullset.yaml
elasticsearch-svc.yaml
elasticsearch.yaml
example-ingress.yaml
filebeat-agent.yaml
fluentd.yaml
root@master:~/Kubernetes#
foo-allow-to-hello.yaml
guestbook-frontend-svc.yaml
guestbook-frontend.yaml
headlesservice.yaml
hello-allow-from-foo.yaml
ingress-app1.yaml
ingress-app2.yaml
ingress-route.yaml
initcontainer.yaml
job-mq.yaml
job-tmpl.yaml
job.yaml
kibana-elk.yaml
kibana.yaml
label-deployment.yaml
liveness-pod.yaml
logstash-configmap.yaml
logstash-deployment.yaml
logstash-svc.yaml
metrics-server.yaml
multi-container.yaml
multi-pod-configmap.yaml
multi-pod-nginx.yaml
multi-prod-consumer.yaml
namespace.yaml
network-policy.yaml
nfs-pv.yaml
nfs-pvc.yaml
nfs-pv-pod.yaml
nginx-deployment.yaml
nginx-hpa.yaml
nginx-svc.yaml
nodeaffinity-deployment.yaml
nodeaffinity1-deployment.yaml
nodeanti-affinity-deployment.yaml
nodeanti-affinity1-deployment.yaml
oke-admin-service-account.yaml
pod-dynamicpv-oci.yaml
pod-dynamicpv.yaml
podaffinity-deployment.yaml
podaffinity1-deployment.yaml
podanti-affinity-deployment.yaml
podanti-affinity1-deployment.yaml
priv-reg-pod.yaml
pvc-oci.yaml
pvc.yaml
quota-pod.yaml
quota-pod1.yaml
quota.yaml
rabbitmq-deployment.yaml
rabbitmq-service.yaml
readiness-pod.yaml
readiness-svc.yaml
redis-cm.yaml
redis-master-svc.yaml
redis-master.yaml
redis-pod.yaml
redis-slave-svc.yaml
redis-slave.yaml
requirements.txt
role-dev.yaml
rolebind.yaml
script.sh
security-cxt-nonroot.yaml
security-cxt-priv.yaml
security-cxt-readonly.yaml
security-cxt-rmcap.yaml
security-cxt-time.yaml
security-cxt.yaml
statefulset1.yaml
tt-pod.yaml
tt-pod1.yaml
web.yaml
worker.py
```

4.1 Adding Label to the nodes in the cluster

1. Check for the default labels of all the nodes. We would use one of the worker node and label going further

```
$ kubectl get nodes --show-labels
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl get nodes --show-labels
NAME          STATUS    ROLES    AGE   VERSION   LABELS
kubeadm-master Ready    master   5d11h v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubeadm-master,kubernetes.io/os=linux,node-role.kubernetes.io/master=
worker1       Ready    <none>    5d11h v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker1,kubernetes.io/os=linux
worker2       Ready    <none>    5d11h v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker2,kubernetes.io/os=linux
root@kubeadm-master:/home/ubuntu#
```

2. Add label to the worker1 node

```
$ kubectl label nodes worker1 disktype=ssd
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl label nodes worker1 disktype=ssd
node/worker1 labeled
root@kubeadm-master:/home/ubuntu#
```

- View the labels of the nodes again to verify labelling was done as expected

```
$ kubectl get nodes --show-labels
```

```
root@kubeadm-master:/home/ubuntu# kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE   VERSION   LABELS
kubeadm-master      Ready    master   5d11h   v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=kubeadm-master,kubernetes.io/os=linux,node-role.kubernetes.io/master=
worker1             Ready    <none>    5d11h   v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,disktype=ssd,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker1,kubernetes.io/os=linux
worker2             Ready    <none>    5d11h   v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker2,kubernetes.io/os=linux
root@kubeadm-master:/home/ubuntu#
```

4.2 Create Deployment With Node Constraints

- Create deployment with 2 replicas and specify the constraint with **nodeSelector** label of disktype: ssd
- Check the content of label-deployment.yaml file

```
$ vim label-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.12
          ports:
            - containerPort: 80
      nodeSelector:
        disktype: ssd
```

- Create deployment with kubectl create command

```
$ kubectl create -f label-deployment.yaml
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl create -f label-deployment.yaml
deployment.apps/nginx-deployment created
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

- Grep the node details for ssd labelled node


```
$ kubectl get nodes --show-labels | grep ssd
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get nodes --show-labels | grep ssd
worker1      Ready    <none>   5d12h   v1.18.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,disktype=ssd,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker1,kubernetes.io/os=linux
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

5. Verify that both the replicas of the nginx-deployment is scheduled on worker node "worker1" which was labelled as disktype: ssd

```
$ kubectl get pods -o wide
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES
nginx-deployment-5cdb5745d-c7c7p    1/1     Running   0           68s   10.46.0.3     worker1   <none>            <none>
nginx-deployment-5cdb5745d-pd4md     1/1     Running   0           68s   10.46.0.2     worker1   <none>            <none>
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

4.3 Clean-Up Resources & Labels

1. Delete the deployment using kubectl delete command with filename

```
$ kubectl delete -f label-deployment.yaml
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl delete -f label-deployment.yaml
deployment.apps "nginx-deployment" deleted
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES
nginx-deployment-5cdb5745d-c7c7p    0/1     Terminating   0           2m3s   10.46.0.3     worker1   <none>            <none>
nginx-deployment-5cdb5745d-pd4md     1/1     Terminating   0           2m3s   10.46.0.2     worker1   <none>            <none>
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

2. Remove the label added to worker1 node with kubectl label command and verify the label is removed

```
$ kubectl label nodes worker1 disktype-
```

```
$ kubectl get nodes --show-labels | grep ssd
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl label nodes worker1 disktype-
node/worker1 labeled
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get nodes --show-labels | grep ssd
root@kubeadm-master:/home/ubuntu/Kubernetes#
```


5 CONSTRAINING PODS WITH NODE AFFINITY

Note: In below Sections we are going to use YAML files no need write complete yaml file because in CKA exam you can official Kubernetes documentation use Below GIT url to clone repo and use yaml files

```
$ git clone https://github.com/k21academyuk/Kubernetes
$ cd Kubernetes
```

```
root@master:~# cd Kubernetes/
root@master:~/Kubernetes# ls
Dockerfile
README.md
__pycache__
adapter-configmap.yaml
adapter-pod.yaml
app.py
apple.yaml
banana.yaml
config-map.yaml
configmap-pod.yaml
counter-pod.yaml
cron.yaml
daemonset.yaml
demo-pod.yaml
docker-compose.yaml
docker-registry-secret.yaml
dockerfile-mq
elasticsearch-rbac.yaml
elasticsearch-stfullset-oci.yaml
elasticsearch-stfullset.yaml
elasticsearch-svc.yaml
elasticsearch.yaml
example-ingress.yaml
filebeat-agent.yaml
fluentd.yaml
root@master:~/Kubernetes#
foo-allow-to-hello.yaml
guestbook-frontend-svc.yaml
guestbook-frontend.yaml
headlesservice.yaml
hello-allow-from-foo.yaml
ingress-app1.yaml
ingress-app2.yaml
ingress-route.yaml
initcontainer.yaml
job-mq.yaml
job-tmpl.yaml
job.yaml
kibana-elk.yaml
kibana.yaml
label-deployment.yaml
liveness-pod.yaml
logstash-configmap.yaml
logstash-deployment.yaml
logstash-svc.yaml
metrics-server.yaml
multi-container.yaml
multi-pod-configmap.yaml
multi-pod-nginx.yaml
multi-prod-consumer.yaml
namespace.yaml
network-policy.yaml
nfs-pv.yaml
nfs-pvc.yaml
nfs-pv-pod.yaml
nginx-deployment.yaml
nginx-hpa.yaml
nginx-svc.yaml
nodeaffinity-deployment.yaml
nodeaffinity1-deployment.yaml
nodeanti-affinity-deployment.yaml
nodeanti-affinity1-deployment.yaml
oke-admin-service-account.yaml
pod-dynamicpv-oci.yaml
pod-dynamicpv.yaml
podaffinity-deployment.yaml
podaffinity1-deployment.yaml
podanti-affinity-deployment.yaml
podanti-affinity1-deployment.yaml
priv-reg-pod.yaml
pvc-oci.yaml
pvc.yaml
quota-pod1.yaml
quota.yaml
rabbitmq-deployment.yaml
rabbitmq-service.yaml
readiness-pod.yaml
readiness-svc.yaml
redis-cm.yaml
redis-master-svc.yaml
redis-master.yaml
redis-pod.yaml
redis-slave-svc.yaml
redis-slave.yaml
requirements.txt
role-dev.yaml
rolebind.yaml
script.sh
security-cxt-nonroot.yaml
security-cxt-priv.yaml
security-cxt-readonly.yaml
security-cxt-rmcap.yaml
security-cxt-time.yaml
security-cxt.yaml
statefulset1.yaml
tt-pod.yaml
tt-pod1.yaml
web.yaml
worker.py
```

5.1 Create Deployment with Node Affinity preferred constraint

1. Create deployment with 2 replicas and specify the constraint with node affinity constraint defined
2. Check the content of nodeaffinity-deployment.yaml file and see the constraint is **“preferredDuringSchedulingIgnoredDuringExecution”**

```
$ vim nodeaffinity-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.12
        ports:
        - containerPort: 80
      affinity:
        nodeAffinity:
          preferredDuringSchedulingIgnoredDuringExecution:
          - weight: 1
            preference:
              matchExpressions:
              - key: disktype
                operator: In
                values:
                - ssd
```

"nodeaffinity-deployment.yaml" 31L, 619C

3. Create deployment with kubectl create command and verify

```
$ kubectl create -f nodeaffinity-deployment.yaml
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl create -f nodeaffinity-deployment.yaml
deployment.apps/nginx-deployment created
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

```
$ kubectl get deployment
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    2/2     2             2           12s
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

4. List the pods and notice that it has been created despite none of the nodes had the specified label

```
$ kubectl get pods -o wide
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES
nginx-deployment-b956d8fb7-6289t    1/1     Running   0           108s   10.46.0.2     worker1   <none>            <none>
nginx-deployment-b956d8fb7-dv1l8    1/1     Running   0           108s   10.40.0.2     worker2   <none>            <none>
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

5. Delete Deployment

```
$ kubectl delete -f nodeaffinity-deployment.yaml
```

5.2 Creating Pod with Node Affinity required constraint

1. Define “**requiredDuringSchedulingIgnoredDuringExecution**” Constraint
 - a. Check the content of nodeaffinity1-deployment.yaml file and see the constraint is “requiredDuringSchedulingIgnoredDuringExecution”

```
$ vim nodeaffinity1-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.12
        ports:
        - containerPort: 80
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
            - matchExpressions:
              - key: disktype
                operator: In
                values:
                - ssd
```

- b. Create deployment with kubectl create command. List the deployment and pods, see that the pods are in **pending state** as none of the nodes have the required label

```
$ kubectl create -f nodeaffinity1-deployment.yaml
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl create -f nodeaffinity1-deployment.yaml
deployment.apps/nginx-deployment created
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

```
$ kubectl get deployment
```

```
$ kubectl get pods -o wide
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment  0/2     2            0           12s
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP          NODE          NOMINATED NODE   READINESS GATES
nginx-deployment-6d46875998-48649  0/1     Pending   0           25s   <none>      <none>        <none>           <none>
nginx-deployment-6d46875998-7j8ks  0/1     Pending   0           25s   <none>      <none>        <none>           <none>
```

5.3 Verify pod scheduling

1. Label node "worker2" and notice that pending pods get scheduled on the labelled node
2. Label worker node "worker2" disk type=ssd

```
$ kubectl label nodes worker2 disktype=ssd
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl label nodes worker2 disktype=ssd
node/worker2 labeled
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

3. List the deployment and pods to verify pods get scheduled on worker2 node

```
$ kubectl get deployment
```

```
$ kubectl get pods -o wide
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment  2/2     2            2           68s
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP          NODE          NOMINATED NODE   READINESS GATES
nginx-deployment-6d46875998-48649  1/1     Running   0           71s   10.40.0.2   worker2        <none>           <none>
nginx-deployment-6d46875998-7j8ks  1/1     Running   0           71s   10.40.0.3   worker2        <none>           <none>
```

5.4 Clean-Up Resources & Label added in this Exercise

1. Delete the deployment

```
$ kubectl delete -f nodeaffinity1-deployment.yaml
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl delete -f nodeaffinity1-deployment.yaml
deployment.apps "nginx-deployment" deleted
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

2. Remove the label added to worker2 node with kubectl label command and verify the label is removed

```
$ kubectl label nodes worker2 disktype-
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl label nodes worker2 disktype-
node/worker2 labeled
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

6 CONSTRAINING PODS WITH NODE ANTI-AFFINITY

Note: In below Sections we are going to use YAML files no need write complete yaml file because in CKA exam you can official Kubernetes documentation use Below GIT url to clone repo and use yaml files

```
$ git clone https://github.com/k21academyuk/Kubernetes
$ cd Kubernetes
```

```
root@master:~# cd Kubernetes/
root@master:~/Kubernetes# ls
Dockerfile
README.md
__pycache__
adapter-configmap.yaml
adapter-pod.yaml
app.py
apple.yaml
banana.yaml
config-map.yaml
configmap-pod.yaml
counter-pod.yaml
cron.yaml
daemonset.yaml
demo-pod.yaml
docker-compose.yaml
docker-registry-secret.yaml
dockerfile-mq
elasticsearch-rbac.yaml
elasticsearch-stfullset-oci.yaml
elasticsearch-stfullset.yaml
elasticsearch-svc.yaml
elasticsearch.yaml
example-ingress.yaml
filebeat-agent.yaml
fluentd.yaml
root@master:~/Kubernetes#
foo-allow-to-hello.yaml
guestbook-frontend-svc.yaml
guestbook-frontend.yaml
headlesservice.yaml
hello-allow-from-foo.yaml
ingress-app1.yaml
ingress-app2.yaml
ingress-route.yaml
initcontainer.yaml
job-mq.yaml
job-tmpl.yaml
job.yaml
kibana-elk.yaml
kibana.yaml
label-deployment.yaml
liveness-pod.yaml
logstash-configmap.yaml
logstash-deployment.yaml
logstash-svc.yaml
metrics-server.yaml
multi-container.yaml
multi-pod-configmap.yaml
multi-pod-nginx.yaml
multi-prod-consumer.yaml
namespace.yaml
network-policy.yaml
nfs-pv.yaml
nfs-pvc.yaml
nfs-pv-pod.yaml
nginx-deployment.yaml
nginx-hpa.yaml
nginx-svc.yaml
nodeaffinity-deployment.yaml
nodeaffinity1-deployment.yaml
nodeanti-affinity-deployment.yaml
nodeanti-affinity1-deployment.yaml
oke-admin-service-account.yaml
pod-dynamicpv-oci.yaml
pod-dynamicpv.yaml
podaffinity-deployment.yaml
podaffinity1-deployment.yaml
podanti-affinity-deployment.yaml
podanti-affinity1-deployment.yaml
priv-reg-pod.yaml
pvc-oci.yaml
pvc.yaml
quota-pod1.yaml
quota.yaml
rabbitmq-deployment.yaml
rabbitmq-service.yaml
readiness-pod.yaml
readiness-svc.yaml
redis-cm.yaml
redis-master-svc.yaml
redis-master.yaml
redis-pod.yaml
redis-slave-svc.yaml
redis-slave.yaml
requirements.txt
role-dev.yaml
rolebind.yaml
script.sh
security-cxt-nonroot.yaml
security-cxt-priv.yaml
security-cxt-readonly.yaml
security-cxt-rmcap.yaml
security-cxt-time.yaml
security-cxt.yaml
statefulset1.yaml
tt-pod.yaml
tt-pod1.yaml
web.yaml
worker.py
```

6.1 Label nodes “worker2” and “worker1”

1. Label worker node “worker1” and “worker2” disktype=ssd

```
$ kubectl label nodes worker1 disktype=ssd
$ kubectl label nodes worker2 disktype=ssd
```

```
root@kubeadm-master:~/Kubernetes#
root@kubeadm-master:~/Kubernetes# kubectl label nodes worker1 disktype=ssd
node/worker1 labeled
root@kubeadm-master:~/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl label nodes worker2 disktype=ssd
node/worker2 labeled
root@kubeadm-master:/home/ubuntu/Kubernetes#
```


6.2 Create Deployment with Node Anti-Affinity preferred constraint

1. Create deployment with 2 replicas and specify the constraint with node anti-affinity constraint defined
2. Check the content of nodeanti-affinity-deployment.yaml file and see the constraint is **"preferredDuringSchedulingIgnoredDuringExecution"**

```
$ vim nodeanti-affinity-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.12
          ports:
            - containerPort: 80
      affinity:
        nodeAffinity:
          preferredDuringSchedulingIgnoredDuringExecution:
            - weight: 1
              preference:
                matchExpressions:
                  - key: disktype
                    operator: NotIn
                    values:
                      - ssd
```

3. Create deployment with kubectl create command and verify that inspite of the node label pod gets placed as the condition is preferred one

```
$ kubectl create -f nodeanti-affinity-deployment.yaml
```

```
root@kubeadm-master:~/Kubernetes# kubectl create -f nodeanti-affinity-deployment.yaml
deployment.apps/nginx-deployment created
root@kubeadm-master:~/Kubernetes#
```

```
$ kubectl get deployment
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    2/2     2             2           12s
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

- List the pods and notice that it has been created despite none of the nodes had the specified label

```
$ kubectl get pods -o wide
```

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES
nginx-deployment-b956d8fb7-6289t    1/1     Running   0           108s  10.46.0.2     worker1  <none>            <none>
nginx-deployment-b956d8fb7-dv1l8    1/1     Running   0           108s  10.40.0.2     worker2  <none>            <none>
```

6.3 Delete the deployment

```
$ kubectl delete -f nodeanti-affinity-deployment.yaml
```

6.4 Creating Pod with Node Anti-Affinity required constraint

- Check the content of nodeaffinity1-deployment.yaml file and see the constraint is "requiredDuringSchedulingIgnoredDuringExecution"

```
$ vim nodeanti-affinity1-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.12
          ports:
            - containerPort: 80
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: disktype
                    operator: NotIn
                    values:
                      - ssd
```

- Create deployment with kubectl create command. List the deployment and pods, see that the pods are in pending state as all the nodes in the cluster have the label

```
$ kubectl create -f nodeanti-affinity1-deployment.yaml
```



```
root@kubeadm-master:/home/ubuntu/Kubernetes#
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl create -f nodeaffinity1-deployment.yaml
deployment.apps/nginx-deployment created
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

\$ kubectl get deployment

\$ kubectl get pods -o wide

```
root@kubeadm-master:~/Kubernetes# kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment 0/2     2            0           101s
root@kubeadm-master:~/Kubernetes# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP          NODE          NOMINATED NODE   READINESS GATES
nginx-deployment-6ff667f855-ddhxm  0/1     Pending   0           103s  <none>      <none>        <none>           <none>
nginx-deployment-6ff667f855-mjj4v  0/1     Pending   0           103s  <none>      <none>        <none>           <none>
root@kubeadm-master:~/Kubernetes#
```

6.5 Verify Pod Scheduling

1. Remove label from node "worker2" and notice that pending pods get scheduled on it, as its not labelled
2. Remove label from worker node "worker2" disktype-

\$ kubectl label nodes worker2 disktype-

3. List the deployment and pods to verify pods get scheduled on worker2 node

\$ kubectl get pods -o wide

```
root@kubeadm-master:~/Kubernetes# kubectl label nodes worker2 disktype-
node/worker2 labeled
root@kubeadm-master:~/Kubernetes# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP          NODE          NOMINATED NODE   READINESS GATES
nginx-deployment-6ff667f855-ddhxm  1/1     Running   0           3m2s  10.38.0.3   worker2       <none>           <none>
nginx-deployment-6ff667f855-mjj4v  1/1     Running   0           3m2s  10.38.0.2   worker2       <none>           <none>
root@kubeadm-master:~/Kubernetes#
```

6.6 Clean-Up Resources & Label added in this Exercise

1. Delete the deployment

\$ kubectl delete -f nodeanti-affinity1-deployment.yaml

```
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl delete -f nodeaffinity1-deployment.yaml
deployment.apps "nginx-deployment" deleted
root@kubeadm-master:/home/ubuntu/Kubernetes#
```

2. Remove the label added to worker1 node with kubectl label command and verify the label is removed

```
$ kubectl label nodes worker1 disktype-
```

7 ADVANCED SCHEDULING WITH TAINT AND TOLERATIONS

Note: In below Sections we are going to use YAML files no need write complete yaml file because in CKA exam you can official Kubernetes documentation use Below GIT url to clone repo and use yaml files

```
$ git clone https://github.com/k21academyuk/Kubernetes
$ cd Kubernetes
```

```
root@master:~# cd Kubernetes/
root@master:~/Kubernetes# ls
Dockerfile
README.md
__pycache__
adapter-configmap.yaml
adapter-pod.yaml
app.py
apple.yaml
banana.yaml
config-map.yaml
configmap-pod.yaml
counter-pod.yaml
cron.yaml
daemonset.yaml
demo-pod.yaml
docker-compose.yaml
docker-registry-secret.yaml
dockerfile-mq
elasticsearch-rbac.yaml
elasticsearch-stfullset-oci.yaml
elasticsearch-stfullset.yaml
elasticsearch-svc.yaml
elasticsearch.yaml
example-ingress.yaml
filebeat-agent.yaml
fluentd.yaml
foo-allow-to-hello.yaml
guestbook-frontend-svc.yaml
guestbook-frontend.yaml
headlesservice.yaml
hello-allow-from-foo.yaml
ingress-app1.yaml
ingress-app2.yaml
ingress-route.yaml
initcontainer.yaml
job-mq.yaml
job-tmpl.yaml
job.yaml
kibana-elk.yaml
kibana.yaml
label-deployment.yaml
liveness-pod.yaml
logstash-configmap.yaml
logstash-deployment.yaml
logstash-svc.yaml
metrics-server.yaml
multi-container.yaml
multi-pod-configmap.yaml
multi-pod-nginx.yaml
multi-prod-consumer.yaml
namespace.yaml
network-policy.yaml
nfs-pv.yaml
nfs-pvc.yaml
nfs-pv-pod.yaml
nginx-deployment.yaml
nginx-hpa.yaml
nginx-svc.yaml
nodeaffinity-deployment.yaml
nodeaffinity1-deployment.yaml
nodeanti-affinity-deployment.yaml
nodeanti-affinity1-deployment.yaml
oke-admin-service-account.yaml
pod-dynamicpv-oci.yaml
pod-dynamicpv.yaml
podaffinity-deployment.yaml
podaffinity1-deployment.yaml
podanti-affinity-deployment.yaml
podanti-affinity1-deployment.yaml
priv-reg-pod.yaml
pvc-oci.yaml
pvc.yaml
quota-pod.yaml
quota-pod1.yaml
quota.yaml
rabbitmq-deployment.yaml
rabbitmq-service.yaml
readiness-pod.yaml
readiness-svc.yaml
redis-cm.yaml
redis-master-svc.yaml
redis-master.yaml
redis-pod.yaml
redis-slave-svc.yaml
redis-slave.yaml
requirements.txt
role-dev.yaml
rolebind.yaml
script.sh
security-cxt-nonroot.yaml
security-cxt-priv.yaml
security-cxt-readonly.yaml
security-cxt-rmcap.yaml
security-cxt-time.yaml
security-cxt.yaml
statefulset1.yaml
tt-pod.yaml
tt-pod1.yaml
web.yaml
worker.py
```

7.1 Tainting a Node to Simulate Advanced Scheduling

1. View all the nodes in the cluster

```
$ kubectl get nodes
```

```
$ kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
aks-agentpool-40017546-vmss000000  Ready    agent    148m   v1.15.10
aks-agentpool-40017546-vmss000001  Ready    agent    148m   v1.15.10
$
```

2. Taint one of the nodes by using its name

```
$ kubectl taint node aks-agentpool-40017546-vmss000001
disktype=magnetic:NoSchedule
```

```
$
$ kubectl taint node aks-agentpool-40017546-vmss000001 disktype=magnetic:NoSchedule
node/aks-agentpool-40017546-vmss000001 tainted
```

3. Verify that the taint was applied to the desired node

```
$ kubectl describe node aks-agentpool-40017546-vmss000001 | grep -i "taints"
```

```
$ kubectl describe node aks-agentpool-40017546-vmss000001 | grep -i "taints"
Taints:          disktype=magnetic:NoSchedule
$
```

7.2 Creating Pod without Toleration

1. View the content of tt-pod.yaml file and create pod using the yaml file

```
$ vi tt-pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: tt-pod
spec:
  containers:
  - name: nginx
    image: nginx
~
~
~
~
~
```

```
$ kubectl create -f tt-pod.yaml
```

```
$
$ kubectl create -f tt-pod.yaml
pod/tt-pod created
$
```

2. Verify the pod status. Notice that it was scheduled on the node which is not tainted

```
$ kubectl get pods -o wide
```

```
$
$ kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP          NODE                                     NOMINATED NODE   READINESS GATES
tt-pod    1/1     Running   0           83s   10.244.0.13 aks-agentpool-40017546-vmss000000    <none>           <none>
```

3. Delete the pod created in this task

```
$ kubectl delete -f tt-pod.yaml
```

```
$
$ kubectl delete -f tt-pod.yaml
pod "tt-pod" deleted
$
```

7.3 Creating a Pod with Toleration

1. View the content of tt-pod1.yaml file and create pod using the yaml file

```
$ vi tt-pod1.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: tt-pod1
spec:
  containers:
  - name: nginx
    image: nginx
  tolerations:
  - key: "disktype"
    operator: "Equal"
    value: "magnetic"
    effect: "NoSchedule"
~
~
~
~
```

```
$ kubectl create -f tt-pod1.yaml
```

```
$
$ kubectl create -f tt-pod1.yaml
pod/tt-pod1 created
$
```

2. Verify the pod status. Notice that it was scheduled on the tainted node

```
$ kubectl get pods -o wide
```

```
$ kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE                                     NOMINATED NODE   READINESS GATES
tt-pod1   1/1     Running   0           8s    10.244.1.16  aks-agentpool-40017546-vmss000001    <none>           <none>
```

3. Delete the pod created in this task

```
$ kubectl delete -f tt-pod1.yaml
```

4. Delete the taint from the node

```
$ kubectl taint node <node_name> disktype-
$ kubectl describe nodes <node_name> | grep -i taint
```

7.4 Simulate eviction of Pod using NoSchedule effect

1. Again create a pod using tt-pod.yaml file. It doesn't have any toleration defined

```
$ kubectl create -f tt-pod.yaml
```

```
$ kubectl create -f tt-pod.yaml
pod/tt-pod created
$
```

2. Taint the node on which the Pod was scheduled

```
$ kubectl get pods -o wide
```

```
$ kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE                                     NOMINATED NODE   READINESS GATES
tt-pod    1/1     Running   0           10s   10.244.0.14   aks-agentpool-40017546-vmss000000    <none>           <none>
```

```
$ kubectl taint node <node_name> disktype=magnetic:NoExecute
```

```
$ kubectl taint node aks-agentpool-40017546-vmss000000 disktype=magnetic:NoExecute
node/aks-agentpool-40017546-vmss000000 tainted
$
```

```
$ kubectl describe node aks-agentpool-40017546-vmss000000 | grep -i "taints"
Taints:          disktype=magnetic:NoExecute
$
```

3. Verify the pods status again and see that the pod is evicted

```
$ kubectl get pods -o wide
```

4. View recent events to see that the pod was evicted due to the taint

```
$ kubectl get events
```

```
$ kubectl get events | grep tt-pod
45m      Normal      Scheduled      pod/tt-pod      Successfully assigned default/tt-pod to aks-agentpool-40017546-vmss000001
45m      Normal      Pulling        pod/tt-pod      Pulling image "nginx"
45m      Normal      Pulled         pod/tt-pod      Successfully pulled image "nginx"
45m      Normal      Created        pod/tt-pod      Created container nginx
45m      Normal      Started        pod/tt-pod      Started container nginx
44m      Normal      Killing        pod/tt-pod      Stopping container nginx
41m      Normal      Scheduled      pod/tt-pod      Successfully assigned default/tt-pod to aks-agentpool-40017546-vmss000001
41m      Normal      Pulling        pod/tt-pod      Pulling image "nginx"
41m      Normal      Pulled         pod/tt-pod      Successfully pulled image "nginx"
41m      Normal      Created        pod/tt-pod      Created container nginx
41m      Normal      Started        pod/tt-pod      Started container nginx
36m      Normal      Killing        pod/tt-pod      Stopping container nginx
35m      Normal      Scheduled      pod/tt-pod      Successfully assigned default/tt-pod to aks-agentpool-40017546-vmss000001
35m      Normal      Pulling        pod/tt-pod      Pulling image "nginx"
34m      Normal      Pulled         pod/tt-pod      Successfully pulled image "nginx"
34m      Normal      Created        pod/tt-pod      Created container nginx
34m      Normal      Started        pod/tt-pod      Started container nginx
2m13s    Normal      TaintManagerEviction pod/tt-pod      Marking for deletion Pod default/tt-pod
33m      Normal      Killing        pod/tt-pod      Stopping container nginx
```

5. Delete the pod and taint from the node

```
$ kubectl delete -f tt-pod.yaml  
$ kubectl taint node <node_name> disktype-  
$ kubectl describe nodes <node_name> | grep -i taint
```


8 TROUBLESHOOTING

8.1 Nodes are remain tainted even after untaint them

Issue: Nodes are remain tainted even after removing the taint.

```
root@master:/home/azureuser/Kubernetes# kubectl taint nodes worker1
node.kubernetes.io/unreachable:NoExecute-
node/worker1 untainted
root@master:/home/azureuser/Kubernetes# kubectl describe nodes worker1 | grep -i Taints
Taints:          node.kubernetes.io/unreachable:NoExecute
```

Reason: Tainted nodes are down

Fix: Tainted nodes are down, make sure they are in running state. Sometimes it takes few minutes to come up.

9 SUMMARY

In this guide we Covered:

- Constraining pods with node selector
- Constraining pods with node affinity
- Constraining pods with node anti-affinity
- Constraining pods with taint and toleration