

Network Policies and User Security Context

[Edition 4]

[Last Update 210216]

For any issues/help contact : support@k21academy.com

Contents

1	Introduction	3
2	Documentation.....	4
2.1	Kubernetes Documentation	4
2.2	Linux Commands and VIM Commands.....	4
3	Authentication and Authorisation using RBAC.....	5
3.1	Creating Namespace, User & User Credentials.....	5
3.2	Setting up Context for New User	6
3.3	Create RBAC Role and Rolebinding	7
3.4	Clean-up Resources Created this Section	8
4	Configuring Network Policies for Applications	9
4.1	Restrict Incoming Traffic Pods.....	9
4.2	Validating Network Policy	10
4.3	Restrict Outgoing Traffic from Pods	10
4.4	Validating Network Policy	11
4.5	Clean-up the resources created in this Section	12
5	Configuring Container Security Context.....	13
5.1	Defining Security Contexts With Default User.....	13
5.2	Defining Security Contexts With Specific User.....	13
5.3	Defining Security Contexts With non-root User.....	14
5.4	Defining Security Contexts With Privileged Container POD	16
5.5	Defining Security Contexts With Privileged Container POD – add Capability	18
5.6	Defining security contexts with privileged container pod – remove capability	19
5.7	Defining security contexts with privileged container pod – ReadOnly	20
6	Summary.....	22

1 INTRODUCTION

In a cluster using a Kubernetes Container Network Interface (CNI) plug-in that supports Kubernetes network policy, network isolation is controlled entirely by NetworkPolicy Custom Resource (CR) objects. In OpenShift Container Platform 4.1, OpenShift SDN supports using NetworkPolicy in its default network isolation mode.

To specify security settings for a Container, include the securityContext field in the Container manifest. The securityContext field is a SecurityContext object. Security settings that you specify for a Container apply only to the individual Container, and they override settings made at the Pod level when there is overlap. Container settings do not affect the Pod's Volumes.

Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within your organization.

This guide Covers:

- Configuring Network Policies for Applications
- Configuring Container Security Context
- Authentication and Authorisation using RBAC

2 DOCUMENTATION

2.1 Kubernetes Documentation

1. Network Policies

<https://kubernetes.io/docs/concepts/services-networking/network-policies/>

2. Configure a Security Context for a Pod or Container

<https://kubernetes.io/docs/tasks/configure-pod-container/security-context/>

3. Using RBAC Authorization

<https://kubernetes.io/docs/reference/access-authn-authz/rbac/>

4. Authorization Overview

<https://kubernetes.io/docs/reference/access-authn-authz/authorization/>

2.2 Linux Commands and VIM Commands

1. Basic Linux Commands

<https://maker.pro/linux/tutorial/basic-linux-commands-for-beginners>

<https://www.hostinger.in/tutorials/linux-commands>

2. Basic VIM Commands

<https://coderwall.com/p/adv71w/basic-vim-commands-for-getting-started>

3. Popular VIM Commands

<https://www.keycdn.com/blog/vim-commands>

3 AUTHENTICATION AND AUTHORISATION USING RBAC

3.1 Creating Namespace, User & User Credentials

1. Create a new namespace named development

```
$ kubectl create ns development
```

```
root@kubeadm-master:/home/ubuntu#  
root@kubeadm-master:/home/ubuntu# kubectl create ns development  
namespace/development created
```

2. View the current clusters and context available. The context allows us to configure the cluster to use, namespace and user for kubectl commands in an easy and consistent manner.

```
$ kubectl config get-contexts
```

```
root@kubeadm-master:/home/ubuntu# kubectl config get-contexts  
CURRENT   NAME                                CLUSTER   AUTHINFO           NAMESPACE  
*          kubernetes-admin@kubernetes        kubernetes kubernetes-admin  
root@kubeadm-master:/home/ubuntu#
```

3. Create a new user **DevDan** and assign a password to him

```
$ sudo useradd -s /bin/bash DevDan
```

```
$ sudo passwd DevDan
```

```
root@kubeadm-master:/home/ubuntu# sudo useradd -s /bin/bash DevDan  
root@kubeadm-master:/home/ubuntu# sudo passwd DevDan  
Enter new UNIX password:  
Retype new UNIX password:  
passwd: password updated successfully
```

4. Generate a private key for DevDan and Certificate Signing Request (CSR) for DevDan

```
$ openssl genrsa -out DevDan.key 2048
```

```
root@kubeadm-master:/home/ubuntu# openssl genrsa -out DevDan.key 2048  
Generating RSA private key, 2048 bit long modulus (2 primes)  
.....+++++  
.+++++
```

```
$ openssl req -new -key DevDan.key \  
-out DevDan.csr -subj "/CN=DevDan/O=development"
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# openssl req -new -key DevDan.key -out DevDan.csr -subj "/CN=DevDan/O=development"  
Can't load /root/.rnd into RNG  
139953709658560:error:2406F079:random number generator:RAND_load_file:Cannot open file:../crypto/rand/randfile.  
c:88:Filename=/root/.rnd
```

5. Generate a self-signed certificate. Use the CA keys for the Kubernetes cluster and set the certificate expiration.

```
$ sudo openssl x509 -req -in DevDan.csr \  
-CA /etc/kubernetes/pki/ca.crt \  
-CAkey /etc/kubernetes/pki/ca.key \  
-CAcreateserial \  
-out DevDan.crt -days 45
```

```
root@kubeadm-master:/home/ubuntu# sudo openssl x509 -req -in DevDan.csr \  
> -CA /etc/kubernetes/pki/ca.crt \  
> -CAkey /etc/kubernetes/pki/ca.key \  
> -CAcreateserial \  
> -out DevDan.crt -days 45  
Signature ok  
subject=CN = DevDan, O = development  
Getting CA Private Key
```

6. Update the access config file to reference the new key and certificate.

```
$ kubectl config set-credentials DevDan \  
--client-certificate=DevDan.crt \  
--client-key=DevDan.key
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl config set-credentials DevDan \  
> --client-certificate=DevDan.crt \  
> --client-key=DevDan.key  
User "DevDan" set.  
root@kubeadm-master:/home/AzureUser/Kubernetes#  
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

3.2 Setting up Context for New User

1. Create context for DevDan user in the cluster and namespace

```
$ kubectl config set-context DevDan-context \  
--cluster=kubernetes \  
--namespace=development \  
--user=DevDan
```

```
root@kubeadm-master:/home/ubuntu# kubectl config set-context DevDan-context \  
> --cluster=kubernetes \  
> --namespace=development \  
> --user=DevDan  
Context "DevDan-context" created.
```

- Verify the context has been properly set. Attempt to view the Pods inside the DevDan-context. Be aware you will get an error.

```
$ kubectl config get-contexts
```

```
$ kubectl --context=DevDan-context get pods
```

```
root@kubeadm-master:/home/ubuntu# kubectl config get-contexts
```

CURRENT	NAME	CLUSTER	AUTHINFO	NAMESPACE
	DevDan-context	kubernetes	DevDan	development
*	kubernetes-admin@kubernetes	kubernetes	kubernetes-admin	

```
root@kubeadm-master:/home/ubuntu# kubectl --context=DevDan-context get pods
Error from server (Forbidden): pods is forbidden: User "DevDan" cannot list resource "pods" in API group "" in the namespace "development"
```

3.3 Create RBAC Role and Rolebinding

- Create a YAML file to associate RBAC rights to a particular namespace and Role. Create the object. Check white space and for typos if you encounter errors.

```
$ kubectl create -f role-dev.yaml
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl create -f role-dev.yaml
role.rbac.authorization.k8s.io/developer created
```

- Then we create will a RoleBinding to associate the Role we just created with a user. Create the object from the rolebind.yaml file.

```
$ kubectl create -f rolebind.yaml
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl apply -f rolebind.yaml
rolebinding.rbac.authorization.k8s.io/developer-role-binding created
```

- Now let's try list pods and then creating a pod using DevDan-context

```
$ kubectl --context=DevDan-context get pods
```

```
$ kubectl --context=DevDan-context run nginx --image=nginx
```

```
$ kubectl --context=DevDan-context get pods
```

```
root@kubeadm-master:/home/ubuntu# kubectl --context=DevDan-context get pods
No resources found in development namespace.
root@kubeadm-master:/home/ubuntu# kubectl --context=DevDan-context run nginx --image=nginx
pod/nginx created
root@kubeadm-master:/home/ubuntu# kubectl --context=DevDan-context get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	ContainerCreating	0	9s

3.4 Clean-up Resources Created this Section

```
$ kubectl delete pod nginx -n development
$ kubectl delete -f rolebind.yaml
$ kubectl delete -f role-dev.yaml
$ kubectl delete ns development
```


4 CONFIGURING NETWORK POLICIES FOR APPLICATIONS

4.1 Restrict Incoming Traffic Pods

1. Run a simple web server application with label app=hello and expose it internally in the cluster

```
$ kubectl run hello-web --labels app=hello --image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose
```

```
root@kubeadm-master:/home/ubuntu#  
root@kubeadm-master:/home/ubuntu# kubectl run hello-web --labels app=hello \  
> --image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose  
service/hello-web created  
pod/hello-web created
```

2. All inbound traffic by default is allowed. So let's configure a NetworkPolicy to allow traffic to hello-web pods from only pods with label app=foo. All other incoming traffic will be blocked.

```
$ vi hello-allow-from-foo.yaml
```

```
kind: NetworkPolicy  
apiVersion: networking.k8s.io/v1  
metadata:  
  name: hello-allow-from-foo  
spec:  
  policyTypes:  
  - Ingress  
  podSelector:  
    matchLabels:  
      app: hello  
  ingress:  
  - from:  
    - podSelector:  
      matchLabels:  
        app: foo
```

3. Apply this policy to the cluster with kubectl command

```
$ kubectl apply -f hello-allow-from-foo.yaml
```

```
root@kubeadm-master:/home/ubuntu# vim hello-allow-from-foo.yaml
root@kubeadm-master:/home/ubuntu# kubectl apply -f hello-allow-from-foo.yaml
networkpolicy.networking.k8s.io/hello-allow-from-foo created
```

4.2 Validating Network Policy

1. Run a temporary Pod with a different label (**app=other**) and get a shell inside the Pod. Observe that the traffic is **not allowed** and therefore the request times out

```
$ kubectl run -l app=other --image=alpine --restart=Never --rm -i -t test-1
# wget -qO- --timeout=2 http://hello-web:8080
# exit
```

```
root@kubeadm-master:/home/ubuntu# kubectl run -l app=other --image=alpine --restart=Never --rm -i -t test-1
If you don't see a command prompt, try pressing enter.
/ # wget -qO- --timeout=2 http://hello-web:8080
wget: download timed out
/ # exit
pod "test-1" deleted
pod default/test-1 terminated (Error)
```

2. Run a temporary Pod with a different label (app=foo) and get a shell inside the Pod. Observe that the traffic is **allowed**

```
$ kubectl run -l app=foo --image=alpine --restart=Never --rm -i -t test-1
# wget -qO- --timeout=2 http://hello-web:8080
# exit
```

```
root@kubeadm-master:/home/ubuntu# kubectl run -l app=foo --image=alpine --restart=Never --rm -i -t test-1
If you don't see a command prompt, try pressing enter.
/ # wget -qO- --timeout=2 http://hello-web:8080
Hello, world!
Version: 1.0.0
Hostname: hello-web
/ # exit
pod "test-1" deleted
```

4.3 Restrict Outgoing Traffic from Pods

1. All outbound traffic by default is allowed. So let's configure a NetworkPolicy to allow traffic only from pods labelled as app=foo to send traffic only to pods with label app=hello. All other outgoing traffic from app=foo will be blocked.

```
$ vi foo-allow-to-hello.yaml
```

2. Apply this policy to the cluster with kubectl command

```
$ kubectl apply -f foo-allow-to-hello.yaml
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# vim foo-allow-to-hello.yaml
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl apply -f foo-allow-to-hello.yaml
networkpolicy.networking.k8s.io/foo-allow-to-hello created
root@kubeadm-master:/home/ubuntu#
```

4.4 Validating Network Policy

1. Run a temporary Pod with a different label (app=hello-2)

```
$ kubectl run hello-web-2 --labels app=hello-2 \
--image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl run hello-web-2 --labels app=hello-2 \
> --image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose
service/hello-web-2 created
pod/hello-web-2 created
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu#
```

2. Next, run a temporary Pod with app=foo label and get a shell prompt inside the container:

```
$ kubectl run -l app=foo --image=alpine --rm -i -t --restart=Never test-3
```

3. Validate that the Pod can establish connections to hello-web:8080:

```
# wget -qO- --timeout=2 http://hello-web:8080
```

4. Validate that the Pod **cannot** establish connections to hello-web-2:8080:

```
# wget -qO- --timeout=2 http://hello-web-2:8080
```

```
root@kubeadm-master:/home/ubuntu#
root@kubeadm-master:/home/ubuntu# kubectl run -l app=foo --image=alpine --rm -i -t --restart=Never test-3
If you don't see a command prompt, try pressing enter.
/ # wget -qO- --timeout=2 http://hello-web:8080
Hello, world!
Version: 1.0.0
Hostname: hello-web
/ # wget -qO- --timeout=2 http://hello-web-2:8080
wget: download timed out
/ #
```

4.5 Clean-up the resources created in this Section

```
$ kubectl delete -f foo-allow-to-hello.yaml  
$ kubectl delete -f hello-allow-from-foo.yaml
```

5 CONFIGURING CONTAINER SECURITY CONTEXT

5.1 Defining Security Contexts With Default User

Note: It allows you to lock down your containers, so that only certain processes can do certain things. This ensures the stability of your containers and allows you to give control or take it away. In this lesson, we'll go through how to set the security context at the container level and the pod level.

1. Run an alpine container with default security

```
$ kubectl run pod-with-defaults --image alpine --restart Never -- /bin/sleep 999999
```

```
root@kubeadm-master:/#  
root@kubeadm-master:/# kubectl run pod-with-defaults --image alpine --restart Never -- /bin/sleep 999999  
pod/pod-with-defaults created  
root@kubeadm-master:/#
```

2. Check the ID on the container:

```
$ kubectl exec pod-with-defaults id
```

```
root@kubeadm-master:/# kubectl exec pod-with-defaults id  
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.  
uid=0(root) gid=0(root) groups=0(root),1(bin),2(daemon),3(sys),4(adm),6(disk),10(wheel),11(floppy),20(dialout),26(tape),27(video)  
root@kubeadm-master:/#
```

5.2 Defining Security Contexts With Specific User

1. The YAML for a container that runs as a user. View the file security-cxt.yaml

```
$ vim security-cxt.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine-user-context
spec:
  containers:
  - name: main
    image: alpine
    command: ["/bin/sleep", "999999"]
    securityContext:
      runAsUser: 405
~
~
~
~
~
```

2. Create the resource from above yaml file

```
$ kubectl apply -f security-cxt.yaml
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl apply -f security-cxt.yaml
pod/alpine-user-context created
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

3. Check the user context

```
$ kubectl exec alpine-user-context id
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec alpine-user-context id
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.
uid=405(guest) gid=100(users)
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

5.3 Defining Security Contexts With non-root User

1. The YAML for a pod that runs the container as non-root:

```
$ vim security-cxt-nonroot.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine-nonroot
spec:
  containers:
  - name: main
    image: alpine
    command: ["/bin/sleep", "999999"]
    securityContext:
      runAsNonRoot: true
~
~
~
~
~
~
```

2. Create a pod that runs the container as non-root:

```
$ kubectl apply -f security-cxt-nonroot.yaml
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl apply -f security-cxt-nonroot.yaml
pod/alpine-nonroot created
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

3. View more information about the pod error:

```
$ kubectl describe pod alpine-nonroot
```

```
$ kubectl get pods
```



```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl describe pod alpine-nonroot
Name:          alpine-nonroot
Namespace:     default
Priority:       0
Node:          worker2/10.0.0.6
Start Time:    Wed, 10 Jun 2020 02:24:07 +0000
Labels:        <none>
Annotations:   Status: Pending
IP:            10.40.0.4
IPs:
  IP: 10.40.0.4
Containers:
  main:
    Container ID:
    Image:        alpine
    Image ID:
    Port:         <none>
    Host Port:    <none>
    Command:
      /bin/sleep
      999999
    State:        Waiting
      Reason:      CreateContainerConfigError
    Ready:        False
    Restart Count: 0
    Environment:  <none>
    Mounts:
```

```
Events:
  Type     Reason      Age          From              Message
  ----     -
  Normal   Scheduled   23s         default-scheduler Successfully assigned default/alpine-nonroot to worker2
  Normal   Pulling     8s (x3 over 22s) kubelet, worker2 Pulling image "alpine"
  Normal   Pulled      7s (x3 over 21s) kubelet, worker2 Successfully pulled image "alpine"
  Warning   Failed      7s (x3 over 21s) kubelet, worker2 Error: container has runAsNonRoot and image will run as root
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl get pods
NAME          READY   STATUS              RESTARTS   AGE
alpine-nonroot 0/1     CreateContainerConfigError 0           29m
```

5.4 Defining Security Contexts With Privileged Container POD

1. The YAML for a privileged container pod:

```
$ vim security-cxt-priv.yaml
```



```
apiVersion: v1
kind: Pod
metadata:
  name: privileged-pod
spec:
  containers:
  - name: main
    image: alpine
    command: ["/bin/sleep", "999999"]
    securityContext:
      privileged: true
~
~
```

2. Create the privileged container pod:

```
kubectl apply -f security-cxt-priv.yaml
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl create -f security-cxt-priv.yaml
pod/privileged-pod created
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

3. View the devices on the default container:

```
$ kubectl exec -it pod-with-defaults ls /dev
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it pod-with-defaults ls /dev
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.
core          null          shm           termination-log
fd            ptmx          stderr        tty
full          pts           stdin         urandom
mqueue        random        stdout        zero
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

4. View the devices on the privileged pod container:

```
$ kubectl exec -it privileged-pod ls /dev
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it privileged-pod ls /dev
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.
autofs        tty12         ttyS11
bsg           tty13         ttyS12
btrfs-control tty14         ttyS13
core          tty15         ttyS14
cpu_dma_latency tty16         ttyS15
cuse          tty17         ttyS16
ecryptfs      tty18         ttyS17
fb0           tty19         ttyS18
fd            tty2          ttyS19
full          tty20         ttyS2
fuse          tty21         ttyS20
hpet          tty22         ttyS21
hwrng         tty23         ttyS22
input         tty24         ttyS23
kmsg          tty25         ttyS24
kvm           tty26         ttyS25
loong-control tty27         ttyS26
```

5. Try to change the time on a default container pod:

```
$ kubectl exec -it pod-with-defaults -- date +%T -s "12:00:00"
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#  
root@kubeadm-master:/home/AzureUser/Kubernetes#  
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it pod-with-defaults -- date +%T -s "12:00:00"  
date: can't set date: Operation not permitted  
12:00:00  
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

5.5 Defining Security Contexts With Privileged Container POD – add Capability

1. The YAML for a container that will allow you to change the time:

```
$ vim security-cxt-time.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: kernelchange-pod  
spec:  
  containers:  
  - name: main  
    image: alpine  
    command: ["/bin/sleep", "999999"]  
    securityContext:  
      capabilities:  
        add:  
        - SYS_TIME
```

2. Create the pod that will allow you to change the container's time:

```
$ kubectl create -f security-cxt-time.yaml
```

3. Change the time on a container:

```
$ kubectl exec -it kernelchange-pod -- date +%T -s "12:00:00"
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#  
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it kernelchange-pod -- date +%T -s "12:00:00"  
12:00:00
```

5.6 Defining security contexts with privileged container pod – remove capability

1. The YAML for a container that removes capabilities:

```
$ vim security-cxt-rmcap.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: remove-capabilities  
spec:  
  containers:  
  - name: main  
    image: alpine  
    command: ["/bin/sleep", "999999"]  
    securityContext:  
      capabilities:  
        drop:  
        - CHOWN  
~  
~  
~  
~  
~  
~  
~
```

2. Create a pod that's container has capabilities removed:

```
$ kubectl apply -f security-cxt-rmcap.yaml
```

3. Try to change the ownership of a container with removed capability:

```
$ kubectl exec remove-capabilities chown guest /tmp
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#  
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec remove-capabilities chown guest /tmp  
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.  
chown: /tmp: Operation not permitted  
command terminated with exit code 1  
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

5.7 Defining security contexts with privileged container pod – ReadOnly

1. The YAML for a pod container that can't write to the local filesystem:

```
$ vim security-cxt-readonly.yaml
```

2. Create a pod that will not allow you to write to the local container filesystem:

```
$ kubectl apply -f security-cxt-readonly.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: readonly-pod
spec:
  containers:
  - name: main
    image: alpine
    command: ["/bin/sleep", "999999"]
    securityContext:
      readOnlyRootFilesystem: true
    volumeMounts:
    - name: my-volume
      mountPath: /volume
      readOnly: false
  volumes:
  - name: my-volume
    emptyDir:
      ~
      ~
      ~
      ~
      ~
```

3. Try to write to the container filesystem:

```
$ kubectl exec -it readonly-pod touch /new-file
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it readonly-pod touch /new-file
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.
touch: /new-file: Read-only file system
command terminated with exit code 1
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

4. Create a file on the volume mounted to the container:

```
$ kubectl exec -it readonly-pod touch /volume/newfile
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it readonly-pod touch /volume/newfile
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl kubectl exec [POD] -- [COMMAND] instead.
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

5. View the file on the volume that's mounted:

```
$ kubectl exec -it readonly-pod -- ls -la /volume/newfile
```

```
root@kubeadm-master:/home/AzureUser/Kubernetes#
root@kubeadm-master:/home/AzureUser/Kubernetes# kubectl exec -it readonly-pod -- ls -la /volume/newfile
-rw-r--r-- 1 root root 0 Jun 10 03:12 /volume/newfile
root@kubeadm-master:/home/AzureUser/Kubernetes#
```

6 SUMMARY

In this guide we Covered:

- Configuring Network Policies for Applications
- Configuring Container Security Context
- Authentication and Authorisation using RBAC