

# **Linux Foundation CKAD Mock Question**

## 1 INTRODUCTION

In this guide there are mock questions for CKAD exam from **Linux Foundation**. These questions are just to give you an overview about what kind of questions you can expect in the CKAD exam. Also we have practice and mock questions for CKAD exam on our portal. You can access them by going to the portal.

Please refer to the below images and links to find the Practice questions.

1. Easy remember link for the portal: <https://k21academy.com/kubernetesdevself>
2. Kubernetes Module 14: Certified Kubernetes Application Developer (CKAD) Practice Questions and Answers: <https://k21academy.com/ckadm14>

# Kubernetes Module 14: Certified Kubernetes Application Developer (CKAD) Practice Questions and Answers

[< Back to Course](#)

### Lesson Topics

☐ [PDF] Practice Questions and Answers

☐ [PDF] Mock Exam Question & Answers

## Q1 | Namespaces

The DevOps team would like to get the list of all Namespaces in the cluster. Get the list and save it to **/opt/course/1/namespaces**.

**Ans.**

```
$ k get ns > /opt/course/1/namespaces
```

The content should then look like:

```
# /opt/course/1/namespaces
NAME      STATUS AGE
default   Active 150m
earth     Active 76m
jupiter   Active 76m
kube-public Active 150m
kube-system Active 150m
mars       Active 76m
mercury    Active 76m
moon       Active 76m
neptune    Active 76m
pluto      Active 76m
saturn     Active 76m
shell-intern Active 76m
sun        Active 76m
venus      Active 76m
```

## Q2 | Pods

Create a single Pod of image **httpd:2.4.41-alpine** in Namespace **default**. The Pod should be named **pod1** and the container should be named **pod1-container**.

Your manager would like to run a command manually on occasion to output the status of that exact Pod. Please write a command that does this into **/opt/course/2/pod1-status-command.sh**. The command should use **kubectl**.

**Ans.**

```
$ k run # help
# check the export on the very top of this document so we can use $do
$ k run pod1 --image=httpd:2.4.41-alpine $do > 2.yaml
```

```
$ vim 2.yaml
```

Change the container name in **2.yaml** to **pod1-container**:

```
# 2.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod1
  name: pod1
spec:
  containers:
  - image: httpd:2.4.41-alpine
    name: pod1-container # change
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

Then run:

```
$ k create -f 2.yaml
```

```
pod/pod1 created
```

```
$ k get pod
```

NAME	READY	STATUS	RESTARTS	AGE
pod1	0/1	ContainerCreating	0	6s

```
$ k get pod
```

NAME	READY	STATUS	RESTARTS	AGE
pod1	1/1	Running	0	30s

Next create the requested command:

```
vim /opt/course/2/pod1-status-command.sh
```

The content of the command file could look like:

```
# /opt/course/2/pod1-status-command.sh
kubectl -n default describe pod pod1 | grep -i status:
```

Another solution would be using jsonpath:

```
# /opt/course/2/pod1-status-command.sh
kubectl -n default get pod pod1 -o jsonpath="{.status.phase}"
```

To test the command:

```
$ sh /opt/course/2/pod1-status-command.sh
```

Running

### Q3 | Job

Team Neptune needs a Job template located at **/opt/course/3/job.yaml**. This Job should run image **busybox:1.31.0** and execute **sleep 2 && echo done**. It should be in namespace **neptune**, run a total of 3 times and should execute 2 runs in parallel.

Start the Job and check its history. Each pod created by the Job should have the label **id: awesome-job**. The job should be named **neb-new-job** and the container **neb-new-job-container**.

Ans.

```
$ k -n neptun create job -h
```

```
# check the export on the very top of this document so we can use $do
```

```
$ k -n neptune create job neb-new-job --image=busybox:1.31.0 $do > /opt/course/3/job.yaml -- sh
-c "sleep 2 && echo done"
```

```
$ vim /opt/course/3/job.yaml
```

**Note:** Before v1.18 it used to be possible to create a Job with `k -n neptune run neb-new-job --image=busybox:1.31.0 --restart=OnFailure`, but not any longer.

Make the required changes in the yaml:

```
# /opt/course/3/job.yaml
apiVersion: batch/v1
kind: Job
metadata:
  creationTimestamp: null
  name: neb-new-job
  namespace: neptune    # add
spec:
  completions: 3        # add
  parallelism: 2         # add
  template:
    metadata:
      creationTimestamp: null
      labels:             # add
      id: awesome-job    # add
    spec:
      containers:
      - command:
        - sh
        - -c
        - sleep 2 && echo done
        image: busybox:1.31.0
        name: neb-new-job-container # update
        resources: {}
      restartPolicy: Never
status: {}
```

Then to create it:

```
$ k -f /opt/course/3/job.yaml create
# namespace already set in yaml
```

Check Job and Pods, you should see two running parallel at most but three in total:

```
$ k -n neptune get pod,job | grep neb-new-job
```

```
pod/neb-new-job-jhq2g      0/1   ContainerCreating  0    4s
pod/neb-new-job-vf6ts      0/1   ContainerCreating  0    4s
job.batch/neb-new-job      0/3                   4s    5s
```

```
$ k -n neptune get pod,job | grep neb-new-job
```

```
pod/neb-new-job-gm8sz      0/1   ContainerCreating  0    0s
pod/neb-new-job-jhq2g      0/1   Completed          0   10s
pod/neb-new-job-vf6ts      1/1   Running            0   10s
job.batch/neb-new-job      1/3                   10s   11s
```

```
$ k -n neptune get pod,job | grep neb-new-job
```

```
pod/neb-new-job-gm8sz      0/1   ContainerCreating  0    5s
pod/neb-new-job-jhq2g      0/1   Completed          0   15s
pod/neb-new-job-vf6ts      0/1   Completed          0   15s
job.batch/neb-new-job      2/3                   15s   16s
```

```
$ k -n neptune get pod,job | grep neb-new-job
```

```
pod/neb-new-job-gm8sz      0/1   Completed          0   12s
pod/neb-new-job-jhq2g      0/1   Completed          0   22s
pod/neb-new-job-vf6ts      0/1   Completed          0   22s
job.batch/neb-new-job      3/3                   21s   23s
```

**Check history:**

```
$ k -n neptune describe job neb-new-job
```

```
...
Events:
  Type            Reason             Age          From              Message
  ----            -
  Normal          SuccessfulCreate    2m52s        job-controller    Created pod: neb-new-job-jhq2g
  Normal          SuccessfulCreate    2m52s        job-controller    Created pod: neb-new-job-vf6ts
```

Normal SuccessfulCreate 2m42s job-controller Created pod: neb-new-job-gm8sz

**Note:** At the age column we can see that two pods run parallel and the third one after that. Just as it was required in the task.

#### Q4 | Requests and Limits, ServiceAccount

Team Neptune needs 3 Pods of image **httpd:2.4-alpine**, create a Deployment named **neptune-10ab** for this. The containers should be named **neptune-pod-10ab**. Each container should have a memory request of 20Mi and a memory limit of 50Mi.

Team Neptune has its own ServiceAccount **neptune-sa-v2** under which the Pods should run. The Deployment should be in Namespace **neptune**.

**Ans.**

```
$ k -n neptune create deployment -h
# help
```

```
$ k -n neptune create deploy -h
# deploy is short for deployment
```

```
# check the export on the very top of this document so we can use $do
$ k -n neptune create deploy neptune-10ab --image=httpd:2.4-alpine $do > 4.yaml
```

```
vim 4.yaml
```

**Note:** Before 1.18 it used to be possible to use the `kubectl run` to create a deployment with any attributes, but not any longer.



Now make the required changes using vim:

```
# 4.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: neptune-10ab
  name: neptune-10ab
  namespace: neptune
spec:
  replicas: 3          # change
  selector:
    matchLabels:
      app: neptune-10ab
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: neptune-10ab
    spec:
      serviceAccountName: neptune-sa-v2 # add
      containers:
      - image: httpd:2.4-alpine
        name: neptune-pod-10ab # change
        resources:           # add
          limits:            # add
            memory: 50Mi     # add
          requests:          # add
            memory: 20Mi     # add
      status: {}
```

If we don't want to write the resources section manually we could run the following command and copy it manually into our yaml file:

```
$ k run tmp --image=busybox $do --requests=memory=20Mi --limits=memory=50Mi
```

Then create the yaml:

[support@k21academy.com](mailto:support@k21academy.com)



```
1sdkwzTmxjblpwWTJWafKyTnZkVzUwTDNOBGNuWnBZMIV0WVdOamlzVnVkQzV1WVcxblElqb
2libVZ3ZEhWdVpTMXpZUZeyTWIJC0ltdDFZbVZ5Ym1WMFpYTXVhVzh2YzJWeWRtbGpaV0Zq
WTI5MWJuUXZjMIZ5ZG1salpTMWhZMk52ZFc1MExuVnBaQ0k2SWpZMlltUmpOak0yTFRKbFI6
TXROREpoWkMwNE9HRTFMV0ZoWXpGbFpqWmxPVFpsTINJc0luTjFZaUk2SW5ONWMzUmxi
VHB6WlhKMmFXTmxZV05qYjNWdWREchVaWEIwZFc1bE9tNWxjSF1xYm1VdGMyRXRkaklpZIE
uVllnYm9NNENUZDBwZENKNzh3alV3bXRhbGgtMnZzS2pBTnlQc2gtNmd1RXdpdFdfcTVGYnc
1WkhQdHZBZHJMbFB6cE9IRWJBZTRIVU05NUJSR1diWUlkd2p1Tjk1SjBENFJORMtWVXQ0OH
R3b2FrUIY3aC1hUHV3c1FYSGhaWnp5NHlpbUZIRzIVZm1zazVZcjRSVmNHNm4xMzd5LUZIMD
hLOHpaaklQQXNLRHFOQIF0eGctbFp2d1ZNaTZ2aUlocnJ6QVFfzME1CT1Y4Mk9KWUd5Mm8tV
1FWYzBVVWFuQ2Y5NFkzZ1QwWVRpcVF2Y3pZTXM2bno5dXQtWGD3aXRyQlk2VGo5QmdQc
HJBOWtfajVxRXhfTFVvVWIVwUEFpRU43T3pka0pzSThjdHRoMTBseXBJMUFIrnl0M3Q2QUx5cl
FvQk0zOWFiRGZxM0Zrc1ltb2NfV013
kind: Secret
```

This shows the base64 encoded token. To get the encoded one we could pipe it manually through **base64 -d** or we simply do:

```
$ k -n neptune describe secret neptune-sa-v2-token-lwhhl
```

```
...
Data
====
token:
eyJhbGciOiJSUzI1NiIsImtpZCI6Im5aZFdkZDJ2aGNvQ3BqWHZOR1g1b3plcm5JZ0hHNWxTZkwzQnFaaTFad2MifQ.eyJpc3MiOiJrdWJlcm5ldGVzL3NlcnZpY2VhY2NvdW50liwia3ViZXJuZXRlcy5pbj9zZXJ2aWNiYWVudC9uYW1lc3BhY2UiOiJuZXB0dW50liwia3ViZXJuZXRlcy5pbj9zZXJ2aWNiYWVudC9uZW50LiwiZm5lcnZpY2VhY2NvdW50L3NlcnZpY2UtYWNjb3VudC5uYW1lloibmVwdHVuZS1zYS12Milslmt1YmVybmV0ZXMuaW8vc2VydmljZWFiY291bnQvc2VydmljZS1hY2NvdW50LnVpZCI6IjY2YmRjNmM2LTJiYzMtNDJhZC04OGE1LWFhYzFIZjZlOTZINSIsInN1Yil6InN5c3RibTpzZXJ2aWNiYWVudDpuZXB0dW50Om5lcHR1bmUtc2EtdjlifQ.VYgboM4CTd0pdCJ78wjUwmtalh-2vsKjANyPsh-6guEwOtWEq5Fbw5ZHPtvAdrLIPzpOHEbAe4eUM95BRGWbYIdwjuN95J0D4RNfKVUt48twoakRV7h-aPuwsQXHhZZzy4yimFHG9Ufmsk5Yr4RVcG6n137y-FH08K8zZjIPAsKDqNBQtgx-IzvVMi6vilhrrzAQs0MBOV82OJYGy2o-WQVc0UUanCf94Y3gT0YTiqQvczYMs6nz9ut-XgwitrBY6Tj9BgPprA9k_j5qEx_LUUZUpPAiEN7OzdkJsI8ctth10lypl1AeFr43t6ALyrQoBM39abDfq3FksR-oc_WMw
ca.crt:    1066 bytes
namespace: 7 bytes
```

Copy the token (part under **token:**) and paste it using vim.

```
vim /opt/course/5/token
```

File **/opt/course/5/token** should contain the token:

[illegible]

## Q6 | ReadinessProbe

Create a single Pod named pod6 in Namespace default of image **busybox:1.31.0**. The Pod should have a readiness-probe executing **cat /tmp/ready**. It should initially wait 5 and periodically wait 10 seconds. This will set the container ready only if the file **/tmp/ready** exists.

The Pod should run the command `touch /tmp/ready && sleep 1d`, which will create the necessary file to be ready and then idles. Create the Pod and confirm it starts.

**Ans.**

```
$ k run pod6 --image=busybox:1.31.0 $do --command -- sh -c "touch /tmp/ready && sleep 1d" > 6.yaml
```

```
$ vim 6.yaml
```

Search for a readiness-probe example on <https://kubernetes.io/docs>, then copy and alter the relevant section for the task:

```
# 6.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod6
  name: pod6
spec:
  containers:
  - args:
    - sh
    - -c
    - touch /tmp/ready && sleep 1d
    image: busybox:1.31.0
    name: pod6
    resources: {}
    readinessProbe:
      exec:
        command:
        - sh
        - -c
        - cat /tmp/ready
      initialDelaySeconds: 5
      periodSeconds: 10
      dnsPolicy: ClusterFirst
      restartPolicy: Always
    status: {}
```

Then:

```
$ k -f 6.yaml create
```

Running `k get pod6` we should see the job being created and completed:

```
$ k get pod pod6
```

NAME	READY	STATUS	RESTARTS	AGE
pod6	0/1	ContainerCreating	0	2s

```
$ k get pod pod6
```

NAME	READY	STATUS	RESTARTS	AGE
pod6	0/1	Running	0	7s

```
$ k get pod pod6
```

NAME	READY	STATUS	RESTARTS	AGE
pod6	1/1	Running	0	15s

We see that the Pod is finally ready.

### Q7 | Pods, Namespaces

The board of Team Neptune decided to take over control of one e-commerce webserver from Team Saturn. The administrator who once setup this webserver is not part of the organisation any longer. All information you could get was that the e-commerce system is called **my-happy-shop**.

Search for the correct Pod in Namespace **saturn** and move it to Namespace **neptune**. It doesn't matter if you shut it down and spin it up again, it probably hasn't any customers anyways.

**Ans.**

Let's see all those Pods:

```
$ k -n saturn get pod
```

NAME	READY	STATUS	RESTARTS	AGE
webserver-sat-001	1/1	Running	0	111m
webserver-sat-002	1/1	Running	0	111m
webserver-sat-003	1/1	Running	0	111m
webserver-sat-004	1/1	Running	0	111m
webserver-sat-005	1/1	Running	0	111m
webserver-sat-006	1/1	Running	0	111m

The Pod names don't reveal any information. We assume the Pod we are searching has a label or annotation with the name **my-happy-shop**, so we search for it:

```
$ k -n saturn describe pod
```

```
# describe all pods, then manually look for it
```

Or

```
$ k -n saturn get pod -o yaml | grep my-happy-shop -A10
```

We see the webserver we're looking for is **webserver-sat-003**

```
$ k -n saturn get pod webserver-sat-003 -o yaml > 7_webserver-sat-003.yaml # export
$ vim 7_webserver-sat-003.yaml
```

Change the Namespace to **neptune**, also remove the status: section, the token volume, the token **volumeMount** and the **nodeName**, else the new Pod won't start. The final file could look as clean like this:

```
# 7_webserver-sat-003.yaml
apiVersion: v1
kind: Pod
metadata:
  annotations:
    description: this is the server for the E-Commerce System my-happy-shop
  labels:
    id: webserver-sat-003
    name: webserver-sat-003
    namespace: neptune # new namespace here
spec:
  containers:
    - image: nginx:1.16.1-alpine
      imagePullPolicy: IfNotPresent
      name: webserver-sat
      restartPolicy: Always
```

Then we execute:

```
$ k -n neptune create -f 7_webserver-sat-003.yaml
$ k -n neptune get pod | grep webserver
```

```
webserver-sat-003      1/1    Running    0      22s
```

It seems the server is running in Namespace **neptune**, so we can do:

```
$ k -n saturn delete pod webserver-sat-003 --force --grace-period=0
```

Let's confirm only one is running:

```
$ k get pod -A | grep webserver-sat-003
```

```
neptune    webserver-sat-003    1/1    Running    0      6s
```

This should list only one pod called **webserver-sat-003** in Namespace **neptune**, status running.

### Q8 | Deployment, Rollouts

There is an existing Deployment named **api-new-c32** in Namespace **neptune**. A developer did make an update to the Deployment but the updated version never came online. Check the Deployment history and find a revision that works, then rollback to it. Could you tell Team Neptune what the error was so it doesn't happen again?

Ans.

```
$ k -n neptune get deploy # overview
$ k -n neptune rollout -h
$ k -n neptune rollout history -h
```

```
$ k -n neptune rollout history deploy api-new-c32
deployment.extensions/api-new-c32
REVISION  CHANGE-CAUSE
1          <none>
2          kubectrl edit deployment api-new-c32 --namespace=neptune
3          kubectrl edit deployment api-new-c32 --namespace=neptune
4          kubectrl edit deployment api-new-c32 --namespace=neptune
5          kubectrl edit deployment api-new-c32 --namespace=neptune
```

We see 5 revisions, let's check Pod and Deployment status:

```
$ k -n neptune get deploy,pod | grep api-new-c32
deployment.extensions/api-new-c32 3/3 1 3 141m

pod/api-new-c32-65d998785d-jtmqq 1/1 Running 0 141m
pod/api-new-c32-686d6f6b65-mj2fp 1/1 Running 0 141m
pod/api-new-c32-6dd45bdb68-2p462 1/1 Running 0 141m
pod/api-new-c32-7d64747c87-zh648 0/1 ImagePullBackOff 0 141m
```

Let's check the pod for errors:

```
$ k -n neptune describe pod api-new-c32-7d64747c87-zh648 | grep -i error
... Error: ImagePullBackOff
```



```
$ k -n neptune describe pod api-new-c32-7d64747c87-zh648 | grep -i image
```

```
Image:      nginx:1.16.3
Image ID:
Reason:     ImagePullBackOff
Warning Failed 4m28s (x616 over 144m) kubelet, gke-s3ef67020-28c5-45f7--default-pool-248abd4f-s010 Error: ImagePullBackOff
```

Someone seems to have added a new image with a spelling mistake in the name `nginx:1.16.3`, that's the reason we can tell Team Neptune!

Now let's revert to the previous version:

```
$ k -n neptune rollout undo deploy api-new-c32
```

Does this one work?

```
$ k -n neptune get deploy api-new-c32
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
api-new-c32	3/3	3	3	146m

Yes! All up-to-date and available.

Also a fast way to get an overview of the ReplicaSets of a Deployment and their images could be done with:

```
$ k -n neptune get rs -o wide | grep api-new-c32
```

## Q9 | Pod -> Deployment

In Namespace `pluto` there is single Pod named **holy-api**. It has been working okay for a while now but Team Pluto needs it to be more reliable. Convert the Pod into a Deployment with 3 replicas and name `holy-api`. The raw Pod template file is available at `/opt/course/9/holy-api-pod.yaml`.

Please create the Deployment and save its yaml under `/opt/course/9/holy-api-deployment.yaml`.

### Ans.

There are multiple ways to do this, one is to copy an Deployment example from <https://kubernetes.io/docs> and then merge it with the existing Pod yaml. That's what we will do now:

```
$ cp /opt/course/9/holy-api-pod.yaml /opt/course/9/holy-api-deployment.yaml # make a copy!
$ vim /opt/course/9/holy-api-deployment.yaml
```

Now copy/use a Deployment example yaml and put the Pod's metadata: and spec: into the Deployment's template: section:

```
# /opt/course/9/holy-api-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: holy-api      # name stays the same
  namespace: pluto    # important
spec:
  replicas: 3         # 3 replicas
  selector:
    matchLabels:
      id: holy-api    # set the correct selector
  template:
    # => from here down its the same as the pods metadata: and spec: sections
    metadata:
      labels:
        id: holy-api
        name: holy-api
    spec:
      containers:
        - env:
            - name: CACHE_KEY_1
              value: b&MTCi0=[T66RXm!jO@
            - name: CACHE_KEY_2
              value: PCAILGej5Ld@Q%{Q1=#
            - name: CACHE_KEY_3
              value: 2qz-]2OJIWDSTn_;RFQ
          image: nginx:1.17.3-alpine
          name: holy-api-container
          volumeMounts:
            - mountPath: /cache1
```

```

    name: cache-volume1
  - mountPath: /cache2
    name: cache-volume2
  - mountPath: /cache3
    name: cache-volume3
volumes:
  - emptyDir: {}
    name: cache-volume1
  - emptyDir: {}
    name: cache-volume2
  - emptyDir: {}
    name: cache-volume3

```

To indent multiple lines using vim you should set the shiftwidth using **:set shiftwidth=2**. Then mark multiple lines using **Shift v** and the up/down keys.

To then indent the marked lines press **>** or **<** and to repeat the action press **.**

Next create the new Deployment:

```
$ k -f /opt/course/9/holy-api-deployment.yaml create
```

and confirm its running:

```
$ k -n pluto get pod | grep holy
```

NAME	READY	STATUS	RESTARTS	AGE
holy-api	1/1	Running	0	19m
holy-api-5dbfdb4569-8qr5x	1/1	Running	0	30s
holy-api-5dbfdb4569-b5clh	1/1	Running	0	30s
holy-api-5dbfdb4569-rj2gz	1/1	Running	0	30s

Finally delete the single Pod:

```
$ k -n pluto delete pod holy-api --force --grace-period=0
```

```
$ k -n pluto get pod,deployment | grep holy
```

pod/holy-api-5dbfdb4569-8qr5x	1/1	Running	0	2m4s
pod/holy-api-5dbfdb4569-b5clh	1/1	Running	0	2m4s
pod/holy-api-5dbfdb4569-rj2gz	1/1	Running	0	2m4s
deployment.extensions/holy-api	3/3	3	3	2m4s

### Q10 | Service, Logs

Team Pluto needs a new cluster internal Service. Create a ClusterIP Service named project-plt-6cc-svc in Namespace pluto. This Service should expose a single Pod named project-plt-6cc-api of image **nginx:1.17.3-alpine**, create that Pod as well. The Pod should be identified by label project: plt-6cc-api. The Service should use tcp port redirection of **3333:80**.

Finally use for example curl from a temporary nginx:alpine Pod to get the response from the Service. Write the response into **/opt/course/10/service\_test.html**. Also check if the logs of Pod **project-plt-6cc-api** show the request and write those into **/opt/course/10/service\_test.log**.

**Ans.**

```
$ k -n pluto run project-plt-6cc-api --image=nginx:1.17.3-alpine --labels project=plt-6cc-api
```

This will create the requested Pod. In yaml it would look like this:

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    project: plt-6cc-api
  name: project-plt-6cc-api
spec:
  containers:
    - image: nginx:1.17.3-alpine
      name: project-plt-6cc-api
      resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

Next we create the service:

```
$ k -n pluto expose pod -h # help
```

```
$ k -n pluto expose pod project-plt-6cc-api --name project-plt-6cc-svc --port 3333 --target-port 80
```

Expose will create a yaml where everything is already set for our case and no need to change anything:

```
apiVersion: v1
```

```
kind: Service
metadata:
  creationTimestamp: null
  labels:
    project: plt-6cc-api
  name: project-plt-6cc-svc # good
  namespace: pluto # great
spec:
  ports:
    - port: 3333 # awesome
      protocol: TCP
      targetPort: 80 # nice
  selector:
    project: plt-6cc-api # beautiful
status:
  loadBalancer: {}
```

We could also use create service but then we would need to change the yaml afterwards:

```
$ k -n pluto create service -h # help
$ k -n pluto create service clusterip -h #help
$ k -n pluto create service clusterip project-plt-6cc-svc --tcp 3333:80 $do
# now we would need to set the correct selector labels
```

Check the Service is running:

```
$ k -n pluto get pod,svc | grep 6cc
```

pod/project-plt-6cc-api	1/1	Running	0	9m42s
service/project-plt-6cc-svc	ClusterIP	10.31.241.234	<none>	3333/TCP 2m24s

Does the Service has one Endpoint?

```
$ k -n pluto describe svc project-plt-6cc-svc
```

```
Name:          project-plt-6cc-svc
Namespace:     pluto
Labels:        project=plt-6cc-api
Annotations:   <none>
Selector:      project=plt-6cc-api
Type:          ClusterIP
```

```
IP:          10.3.244.240
Port:        <unset> 3333/TCP
TargetPort:  80/TCP
Endpoints:   10.28.2.32:80
Session Affinity: None
Events:      <none>
```

```
k -n pluto get ep
```

NAME	ENDPOINTS	AGE
project-plt-6cc-svc	10.28.2.32:80	84m

Yes, endpoint there! Finally we check the connection using a temporary Pod:

```
k run tmp --restart=Never --rm --image=nginx:alpine -i -- curl http://project-plt-6cc-svc.pluto:3333
```

```
% Total  % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed

100  612  100  612    0     0   32210      0 --:--:-- --:--:-- --:--:-- 32210
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
...
```

Great! Notice that we use the Kubernetes Namespace dns resolving (project-plt-6cc-svc.pluto) here. We could only use the Service name if we would also spin up the temporary Pod in Namespace pluto .

And now really finally copy or pipe the html content into **/opt/course/10/service\_test.html**.

```
# /opt/course/10/service_test.html
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
...

```

Also the requested logs:

```
$ k -n pluto logs project-plt-6cc-api > /opt/course/10/service_test.log
# /opt/course/10/service_test.log
10.44.0.0 - - [22/Jan/2021:23:19:55 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.69.1" "-"

```

### Q11 | Labels, Annotations

Team Sunny needs to identify some of their Pods in namespace sun. They ask you to add a new label **protected: true** to all Pods with an existing label **type: worker** or **type: runner**. Also add an annotation **protected: do not delete this pod** to all Pods having the new label **protected: true**.

Ans.

```
$ k -n sun get pod --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
0509649a	1/1	Running	0	25s	type=runner,type_old=messenger
0509649b	1/1	Running	0	24s	type=worker
1428721e	1/1	Running	0	23s	type=worker
1428721f	1/1	Running	0	22s	type=worker
43b9a	1/1	Running	0	22s	type=test
4c09	1/1	Running	0	21s	type=worker
4c35	1/1	Running	0	20s	type=worker
4fe4	1/1	Running	0	19s	type=worker
5555a	1/1	Running	0	19s	type=messenger
86cda	1/1	Running	0	18s	type=runner
8d1c	1/1	Running	0	17s	type=messenger
a004a	1/1	Running	0	16s	type=runner
a94128196	1/1	Running	0	15s	type=runner,type_old=messenger
afd79200c56a	1/1	Running	0	15s	type=worker
b667	1/1	Running	0	14s	type=worker
fdb2	1/1	Running	0	13s	type=worker

If we would only like to get pods with certain labels we can run:

```
$ k -n sun get pod -l type=runner # only pods with label runner
```

We can use this label filtering also when using other commands, like setting new labels:

```
$ k label -h # help
```

```
$ k -n sun label pod -l type=runner protected=true # run for label runner
```

```
$ k -n sun label pod -l type=worker protected=true # run for label worker
```

Or we could run:

```
$ k -n sun label pod -l "type in (worker,runner)" protected=true
```

Let's check the result:

```
k -n sun get pod --show-labels
```

NAME	...	AGE	LABELS
0509649a	...	56s	protected=true,type=runner,type_old=messenger
0509649b	...	55s	protected=true,type=worker
1428721e	...	54s	protected=true,type=worker
1428721f	...	53s	protected=true,type=worker
43b9a	...	53s	type=test
4c09	...	52s	protected=true,type=worker
4c35	...	51s	protected=true,type=worker
4fe4	...	50s	protected=true,type=worker
5555a	...	50s	type=messenger
86cda	...	49s	protected=true,type=runner
8d1c	...	48s	type=messenger
a004a	...	47s	protected=true,type=runner
a94128196	...	46s	protected=true,type=runner,type_old=messenger
afd79200c56a	...	46s	protected=true,type=worker
b667	...	45s	protected=true,type=worker
fdb2	...	44s	protected=true,type=worker

Looking good. Finally we set the annotation using the newly assigned label **protected: true**:

```
$ k -n sun annotate pod -l protected=true protected="do not delete this pod"
```

Not requested in the task but for your own control you could run:



```
$ k -n sun get pod -l protected=true -o yaml | grep -A 8 metadata:
```

### Q12 | Storage, PV, PVC, Pod volume

Create a new PersistentVolume named **earth-project-earthflower-pv**. It should have a capacity of 2Gi, accessMode ReadWriteOnce, hostPath /Volumes/Data and no storageClassName defined.

Next create a new PersistentVolumeClaim in Namespace earth named earth-project-earthflower-pvc . It should request 2Gi storage, accessMode ReadWriteOnce and should not define a storageClassName. The PVC should bound to the PV correctly.

Finally create a new Deployment project-earthflower in Namespace earth which mounts that volume at **/tmp/project-data**. The Pods of that Deployment should be of image **httpd:2.4.41-alpine**.

**Ans.**

```
$ vim 12_pv.yaml
```

Find an example from <https://kubernetes.io/docs> and alter it:

```
# 12_pv.yaml
kind: PersistentVolume
apiVersion: v1
metadata:
  name: earth-project-earthflower-pv
spec:
  capacity:
    storage: 2Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/Volumes/Data"
```

Then create it:

```
$ k -f 12_pv.yaml create
```

Next the PersistentVolumeClaim:

```
$ vim 12_pvc.yaml
```

Find an example from <https://kubernetes.io/docs> and alter it:

```
# 12_pvc.yaml
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: earth-project-earthflower-pvc
  namespace: earth
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
```

Then create:

```
$ k -f 12_pvc.yaml create
```

And check that both have the status Bound:

```
$ k -n earth get pv,pvc
```

NAME	CAPACITY	ACCESS MODES	...	STATUS	CLAIM
persistentvolume/...earthflower-pv	2Gi	RWO	...	Bound	...er-pvc

  

NAME	STATUS	VOLUME	CAPACITY
persistentvolumeclaim/...earthflower-pvc	Bound	earth-project-earthflower-pv	2Gi

Next we create a Deployment and mount that volume:

```
$ k -n earth create deploy project-earthflower --image=httpd:2.4.41-alpine $do > 12_dep.yaml
$ vim 12_dep.yaml
```

Alter the yaml to mount the volume:

```
# 12_dep.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: project-earthflower
  name: project-earthflower
  namespace: earth
spec:
  replicas: 1
```

```
selector:
  matchLabels:
    app: project-earthflower
strategy: {}
template:
  metadata:
    creationTimestamp: null
    labels:
      app: project-earthflower
  spec:
    volumes:
      # add
      - name: data
        # add
        persistentVolumeClaim:
          # add
          claimName: earth-project-earthflower-pvc # add
    containers:
      - image: httpd:2.4.41-alpine
        name: container
        volumeMounts:
          # add
          - name: data
            # add
            mountPath: /tmp/project-data # add
```

```
$ k -f 12_dep.yaml create
```

We can confirm its mounting correctly:

```
$ k -n earth describe pod project-earthflower-d6887f7c5-pn5wv | grep -A2 Mounts:
```

```
Mounts:
  /tmp/project-data from data (rw) # there it is
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-n2sjj (ro)
```

### **Q13 | Storage, StorageClass, PVC**

Team Moonpie, which has the Namespace moon, needs more storage. Create a new PersistentVolumeClaim named moon-pvc-126 in that namespace. This claim should use a new StorageClass moon-retain with the provisioner set to moon-retainer and the reclaimPolicy set to Retain. The claim should request storage of 3Gi, an accessMode of ReadWriteOnce and should use the new StorageClass.

The provisioner moon-retainer will be created by another team, so it's expected that the PVC will not boot yet. Confirm this by writing the log message from the PVC into file **/opt/course/13/pvc-126-reason**.

**Ans.**

```
vim 13_sc.yaml
```

Head to <https://kubernetes.io/docs> , search for "storageclass" and alter the example code to this:

```
# 13_sc.yaml
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: moon-retain
provisioner: moon-retainer
reclaimPolicy: Retain
```

```
$ k -f 13_pvc.yaml create
```

Next we check the status of the PVC :

```
$ k -n moon get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
moon-pvc-126	Pending			moon-retain	2m57s	

```
$ k -n moon describe pvc moon-pvc-126
```

```
Name:      moon-pvc-126
```

```
...
```

```
Status:    Pending
```

```
...
```

```
Events:
```

```
...
```

```
waiting for a volume to be created, either by external provisioner "moon-retainer" or manually created by system administrator
```

This confirms that the PVC waits for the provisioner moon-retainer to be created. Finally we copy or write the event message into the requested location:

```
# /opt/course/13/pvc-126-reason
```

waiting for a volume to be created, either by external provisioner "moon-retainer" or manually created by system administrator

#### Q14 | Secret, Secret-Volume, Secret-Env

You need to make changes on an existing Pod in Namespace moon called secret-handler. Create a new Secret secret1 which contains user=test and pass=pwd. The Secret's content should be available in Pod secret-handler as environment variables **SECRET1\_USER** and **SECRET1\_PASS**. The yaml for Pod secret-handler is available at **/opt/course/14/secret-handler.yaml**.

There is existing yaml for another Secret at **/opt/course/14/secret2.yaml**, create this Secret and mount it inside the same Pod at /tmp/secret2. Your changes should be saved under **/opt/course/14/secret-handler-new.yaml**. Both Secrets should only be available in Namespace moon.

**Ans.**

```
$ k -n moon get pod # show pods
$ k -n moon create secret -h # help
$ k -n moon create secret generic -h # help
$ k -n moon create secret generic secret1 --from-literal user=test --from-literal pass=pwd
```

The last command would generate this yaml:

```
apiVersion: v1
data:
  pass: cHdk
  user: dGVzdA==
kind: Secret
metadata:
  creationTimestamp: null
  name: secret1
  namespace: moon
```

Next we create the second Secret from the given location, making sure it'll be created in Namespace moon:

```
$ k -n moon -f /opt/course/14/secret2.yaml create
$ k -n moon get secret
```

NAME	TYPE	DATA	AGE
------	------	------	-----

```
default-token-rvzcf  kubernetes.io/service-account-token  3  66m
secret1              Opaque                             2  4m3s
secret2              Opaque                             1  8s
```

We will now edit the Pod yaml:

```
$ cp /opt/course/14/secret-handler.yaml /opt/course/14/secret-handler-new.yaml
$ vim /opt/course/14/secret-handler-new.yaml
```

Add the following to the yaml:

```
# /opt/course/14/secret-handler-new.yaml
apiVersion: v1
kind: Pod
metadata:
  labels:
    id: secret-handler
    uuid: 1428721e-8d1c-4c09-b5d6-afd79200c56a
    red_id: 9cf7a7c0-fdb2-4c35-9c13-c2a0bb52b4a9
    type: automatic
  name: secret-handler
  namespace: moon
spec:
  volumes:
    - name: cache-volume1
      emptyDir: {}
    - name: cache-volume2
      emptyDir: {}
    - name: cache-volume3
      emptyDir: {}
    - name: secret2-volume      # add
      secret:                  # add
        secretName: secret2   # add
  containers:
    - name: secret-handler
      image: bash:5.0.11
      args: ['bash', '-c', 'sleep 2d']
      volumeMounts:
        - mountPath: /cache1
          name: cache-volume1
        - mountPath: /cache2
          name: cache-volume2
        - mountPath: /cache3
```

```

name: cache-volume3
- name: secret2-volume      # add
  mountPath: /tmp/secret2   # add
env:
- name: SECRET_KEY_1
  value: ">8$kH#kj..i8}HlmQd{"
- name: SECRET_KEY_2
  value: "IO=a4L/XkRdvN8jM=Y+"
- name: SECRET_KEY_3
  value: "-7PA0_Z]>{pwa43r)___"
- name: SECRET1_USER      # add
  valueFrom:              # add
    secretKeyRef:         # add
      name: secret1       # add
      key: user           # add
- name: SECRET1_PASS      # add
  valueFrom:              # add
    secretKeyRef:         # add
      name: secret1       # add
      key: pass           # add

```

Then we apply the changes:

```

$ k -f /opt/course/14/secret-handler.yaml delete --force --grace-period=0
$ k -f /opt/course/14/secret-handler-new.yaml create

```

Instead of running delete and create we can also use recreate:

```

$ k -f /opt/course/14/secret-handler-new.yaml replace --force --grace-period=0

```

It was not requested directly, but you should always confirm its working:

```

k -n moon exec secret-handler -- env | grep SECRET1

```

```

SECRET1_USER=test
SECRET1_PASS=pwd

```

```

$ k -n moon exec secret-handler -- find /tmp/secret2

```

```

/tmp/secret2
/tmp/secret2/..data
/tmp/secret2/key

```

```
/tmp/secret2/..2019_09_11_09_03_08.147048594
/tmp/secret2/..2019_09_11_09_03_08.147048594/key
```

```
$ k -n moon exec secret-handler -- cat /tmp/secret2/key
12345678
```

### Q15 | ConfigMap, Configmap-Volume

Team Moonpie has a nginx server Deployment called web-moon in Namespace moon. Someone started configuring it but it was never completed. To complete please create a ConfigMap called configmap-web-moon-html containing the content of file **/opt/course/15/web-moon.html** under the data key-name **index.html**.

The Deployment web-moon is already configured to work with this ConfigMap and serve its content. Test the nginx configuration for example using curl from a temporary nginx:alpine Pod.

**Ans.**

Let's check the existing Pods:

```
$ k -n moon get pod
```

NAME	READY	STATUS	RESTARTS	AGE
secret-handler	1/1	Running	0	55m
web-moon-847496c686-2rzj4	0/1	ContainerCreating	0	33s
web-moon-847496c686-9nwwj	0/1	ContainerCreating	0	33s
web-moon-847496c686-cxdbx	0/1	ContainerCreating	0	33s
web-moon-847496c686-hvqlw	0/1	ContainerCreating	0	33s
web-moon-847496c686-tj7ct	0/1	ContainerCreating	0	33s

```
$ k -n moon describe pod web-moon-847496c686-2rzj4
```

```
...
Warning FailedMount 31s (x7 over 63s) kubelet, gke-test-default-pool-ce83a51a-p6s4
MountVolume.SetUp failed for volume "html-volume" : configmaps "configmap-web-moon-html"
not found
```

Good so far, now let's create the missing ConfigMap:

```
$ k -n moon create configmap -h # help
```



```
$ k -n moon create configmap configmap-web-moon-html --from-  
file=index.html=/opt/course/15/web-moon.html # important to set the index.html key
```

This should create a ConfigMap with yaml like:

```
apiVersion: v1  
data:  
  index.html: | # notice the key index.html, this will be the filename when mounted  
    <!DOCTYPE html>  
    <html lang="en">  
    <head>  
      <meta charset="UTF-8">  
      <title>Web Moon Webpage</title>  
    </head>  
    <body>  
      This is some great content.  
    </body>  
    </html>  
kind: ConfigMap  
metadata:  
  creationTimestamp: null  
  name: configmap-web-moon-html  
  namespace: moon
```

After waiting a bit or deleting/recreating (**k -n moon rollout restart deploy web-moon**) the Pods we should see:

```
$ k -n moon get pod
```

NAME	READY	STATUS	RESTARTS	AGE
secret-handler	1/1	Running	0	59m
web-moon-847496c686-2rzj4	1/1	Running	0	4m28s
web-moon-847496c686-9nwwj	1/1	Running	0	4m28s
web-moon-847496c686-cxdbx	1/1	Running	0	4m28s
web-moon-847496c686-hvqlw	1/1	Running	0	4m28s
web-moon-847496c686-tj7ct	1/1	Running	0	4m28s

Looking much better. Finally we check if the nginx returns the correct content:

```
$ k -n moon get pod -o wide # get pod cluster IPs
```

Then use one IP to test the configuration:

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl 10.44.0.78
```

```
% Total    % Received % Xferd Average Speed  Time  Time  Time  Current
           Dload  Upload  Total    Spent  Left  Speed
100 161 100 161    0    0 80500    0 --:--:-- --:--:-- --:--:-- 157k
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>Web Moon Webpage</title>
</head>
<body>
This is some great content.
</body>
```

For debugging or further checks we could find out more about the Pods volume mounts:

```
$ k -n moon describe pod web-moon-c77655cc-dc8v4 | grep -A2 Mounts:
```

```
Mounts:
  /usr/share/nginx/html from html-volume (rw)
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-rvzcf (ro)
```

And check the mounted folder content:

```
$ k -n moon exec web-moon-c77655cc-dc8v4 find /usr/share/nginx/html
```

```
/usr/share/nginx/html
/usr/share/nginx/html/..2019_09_11_10_05_56.336284411
/usr/share/nginx/html/..2019_09_11_10_05_56.336284411/index.html
/usr/share/nginx/html/..data
/usr/share/nginx/html/index.html
```

Here it was important that the file will have the name index.html and not the original one web-moon.html which is controlled through the ConfigMap data key.

## **Q16 | Logging sidecar**

The Tech Lead of Mercury2D decided its time for more logging, to finally fight all these missing data incidents. There is an existing container named cleaner-con in Deployment cleaner in Namespace mercury. This container mounts a volume and writes logs into a file called cleaner.log.

The yaml for the existing Deployment is available at **/opt/course/16/cleaner.yaml**. Persist your changes at **/opt/course/16/cleaner-new.yaml** but also make sure the Deployment is running.

Create a sidecar container named logger-con, image busybox:1.31.0, which mounts the same volume and writes the content of cleaner.log to stdout, you can use the tail **-f** command for this. This way it can be picked up by kubectl logs.

Check if the logs of the new container reveal something about the missing data incidents.

**Ans.**

```
$ cp /opt/course/16/cleaner.yaml /opt/course/16/cleaner-new.yaml
$ vim /opt/course/16/cleaner-new.yaml
```

Add a sidecar container which outputs the log file to stdout:

```
# /opt/course/16/cleaner-new.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  name: cleaner
  namespace: mercury
spec:
  replicas: 2
  selector:
    matchLabels:
      id: cleaner
  template:
    metadata:
      labels:
        id: cleaner
    spec:
      volumes:
      - name: logs
        emptyDir: {}
      initContainers:
      - name: init
        image: bash:5.0.11
```

```
command: ['bash', '-c', 'echo init > /var/log/cleaner/cleaner.log']
volumeMounts:
- name: logs
  mountPath: /var/log/cleaner
containers:
- name: cleaner-con
  image: bash:5.0.11
  args: ['bash', '-c', 'while true; do echo `date`: "remove random file" >>
/var/log/cleaner/cleaner.log; sleep 1; done']
  volumeMounts:
  - name: logs
    mountPath: /var/log/cleaner
- name: logger-con                                # add
  image: busybox:1.31.0                            # add
  command: ["sh", "-c", "tail -f /var/log/cleaner/cleaner.log"] # add
  volumeMounts:                                    # add
  - name: logs                                     # add
    mountPath: /var/log/cleaner                    # add
```

Then apply the changes and check the logs of the sidecar:

```
$ k -f /opt/course/16/cleaner-new.yaml apply
```

This will cause a deployment rollout of which we can get more details:

```
$ k -n mercury rollout history deploy cleaner
$ k -n mercury rollout history deploy cleaner --revision 1
$ k -n mercury rollout history deploy cleaner --revision 2
```

Check Pod statuses:

```
$ k -n mercury get pod
```

NAME	READY	STATUS	RESTARTS	AGE
cleaner-86b7758668-9pw6t	2/2	Running	0	6s
cleaner-86b7758668-qgh4v	0/2	Init:0/1	0	1s

```
$ k -n mercury get pod
```

NAME	READY	STATUS	RESTARTS	AGE
cleaner-86b7758668-9pw6t	2/2	Running	0	14s
cleaner-86b7758668-qgh4v	2/2	Running	0	9s

Finally check the logs of the logging sidecar container:

```
$ k -n mercury logs cleaner-576967576c-cqtgx -c logger-con
```

```
init
Wed Sep 11 10:45:44 UTC 2099: remove random file
Wed Sep 11 10:45:45 UTC 2099: remove random file
...
```

Mystery solved, something is removing files at random. It's important to understand how containers can communicate with each other using volumes.

### Q17 | InitContainer

Last lunch you told your coworker from department Mars Inc how amazing InitContainers are. Now he would like to see one in action. There is a Deployment yaml at **/opt/course/17/test-init-container.yaml**. This Deployment spins up a single Pod of image `nginx:1.17.3-alpine` and serves files from a mounted volume, which is empty right now.

Create an InitContainer named `init-con` which also mounts that volume and creates a file `index.html` with content `check this out!` in the root of the mounted volume. For this test we ignore that it doesn't contain valid html.

The InitContainer should be using image **busybox:1.31.0**. Test your implementation for example using `curl` from a temporary **nginx:alpine** Pod.

**Ans.**

```
$ cp /opt/course/17/test-init-container.yaml ~/17_test-init-container.yaml
$ vim 17_test-init-container.yaml
```

Add the InitContainer:

```
# 17_test-init-container.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: test-init-container
```

```
namespace: mars
spec:
  replicas: 1
  selector:
    matchLabels:
      id: test-init-container
  template:
    metadata:
      labels:
        id: test-init-container
    spec:
      volumes:
        - name: web-content
          emptyDir: {}
      initContainers:
        - name: init-con
          image: busybox:1.31.0
          command: ['sh', '-c', 'echo "check this out!" > /tmp/web-content/index.html']
          volumeMounts:
            - name: web-content
              mountPath: /tmp/web-content # initContainer end
      containers:
        - image: nginx:1.17.3-alpine
          name: nginx
          volumeMounts:
            - name: web-content
              mountPath: /usr/share/nginx/html
          ports:
            - containerPort: 80
```

Then we create the Deployment:

```
$ k -f 17_test-init-container.yaml create
```

Finally, we test the configuration:

```
$ k -n mars get pod -o wide # to get the cluster IP
```

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl 10.0.0.67
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
---------	------------	---------	---------------	------	------	------	---------

Dload Upload Total Spent Left Speed  
check this out!

### Q18 | Service misconfiguration

There seems to be an issue in Namespace mars where the ClusterIP service **manager-api-svc** should make the Pods of Deployment **manager-api-deployment** available inside the cluster.

You can test this with curl **manager-api-svc.mars:4444** from a temporary **nginx:alpine** Pod. Check for the misconfiguration and apply a fix.

**Ans.**

First let's get an overview:

```
$ k -n mars get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/manager-api-deployment-dbcc6657d-bg2hh	1/1	Running	0	98m
pod/manager-api-deployment-dbcc6657d-f5fv4	1/1	Running	0	98m
pod/manager-api-deployment-dbcc6657d-httjv	1/1	Running	0	98m
pod/manager-api-deployment-dbcc6657d-k98xn	1/1	Running	0	98m
pod/test-init-container-5db7c99857-htx6b	1/1	Running	0	2m19s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/manager-api-svc	ClusterIP	10.15.241.159	<none>	4444/TCP	99m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/manager-api-deployment	4/4	4	4	98m
deployment.apps/test-init-container	1/1	1	1	2m19s

...

Everything seems to be running, but we can't seem to get a connection:

```
$ k -n mars run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc:4444
```

If you don't see a command prompt, try pressing enter.

```
0 0 0 0 0 0 0 0 --:--:-- 0:00:01 --:--:-- 0
curl: (28) Connection timed out after 1000 milliseconds
pod "tmp" deleted
pod mars/tmp terminated (Error)
```

Ok, let's try to connect to one pod directly:

```
$ k -n mars get pod -o wide # get cluster IP
```

```
$ k -n mars run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 10.0.1.14
```

```
% Total % Received % Xferd Average Speed Time Time Time Current
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

The Pods itself seem to work. Let's investigate the Service a bit:

```
$ k -n mars describe service manager-api-svc
```

```
Name:          manager-api-svc
Namespace:     mars
Labels:        app=manager-api-svc
...
Endpoints:     <none>
...
```

Endpoint inspection is also possible using:

```
$ k -n mars get ep
```

No endpoints - No good. We check the Service yaml:

```
$ k -n mars edit service manager-api-svc
```

```
apiVersion: v1
kind: Service
metadata:
...
labels:
  app: manager-api-svc
  name: manager-api-svc
  namespace: mars
...
spec:
  clusterIP: 10.3.244.121
  ports:
    - name: 4444-80
```



```
port: 4444
protocol: TCP
targetPort: 80
selector:
  #id: manager-api-deployment # wrong selector, needs to point to pod!
  id: manager-api-pod
sessionAffinity: None
type: ClusterIP
```

Though Pods are usually never created without a Deployment or ReplicaSet, Services always select for Pods directly. This gives great flexibility because Pods could be created through various customized ways. After saving the new selector we check the Service again for endpoints:

```
$ k -n mars get ep
```

NAME	ENDPOINTS	AGE
manager-api-svc	10.0.0.30:80,10.0.1.30:80,10.0.1.31:80 + 1 more...	41m

Endpoints - Good! Now we try connecting again:

```
$ k -n mars run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc:4444
```

```
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left  Speed
100 612 100 612    0     0  99k    0 --:--:-- --:--:-- --:--:--  99k
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

And we fixed it. Good to know is how to be able to use Kubernetes DNS resolution from a different Namespace. Not necessary, but we could spin up the temporary Pod in default Namespace:

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc:4444
```

```
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left  Speed
0   0   0   0    0     0    0    0 --:--:-- --:--:-- --:--:--    0curl: (6) Could not resolve host:
manager-api-svc
pod "tmp" deleted
pod default/tmp terminated (Error)
```

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc.mars:4444
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload Total Spent Left Speed				
100	612	100	612	0	0	68000	0 --:--:-- --:--:-- --:--:-- 68000

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

Short **manager-api-svc.mars** or long **manager-api-svc.mars.svc.cluster.local** work.

### Q19 | Service ClusterIP->NodePort

In Namespace **jupiter** you'll find an apache Deployment (with one replica) named **jupiter-crew-deploy** and a ClusterIP Service called **jupiter-crew-svc** which exposes it. Change this service to a NodePort one to make it available on all nodes on port 30100.

Test the NodePort Service using the internal IP of all available nodes and the port 30100 using curl, you can reach the internal node IPs directly from your main terminal. On which nodes is the Service reachable? On which node is the Pod running?

**Ans.**

First, we get an overview:

```
$ k -n jupiter get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/jupiter-crew-deploy-8cdf99bc9-klwqt	1/1	Running	0	34m

  

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/jupiter-crew-svc	ClusterIP	10.100.254.66	<none>	8080/TCP	34m

...

(Optional) Next, we check if the ClusterIP Service actually works:

```
$ k -n jupiter run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 jupiter-crew-svc:8080
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload Total Spent Left Speed				
100	45	100	45	0	0	5000	0 --:--:-- --:--:-- --:--:-- 5000

```
<html><body><h1>It works!</h1></body></html>
```

The Service is working great. Next we change the Service type to NodePort and set the port:

```
$ k -n jupiter edit service jupiter-crew-svc
```

```
apiVersion: v1
kind: Service
metadata:
  name: jupiter-crew-svc
  namespace: jupiter
...
spec:
  clusterIP: 10.3.245.70
  ports:
    - name: 8080-80
      port: 8080
      protocol: TCP
      targetPort: 80
      nodePort: 30100 # add the nodePort
  selector:
    id: jupiter-crew
  sessionAffinity: None
  #type: ClusterIP
  type: NodePort # change type
status:
  loadBalancer: {}
```

We check if the Service type was updated:

```
$ k -n jupiter get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
jupiter-crew-svc	NodePort	10.3.245.70	<none>	8080:30100/TCP	3m52s

(Optional) And we confirm that the service is still reachable internally:

```
$ k -n jupiter run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 jupiter-crew-svc:8080
```

```
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload Total   Spent    Left  Speed
<html><body><h1>It works!</h1></body></html>
```

Nice. A NodePort Service kind of lies on top of a ClusterIP one, making the ClusterIP Service reachable on the Node IPs (internal and external). Next we get the internal IPs of all nodes to check the connectivity:

```
$ k get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	...
cluster1-master1	Ready	master	18h	v1.21.0	192.168.100.11	...
cluster1-worker1	Ready	<none>	18h	v1.21.0	192.168.100.12	...

On which nodes is the Service reachable?

```
$ curl 192.168.100.11:30100
```

```
<html><body><h1>It works!</h1></body></html>
```

```
$ curl 192.168.100.12:30100
```

```
<html><body><h1>It works!</h1></body></html>
```

On both, even the master. On which node is the Pod running?

```
$ k -n jupiter get pod jupiter-crew-deploy-8cdf99bc9-klwqt -o yaml | grep nodeName
```

```
nodeName: cluster1-worker1
```

```
$ k -n jupiter get pod -o wide # or even shorter
```

In our case on cluster1-worker1, but could be any other worker if more available. Here we hopefully gained some insight into how a NodePort Service works. Although the Pod is just running on one specific node, the Service makes it available through port 30100 on the internal and external IP addresses of all nodes. This is at least the common/default behaviour but can depend on cluster configuration.

## Q20 | Network Policy

In Namespace venus you'll find two Deployments named api and frontend. Both Deployments are exposed inside the cluster using Services. Create a NetworkPolicy named np1 which restricts outgoing tcp connections from Deployment frontend and only allows those going to Deployment api. Make sure the NetworkPolicy still allows outgoing traffic on UDP/TCP ports 53 for DNS resolution.

Test using: `wget www.google.com` and `wget api:2222` from a Pod of Deployment frontend.

**Ans.**

First, we get an overview:

```
k -n venus get all
```

```
NAME                READY STATUS RESTARTS AGE
pod/api-5979b95578-gktxp    1/1   Running 0       57s
pod/api-5979b95578-lhcl5    1/1   Running 0       57s
pod/frontend-789cbdc677-c9v8h 1/1   Running 0       57s
pod/frontend-789cbdc677-npk2m 1/1   Running 0       57s
pod/frontend-789cbdc677-pl67g 1/1   Running 0       57s
pod/frontend-789cbdc677-rjt5r 1/1   Running 0       57s
pod/frontend-789cbdc677-xgf5n 1/1   Running 0       57s

NAME                TYPE        CLUSTER-IP   EXTERNAL-IP  PORT(S)  AGE
service/api          ClusterIP   10.3.255.137 <none>       2222/TCP 37s
service/frontend     ClusterIP   10.3.255.135 <none>       80/TCP   57s
...
```

(Optional) This is not necessary but we could check if the Services are working inside the cluster:

```
k -n venus run tmp --restart=Never --rm -i --image=busybox -i -- wget -O- frontend:80
```

```
Connecting to frontend:80 (10.3.245.9:80)
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

```
$ k -n venus run tmp --restart=Never --rm --image=busybox -i -- wget -O- api:2222
```

```
Connecting to api:2222 (10.3.250.233:2222)
<html><body><h1>It works!</h1></body></html>
```

Then we use any frontend Pod and check if it can reach external names and the api Service:

```
$ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- www.google.com
```

```
Connecting to www.google.com (216.58.205.227:80)
- 100% |*****| 12955 0:00:00 ETA
<!doctype html><html itemscope="" itemtype="http://schema.org/WebPage" lang="en"><head>
...
```

```
$ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- api:2222
```

```
<html><body><h1>It works!</h1></body></html>
Connecting to api:2222 (10.3.255.137:2222)
- 100% |*****| 45 0:00:00 ETA
...
```

We see Pods of frontend can reach the api and external names.

```
$ vim 20_np1.yaml
```

Now we head to <https://kubernetes.io/docs> , search for NetworkPolicy, copy the example code and adjust it to:

```
# 20_np1.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: np1
  namespace: venus
spec:
  podSelector:
    matchLabels:
      id: frontend      # label of the pods this policy should be applied on
  policyTypes:
  - Egress              # we only want to control egress
  egress:
  - to:                 # 1st egress rule
    - podSelector:      # allow egress only to pods with api label
      matchLabels:
        id: api
  - ports:              # 2nd egress rule
    - port: 53           # allow DNS UDP
      protocol: UDP
    - port: 53           # allow DNS TCP
      protocol: TCP
```

Notice that we specify two egress rules in the yaml above. If we specify multiple egress rules then these are connected using a logical OR. So in the example above we do:

```
allow outgoing traffic if
(destination pod has label id:api) OR ((port is 53 UDP) OR (port is 53 TCP))
```

Let's have a look at example code which wouldn't work in our case:

```
# this example does not work in our case
...
egress:
- to:          # 1st AND ONLY egress rule
  podSelector: # allow egress only to pods with api label
    matchLabels:
      id: api
  ports:       # STILL THE SAME RULE but just an additional selector
    - port: 53 # allow DNS UDP
      protocol: UDP
    - port: 53 # allow DNS TCP
      protocol: TCP
```

In the yaml above we only specify one egress rule with two selectors. It can be translated into:

allow outgoing traffic if  
(destination pod has label id:api) AND ((port is 53 UDP) OR (port is 53 TCP))

```
$ k -f 20_np1.yaml create
```

And try again, external is not working any longer:

```
k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- www.google.de
```

```
Connecting to www.google.de:2222 (216.58.207.67:80)
^C
```

```
$ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- -T 5 www.google.de:80
```

```
Connecting to www.google.com (172.217.203.104:80)
wget: download timed out
command terminated with exit code 1
```

Internal connection to api work as before:

```
$ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- api:2222
```

```
<html><body><h1>It works!</h1></body></html>
Connecting to api:2222 (10.3.255.137:2222)
- 100% |*****| 45 0:00:00 ETA
```

## Preview Question 1

In Namespace pluto there is a Deployment named project-23-api. It has been working okay for a while but Team Pluto needs it to be more reliable. Implement a liveness-probe which checks the container to be reachable on port 80. Initially the probe should wait 10, periodically 15 seconds.

The original Deployment yaml is available at /opt/course/p1/project-23-api.yaml. Save your changes at /opt/course/p1/project-23-api-new.yaml and apply the changes.

**Ans.**

First, we get an overview:

```
$ k -n pluto get all -o wide
```

```
NAME                                READY STATUS ... IP      ...
pod/holy-api                        1/1   Running ... 10.12.0.26 ...
pod/project-23-api-784857f54c-dx6h6 1/1   Running ... 10.12.2.15 ...
pod/project-23-api-784857f54c-sj8df 1/1   Running ... 10.12.1.18 ...
pod/project-23-api-784857f54c-t4xmh 1/1   Running ... 10.12.0.23 ...
```

```
NAME                                READY UP-TO-DATE AVAILABLE ...
deployment.apps/project-23-api 3/3    3            3            ...
```

To note: we see another Pod here called holy-api which is part of another section. This is often the case in the provided scenarios, so be careful to only manipulate the resources you need to. Just like in the real world and in the exam.

Next we use nginx:alpine and curl to check if one Pod is accessible on port 80:

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 10.12.2.15
```

```
% Total    % Received % Xferd Average Speed   Time    Time     Time Current
          Dload  Upload   Total   Spent    Left   Speed
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

We could also use busybox and wget for this:



```
$ k run tmp --restart=Never --rm --image=busybox -i -- wget -O- 10.12.2.15
```

```
Connecting to 10.12.2.15 (10.12.2.15:80)
writing to stdout
-          100% |*****| 612 0:00:00 ETA
written to stdout
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

Now that we're sure the Deployment works we can continue with altering the provided yaml:

```
$ cp /opt/course/p1/project-23-api.yaml /opt/course/p1/project-23-api-new.yaml
```

```
$ vim /opt/course/p1/project-23-api-new.yaml
```

Add the liveness-probe to the yaml:

```
# /opt/course/p1/project-23-api-new.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: project-23-api
  namespace: pluto
spec:
  replicas: 3
  selector:
    matchLabels:
      app: project-23-api
  template:
    metadata:
      labels:
        app: project-23-api
    spec:
      volumes:
        - name: cache-volume1
          emptyDir: {}
        - name: cache-volume2
          emptyDir: {}
        - name: cache-volume3
          emptyDir: {}
      containers:
        - image: httpd:2.4-alpine
```

```
name: httpd
volumeMounts:
- mountPath: /cache1
  name: cache-volume1
- mountPath: /cache2
  name: cache-volume2
- mountPath: /cache3
  name: cache-volume3
env:
- name: APP_ENV
  value: "prod"
- name: APP_SECRET_N1
  value: "IO=a4L/XkRdvN8jM=Y+"
- name: APP_SECRET_P1
  value: "-7PA0_Z]>{pwa43r}__"
livenessProbe:      # add
tcpSocket:          # add
  port: 80           # add
initialDelaySeconds: 10 # add
periodSeconds: 15     # add
```

Then let's apply the changes:

```
$ k -f /opt/course/p1/project-23-api-new.yaml apply
```

Next we wait 10 seconds and confirm the Pods are still running:

```
$ k -n pluto get pod
```

NAME	READY	STATUS	RESTARTS	AGE
holy-api	1/1	Running	0	144m
project-23-api-5b4579fd49-8knh8	1/1	Running	0	90s
project-23-api-5b4579fd49-cbgph	1/1	Running	0	88s
project-23-api-5b4579fd49-tcfq5	1/1	Running	0	86s

We can also check the configured liveness-probe settings on a Pod or the Deployment:

```
$ k -n pluto describe pod project-23-api-5b4579fd49-8knh8 | grep Liveness
```

```
Liveness:  tcp-socket :80 delay=10s timeout=1s period=15s #success=1 #failure=3
```

```
$ k -n pluto describe deploy project-23-api | grep Liveness
```

```
Liveness: tcp-socket :80 delay=10s timeout=1s period=15s #success=1 #failure=3
```

## Preview Question 2

Team Sun needs a new Deployment named sunny with 4 replicas of image nginx:1.17.3-alpine in Namespace sun. The Deployment and its Pods should use the existing ServiceAccount sa-sun-deploy.

Expose the Deployment internally using a ClusterIP Service named sun-srv on port 9999. The nginx containers should run as default on port 80. The management of Team Sun would like to execute a command to check that all Pods are running on occasion. Write that command into file /opt/course/p2/sunny\_status\_command.sh. The command should use kubectl.

**Ans.**

```
$ k -n sun create deployment -h #help
```

```
# check the export on the very top of this document so we can use $do
```

```
$ k -n sun create deployment sunny --image=nginx:1.17.3-alpine $do > p2_sunny.yaml
```

```
$ vim p2_sunny.yaml
```

Then alter its yaml to include the requirements:

```
# p2_sunny.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: sunny
    name: sunny
    namespace: sun
spec:
  replicas: 4                # change
  selector:
    matchLabels:
```

```

  app: sunny
strategy: {}
template:
  metadata:
    creationTimestamp: null
    labels:
      app: sunny
  spec:
    serviceAccountName: sa-sun-deploy # add
    containers:
      - image: nginx:1.17.3-alpine
        name: nginx
        resources: {}
status: {}

```

Now create the yaml and confirm its running:

```

k create -f p2_sunny.yaml
deployment.apps/sunny created

```

```

$ k -n sun get pod

```

NAME	READY	STATUS	RESTARTS	AGE
0509649a	1/1	Running	0	149m
0509649b	1/1	Running	0	149m
1428721e	1/1	Running	0	149m
...				
sunny-64df8dbdbb-9mxbw	1/1	Running	0	10s
sunny-64df8dbdbb-mp5cf	1/1	Running	0	10s
sunny-64df8dbdbb-pggdf	1/1	Running	0	6s
sunny-64df8dbdbb-zvqth	1/1	Running	0	7s

Confirmed, the AGE column is always in important information about if changes were applied. Next we expose the Pods by created the Service:

```

$ k -n sun expose -h # help
$ k -n sun expose deployment sunny --name sun-srv --port 9999 --target-port 80

```

Using expose instead of kubectl create service clusterip is faster because it already sets the correct selector-labels. The previous command would produce this yaml:

```
# k -n sun expose deployment sunny --name sun-srv --port 9999 --target-port 80
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: sunny
  name: sun-srv      # required by task
spec:
  ports:
    - port: 9999      # service port
      protocol: TCP
      targetPort: 80   # target port
  selector:
    app: sunny        # selector is important
status:
  loadBalancer: {}
```

Let's test the Service using wget from a temporary Pod:

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 sun-srv.sun:9999
```

```
Connecting to sun-srv.sun:9999 (10.23.253.120:9999)
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

Because the Service is in a different Namespace as our temporary Pod, it is reachable using the names sun-srv.sun or fully: sun-srv.sun.svc.cluster.local.

Finally we need a command which can be executed to check if all Pods are running, this can be done with:

```
$ vim /opt/course/p2/sunny_status_command.sh
```

```
# /opt/course/p2/sunny_status_command.sh
$ kubectl -n sun get deployment sunny
```

To run the command:

```
$ sh /opt/course/p2/sunny_status_command.sh
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
sunny	4/4	4	4	13m

### Preview Question 3

Management of EarthAG recorded that one of their Services stopped working. Dirk, the administrator, left already for the long weekend. All the information they could give you is that it was located in Namespace earth and that it stopped working after the latest rollout. All Services of EarthAG should be reachable from inside the cluster.

Find the Service, fix any issues and confirm its working again. Write the reason of the error into file /opt/course/p3/ticket-654.txt so Dirk knows what the issue was.

**Ans.**

First we get an overview of the resources in Namespace earth:

```
$ k -n earth get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/earth-2x3-api-584df69757-ngnwp	1/1	Running	0	116m
pod/earth-2x3-api-584df69757-ps8cs	1/1	Running	0	116m
pod/earth-2x3-api-584df69757-ww9q8	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-48vjt	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-6mqmb	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-6vjll	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-fnkbp	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-pjm5m	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-pwfvj	1/1	Running	0	116m
pod/earth-3cc-runner-6cb6cc6974-8wm5x	1/1	Running	0	116m
pod/earth-3cc-runner-6cb6cc6974-9fx8b	1/1	Running	0	116m
pod/earth-3cc-runner-6cb6cc6974-b9nrv	1/1	Running	0	116m
pod/earth-3cc-runner-heavy-6bf876f46d-b47vq	1/1	Running	0	116m
pod/earth-3cc-runner-heavy-6bf876f46d-mrzqd	1/1	Running	0	116m
pod/earth-3cc-runner-heavy-6bf876f46d-qkd74	1/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-f74cj	0/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-n4z7z	0/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-rcmxs	0/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-xl467	0/1	Running	0	116m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/earth-2x3-api-svc	ClusterIP	10.3.241.242	<none>	4546/TCP	116m
service/earth-2x3-web-svc	ClusterIP	10.3.250.247	<none>	4545/TCP	116m
service/earth-3cc-web	ClusterIP	10.3.243.24	<none>	6363/TCP	116m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/earth-2x3-api	3/3	3	3	116m
deployment.apps/earth-2x3-web	6/6	6	6	116m
deployment.apps/earth-3cc-runner	3/3	3	3	116m
deployment.apps/earth-3cc-runner-heavy	3/3	3	3	116m
deployment.apps/earth-3cc-web	0/4	4	0	116m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/earth-2x3-api-584df69757	3	3	3	116m
replicaset.apps/earth-2x3-web-85c5b7986c	6	6	6	116m
replicaset.apps/earth-3cc-runner-6cb6cc6974	3	3	3	116m
replicaset.apps/earth-3cc-runner-heavy-6bf876f46d	3	3	3	116m
replicaset.apps/earth-3cc-web-6895587dc7	0	0	0	116m
replicaset.apps/earth-3cc-web-6bfdf8b848	4	4	0	116m
replicaset.apps/earth-3cc-web-d49645966	0	0	0	116m

First impression could be that all Pods are in status RUNNING. But looking closely we see that some of the Pods are not ready, which also confirms what we see about one Deployment and one replicaset. This could be our error to further investigate.

Another approach could be to check the Services for missing endpoints:

```
k -n earth get ep
```

NAME	ENDPOINTS	AGE
earth-2x3-api-svc	10.0.0.10:80,10.0.1.5:80,10.0.2.4:80	116m
earth-2x3-web-svc	10.0.0.11:80,10.0.0.12:80,10.0.1.6:80 + 3 more...	116m
earth-3cc-web		

Service earth-3cc-web doesn't have endpoints. This could be a selector/label misconfiguration or the endpoints are actually not available/ready.

Checking all Services for connectivity should show the same (this step is optional and just for demonstration:

```
k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-2x3-api-svc.earth:4546
```

```
...
<html><body><h1>It works!</h1></body></html>
```

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-2x3-web-svc.earth:4545
```

```
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100  45 100  45  0  0 5000    0 --:--:-- --:--:-- --:--:-- 5000
<html><body><h1>It works!</h1></body></html>
```

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-3cc-web.earth:6363
```

```
If you don't see a command prompt, try pressing enter.
 0  0  0  0  0  0  0  0  0 --:--:-- 0:00:05 --:--:-- 0
curl: (28) Connection timed out after 5000 milliseconds
pod "tmp" deleted
pod default/tmp terminated (Error)
```

Notice that we use here for example earth-2x3-api-svc.earth. We could also spin up a temporary Pod in Namespace earth and connect directly to earth-2x3-api-svc.

We get no connection to earth-3cc-web.earth:6363. Let's look at the Deployment earth-3cc-web. Here we see that the requested amount of replicas is not available/ready:

```
$ k -n earth get deploy earth-3cc-web
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
earth-3cc-web	0/4	4	0	7m18s

To continue we check the Deployment yaml for some misconfiguration:

```
$ k -n earth edit deploy earth-3cc-web
```

```
# k -n earth edit deploy earth-3cc-web
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
...
  generation: 3          # there have been rollouts
  name: earth-3cc-web
  namespace: earth
...
spec:
```



```
...
template:
  metadata:
    creationTimestamp: null
    labels:
      id: earth-3cc-web
  spec:
    containers:
      - image: nginx:1.16.1-alpine
        imagePullPolicy: IfNotPresent
        name: nginx
        readinessProbe:
          failureThreshold: 3
          initialDelaySeconds: 10
          periodSeconds: 20
          successThreshold: 1
          tcpSocket:
            port: 82          # this port doesn't seem to be right, should be 80
          timeoutSeconds: 1
...

```

We change the readiness-probe port, save and check the Pods:

```
$ k -n earth get pod -l id=earth-3cc-web
```

NAME	READY	STATUS	RESTARTS	AGE
earth-3cc-web-d49645966-52vb9	0/1	Running	0	6s
earth-3cc-web-d49645966-5tts6	0/1	Running	0	6s
earth-3cc-web-d49645966-db5gp	0/1	Running	0	6s
earth-3cc-web-d49645966-mk7gr	0/1	Running	0	6s

Running, but still not in ready state. Wait 10 seconds (initialDelaySeconds of readinessProbe) and check again:

```
$ k -n earth get pod -l id=earth-3cc-web
```

NAME	READY	STATUS	RESTARTS	AGE
earth-3cc-web-d49645966-52vb9	1/1	Running	0	32s
earth-3cc-web-d49645966-5tts6	1/1	Running	0	32s
earth-3cc-web-d49645966-db5gp	1/1	Running	0	32s
earth-3cc-web-d49645966-mk7gr	1/1	Running	0	32s

Let's check the service again:

```
$ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-3cc-web.earth:6363
```

```
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload  Total   Spent    Left   Speed
100 612 100 612    0     0 55636    0 --:--:-- --:--:-- --:--:-- 55636
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
...
```

We did it! Finally we write the reason into the requested location:

```
$ vim /opt/course/p3/ticket-654.txt
```

```
# /opt/course/p3/ticket-654.txt
yo Dirk, wrong port for readinessProbe defined!
```